




# ATOMICS, REDUCTIONS, WARP SHUFFLE

Guillaume Barnier, Solutions Architect, Energy  
[gbarnier@nvidia.com](mailto:gbarnier@nvidia.com)




# AGENDA

- Transformations vs. reductions, thread strategy
- Atomics, atomic reductions
- Atomic tips and tricks
- Classical parallel reduction
- Parallel reduction + atomics
- Warp shuffle, reduction with warp shuffle



**ATOMICS**





# MOTIVATING EXAMPLE


## Sum - Reduction

```
const int size = 100000;  
float a[size] = {...};  
float sum = 0;  
for (int i = 0; i < size; i++)  
    sum += a[i];
```

Goal: sum variable contains the sum of all the elements of array a

# TRANSFORMATION VS. REDUCTION


May guide the thread strategy: what will each thread do?



Transformation:

$$c[i] = a[i] + 10;$$

**Thread strategy:** one thread per output point



Reduction:

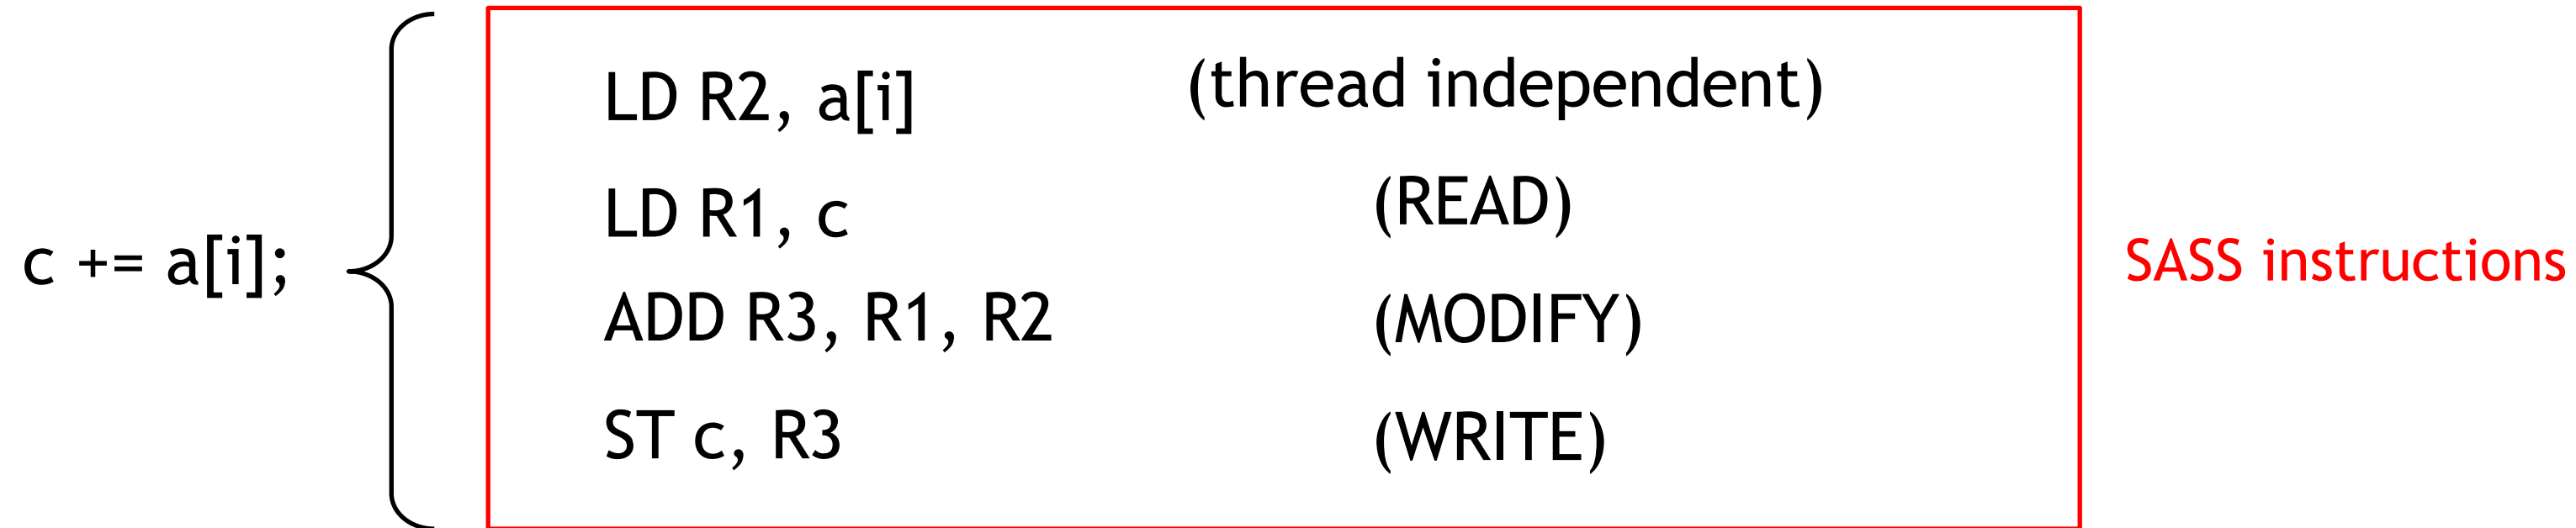
$$c = \sum a[i]$$

**Thread strategy:** ?

# REDUCTION: NAÏVE THREAD STRATEGY

One thread per input point

Doesn't work



- But **every thread** is trying to do this, potentially at the same time
- The CUDA programming model does not enforce any order of thread execution

# ATOMIC TO THE RESCUE

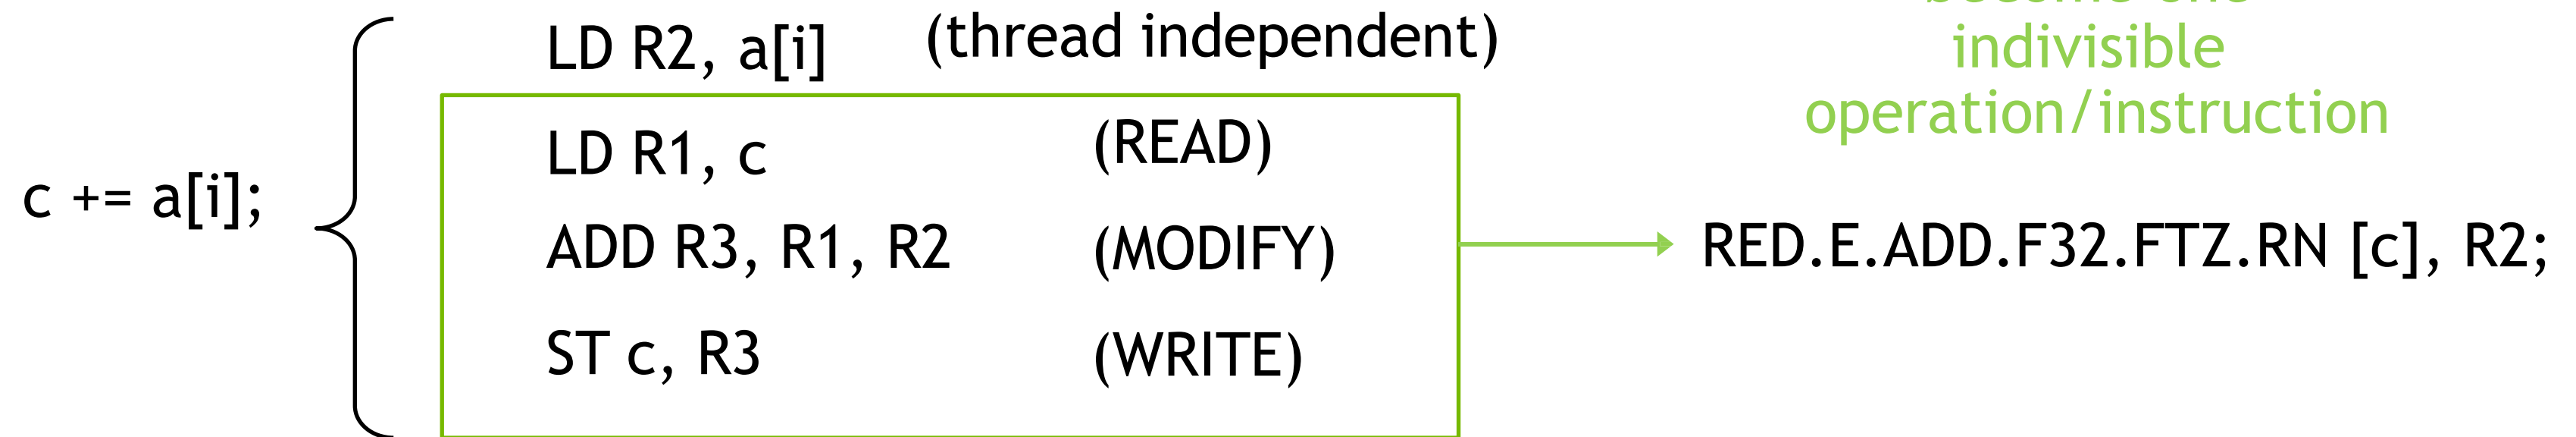


“Indivisible”

# ATOMIC TO THE RESCUE

## Indivisible READ-MODIFY-WRITE

`atomicAdd(&c, a[i]);`



- Facilitated by special hardware in the L2 cache
- May have performance implications



# OTHER ATOMICS

- ▶ `atomicMax/Min` - choose the max (or min)
- ▶ `atomicAdd/Sub` - add to (or subtract from)
- ▶ `atomicInc/Dec` - increment (or decrement) and account for rollover/underflow
- ▶ `atomicExch/CAS` - swap values, or conditionally swap values
- ▶ `atomicAnd/Or/Xor` - bitwise ops
- ▶ atomics have different datatypes they can work on (e.g. int, unsigned, float, etc.)
- ▶ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

# ATOMIC TIPS AND TRICKS

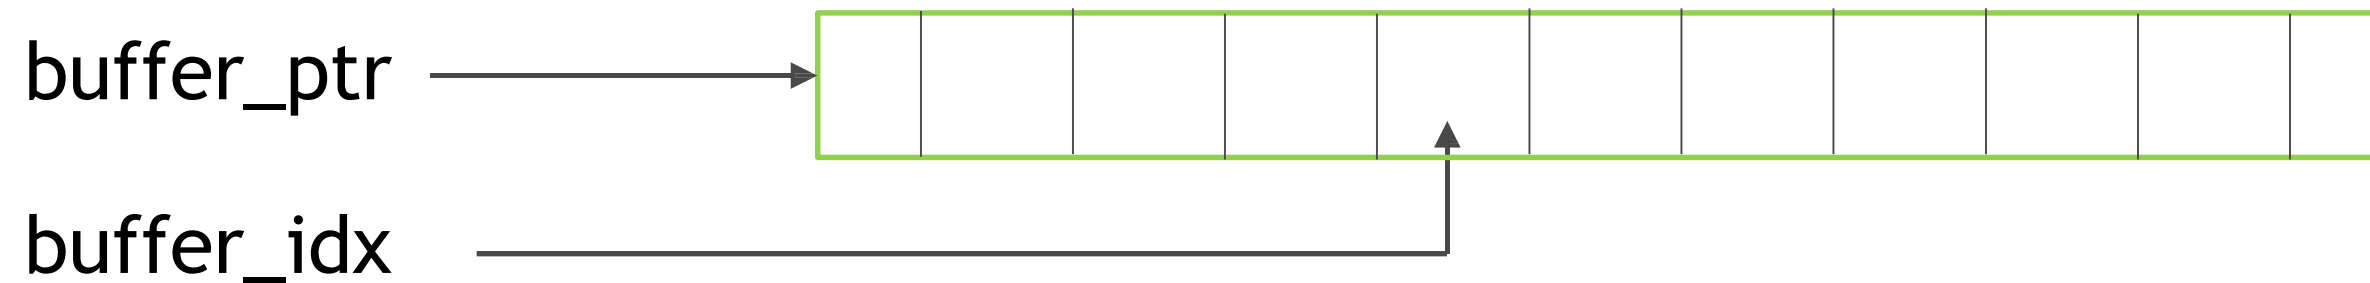
Determine my place in an order

- Could be used to determine next work item, queue slot, etc.
- `int my_position = atomicAdd(order, 1);`
- Most atomics return a value that is the “old” value that was in the location receiving the atomic update

# ATOMIC TIPS AND TRICKS

## Reserve space in a buffer


Each thread in my kernel may produce a variable amount of data. How to collect all of this in one buffer, in parallel?



```
int my_length = var; // Length of array  
float local_buffer[my_length] = {...}; // Allocate and initialize array  
int my_offset = atomicAdd(buffer_idx, my_length); // Increment buffer_idx value, save position of first element  
// buffer_ptr + my_offset now points to the first reserved location, of length "my_length"  
memcpy(buffer_ptr + my_offset, local_buffer, my_length * sizeof(float));
```




# CLASSICAL PARALLEL REDUCTION



# THE CLASSICAL PARALLEL REDUCTION

Atoms don't run at full memory bandwidth

- We would like a reduction method that is not limited by atomic throughput
- We would like to effectively use all threads, as much as possible
- Parallel reduction is a common and important data parallel primitive Naïve
- implementations will often run into bottlenecks
- Basic methodology is a tree-based approach



# PROBLEM: GLOBAL SYNCHRONIZATION

If we could synchronize across all thread blocks, could easily reduce very large arrays, right?

- Global sync after each block produces its result
- Once all blocks reach sync, continue recursively

One possible solution: decompose into multiple kernels


- Kernel launch serves as a global synchronization point
- Kernel launch has low overhead (but not zero)

Other possible solutions:

- Use atomics at the end of threadblock-level reduction
- Use a threadblock-draining approach (see threadFenceReduction sample code)
- Use cooperative groups - cooperative kernel launch



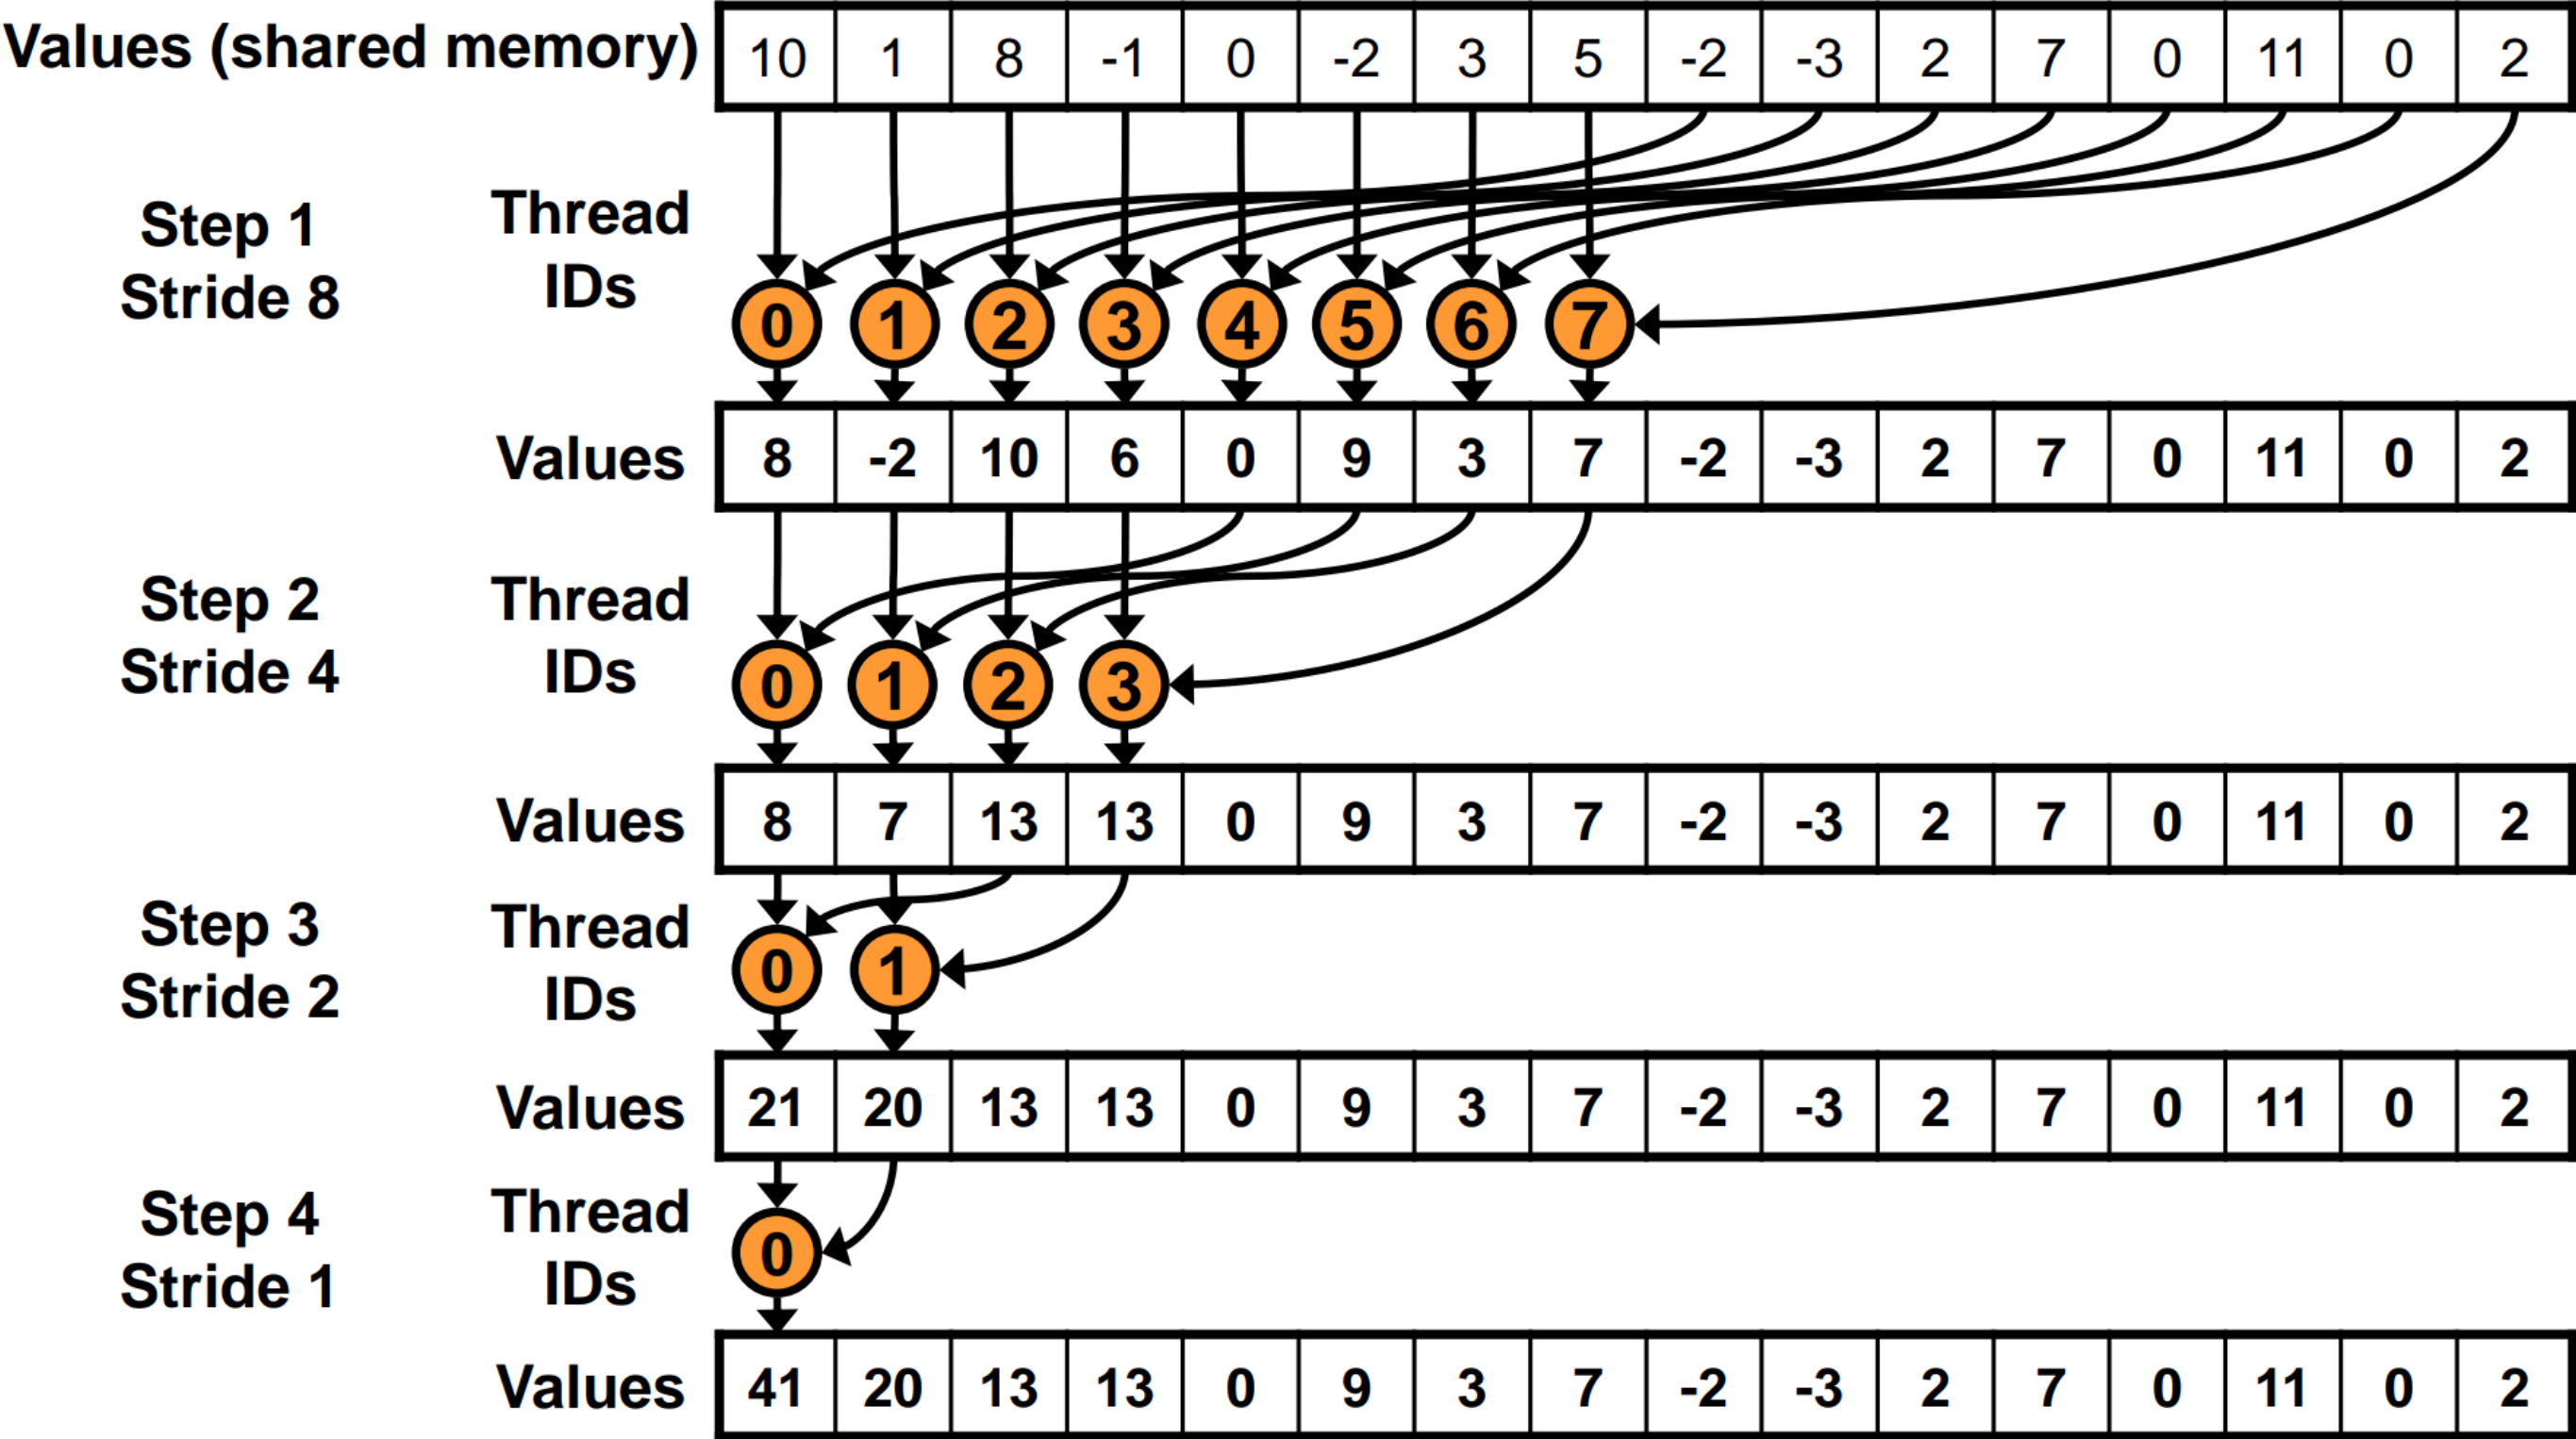
# SOLUTION: KERNEL DECOMPOSITION



- Create global sync by decomposing computations into multiple kernel invocations
- In the case of reductions, code for all levels (invocations) is the same

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
  if (tid < s) {
    sdata[tid] += sdata[tid + s]; }
  __syncthreads(); // outside the if-statement
}
```

# SEQUENTIAL ADDRESSING



Sequential  
addressing is  
bank-conflict  
free

# DETOUR: GRID-STRIDE LOOPS

- ▶ We'd like to be able to design kernels that load and operate on arbitrary data sizes efficiently
- ▶ Want to maintain coalesced loads/stores, efficient use of shared memory
- ▶ Can also be used for ninja-level tuning - choose number of blocks sized to the GPU

▶ `gdata[0..N-1]`:



| grid-width stride | grid-width stride | grid-width stride ...



```
int idx = threadIdx.x+blockDim.x*blockIdx.x;
float sum = 0.f;
while (idx < N) {
    sum += gdata[idx];
    idx += gridDim.x*blockDim.x; // grid width
}
sdata[threadIdx.x] = sum;
```

# PUTTING IT ALL TOGETHER

```
__global__ void reduce(float *gdata, float *out)
{
    __shared__ float sdata[BLOCK_SIZE];
    int tid = threadIdx.x;
    size_t idx = threadIdx.x + blockDim.x * blockIdx.x;

    // Grid stride loop to load data into shared memory
    float sum = 0.f;
    while (idx < N) {
        sum += gdata[idx];
        idx += blockDim.x * blockDim.x;
    }
    sdata[tid] = sum; // Load temporary sum into share mem

    // Parallel sweep reduction using shared memory
    for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
        __syncthreads();
        if (tid < s) sdata[tid] += sdata[tid + s];
    }

    // Thread 0 writes total sum
    if (tid == 0)
        out[blockIdx.x] = sdata[0];
}
```

# GETTING RID OF 2<sup>ND</sup> KERNEL LAUNCH


```
__global__ void reduce(float *gdata, float *out)
{
    __shared__ float sdata[BLOCK_SIZE];
    int tid = threadIdx.x;
    size_t idx = threadIdx.x + blockDim.x * blockIdx.x;

    // Grid stride loop to load data into shared memory
    float sum = 0.f;
    while (idx < N) {
        sum += gdata[idx];
        idx += blockDim.x * blockDim.x;
    }
    sdata[tid] = sum; // Load temporary sum into share mem


    // Parallel sweep reduction using shared memory
    for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
        __syncthreads();
        if (tid < s) sdata[tid] += sdata[tid + s];
    }

    // Thread 0 writes total sum
    if (tid == 0)
        atomicAdd(out, sdata[0]);
}
```






# WARP SHUFFLE






# INTER-THREAD COMMUNICATION: SO FAR

- ▶ Using shared memory:



- ▶ Threads:


- ▶ Wouldn't this be convenient:



- ▶ Threads:

# INTRODUCING WARP SHUFFLE

- ▶ Allows for intra-warp communication
- ▶ Various supported movement patterns:
  - ▶ `__shfl_sync()`: copy from lane ID (arbitrary pattern)
  - ▶ `__shfl_xor_sync()`: copy from calculated lane ID (calculated pattern)
  - ▶ `__shfl_up_sync()`: copy from delta/offset lower lane
  - ▶ `__shfl_down_sync()`: copy from delta/offset higher lane:



- ▶ Both source and destination threads in the warp must “participate”
- ▶ Sync “mask” used to identify and reconverge needed threads

```

__global__ void reduce (float *gdata, float *out)
{
    __shared__ float sdata[32]; // Allocate shared memory
    int tid = threadIdx.x;
    sdata[tid] = 0.f;
    size_t idx = threadIdx.x + blockDim.x * blockIdx.x;
    float val = 0.f;
    unsigned int mask = ~0;
    int lane = threadIdx.x % warpSize;
    int warpID = threadIdx.x / warpSize;

    // Each thread computes sums over grid stride loops
    while (idx < N) {
        val += gdata[idx];
        idx += blockDim.x*blockDim.x;
    }

    // First warp-shuffle
    for (int offset = warpSize/2; offset > 0; offset >>= 1)
        val += __shfl_down_sync(mask, val, offset);

    // Only thread with lane = 0 write their partial sum to shared memory
    if (lane == 0)
        sdata[warpID] = val;
    __syncthreads();
}

```

```

// Only warp 0 is working
if (warpID == 0) {
    val = (tid < blockDim.x / warpSize) ? sdata[lane] : 0.f;

    // Second warp-shuffle
    for (int offset = warpSize/2; offset > 0; offset >>= 1)
        val += __shfl_down_sync(mask, val, offset);

    // Only thread with lane = 0 write their partial sum to shared memory
    if (threadIdx.x == 0)
        atomicAdd(out, val);
}

} // End kernel

```


# WARP SHUFFLE BENEFIT

- Reduce or eliminate shared memory usage
- Single instruction vs. 2 or more instructions
- Reduce level of explicit synchronization

# WARP SHUFFLE TIPS AND TRICKS

What else can we do with it?

- Broadcast a value to all threads in the warp in a single instruction
- Perform a warp-level prefix sum
- Atomic aggregation





# MANAGED MEMORY

- Good for prototyping or for porting applications quickly
- Page fault at the page level
- Managed memory -> you rely on page fault to move the data
- Page fault are expensive, and you only move one page
- One load/pagefault, then next piece of data, page fault, etc.
- The code may be slow if you rely a lot on these page faults
- We have functions such as prefetch to help reducing page faults
- Bigger pages -> fewer page faults
- Recommendation on x86: cudaMalloc and explicitly copy data