Introduction to CUDA Performance Optimization

Athena Elafrou, Guillaume Thomas Collignon, NVIDIA DevTech Compute

GPU Technology Conference, March 18th 2024

Agenda

- GPU Architecture and CUDA Programming Basics
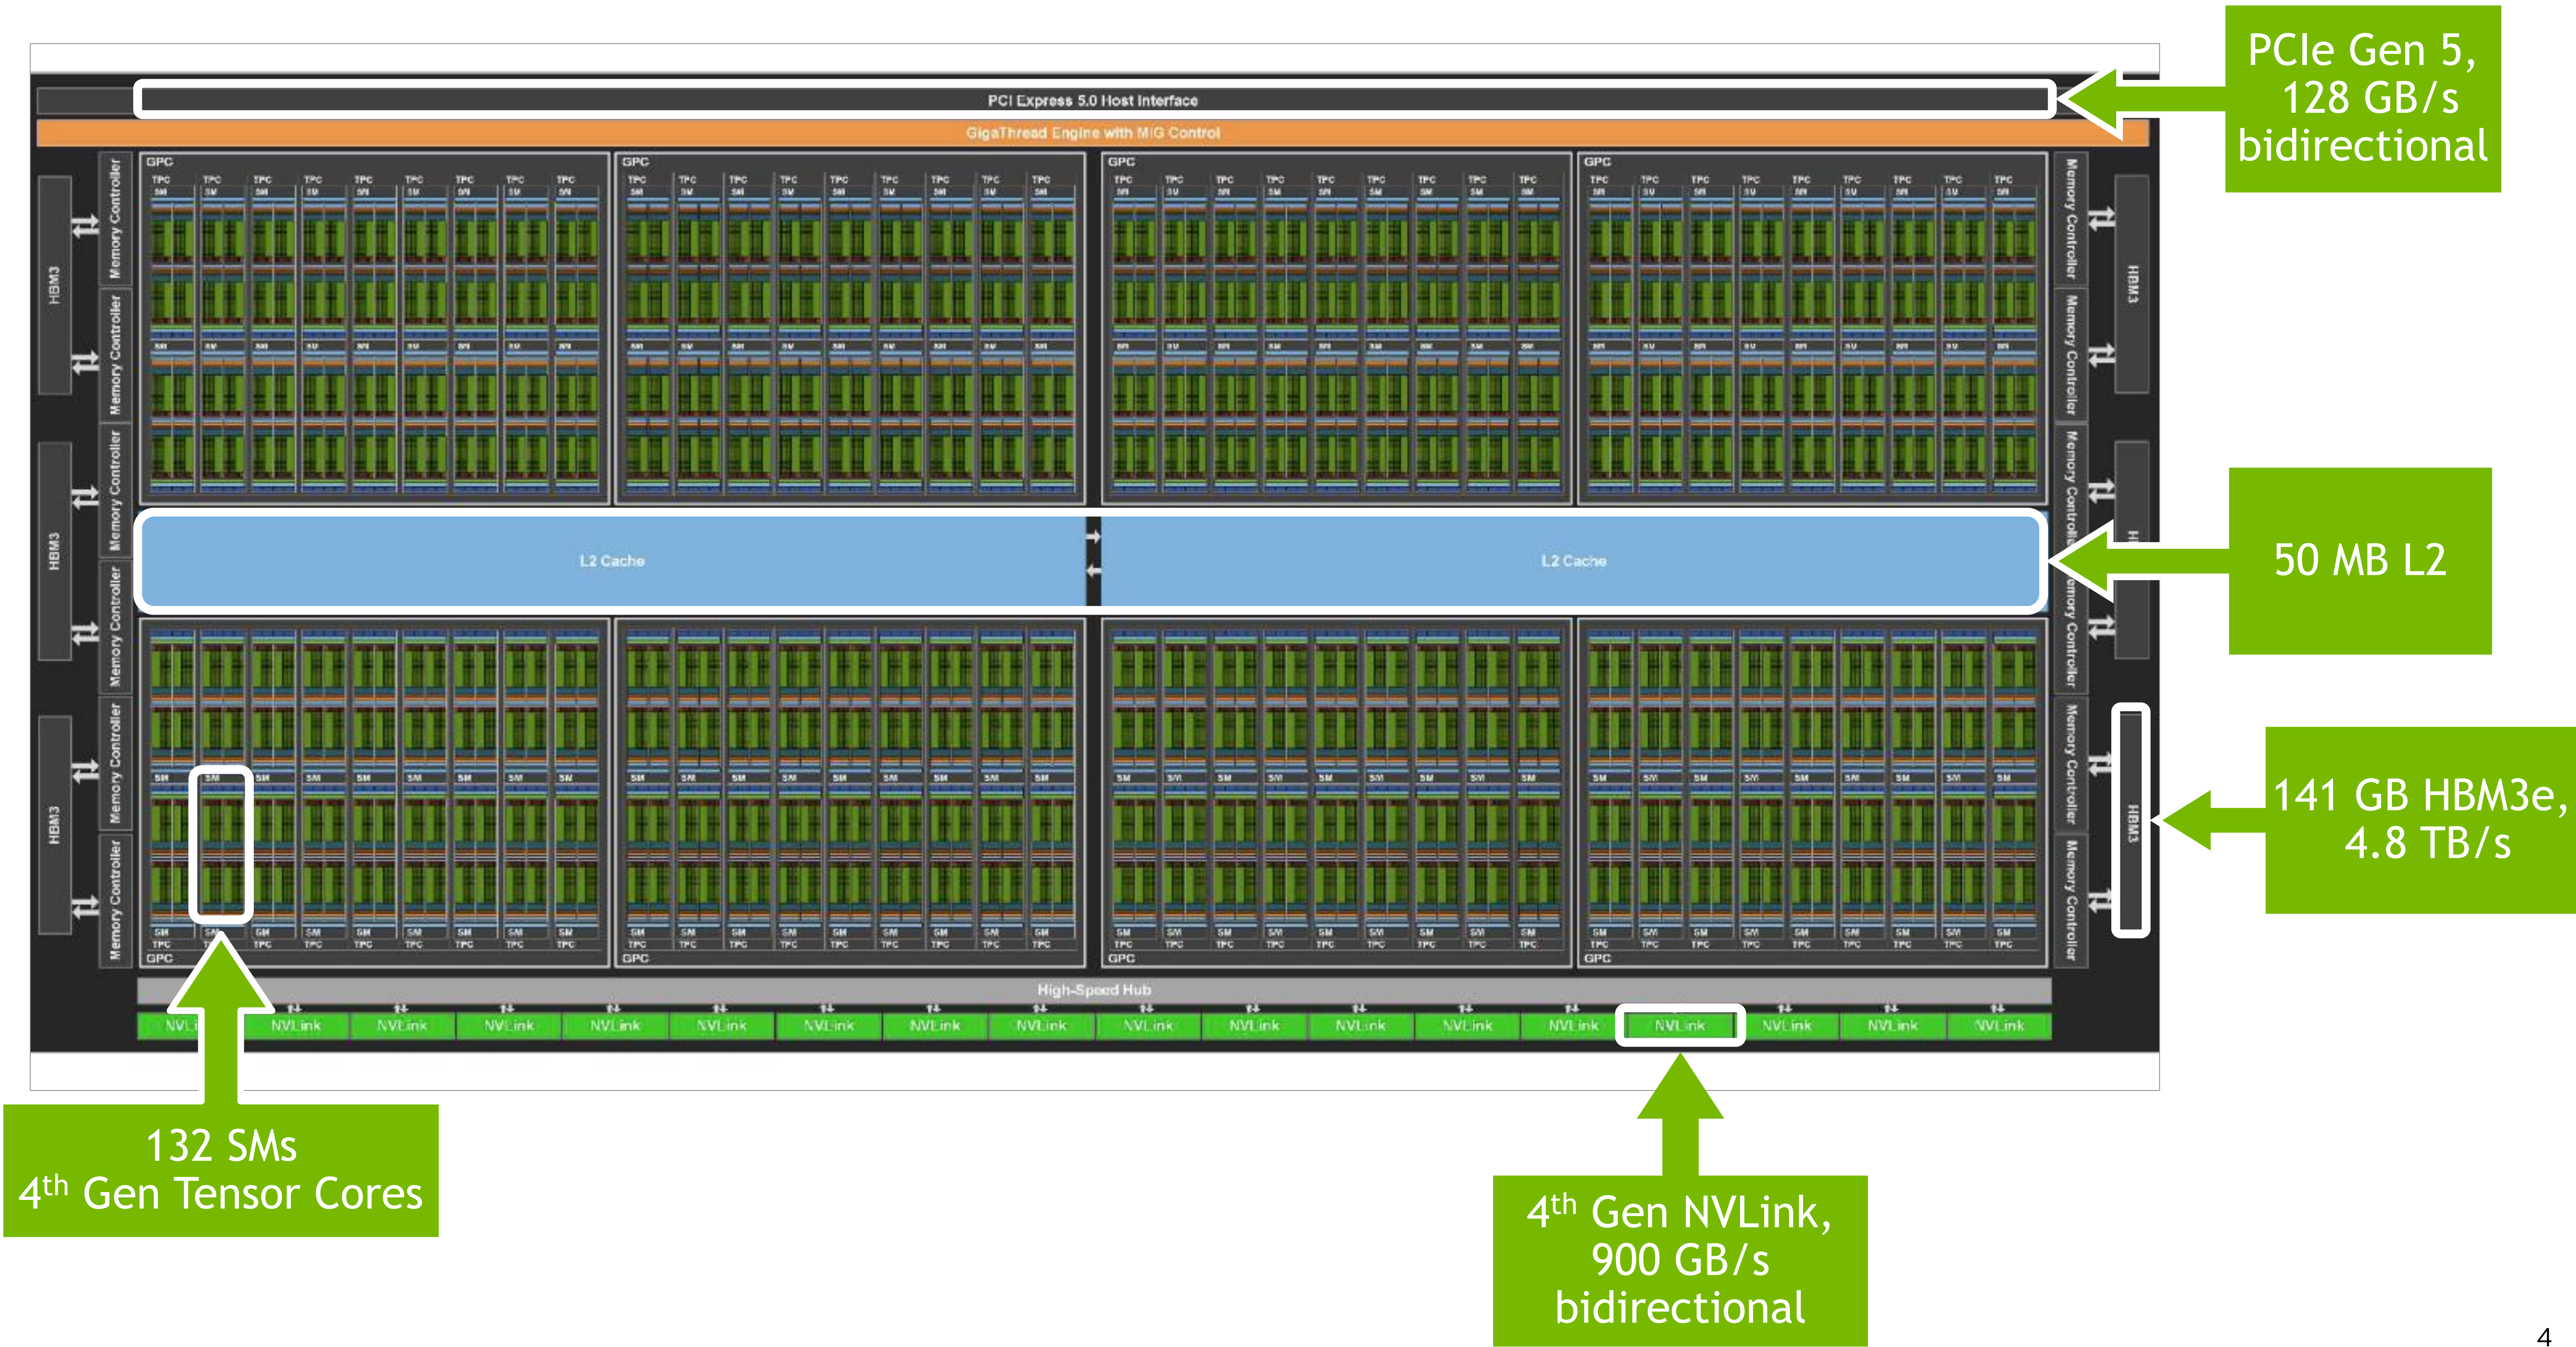
- Fundamental Performance Optimizations

- Summary

NVIDIA.

GPU Architecture and CUDA Programming Basics

GPU Overview

NVIDIA H200 SXM

PCIe Gen 5, 128 GB/s bidirectional

50 MB L2

141 GB HBM3e, 4.8 TB/s

132 SMs 4th Gen Tensor Cores

4th Gen NVLink, 900 GB/s bidirectional

Streaming Multiprocessor (SM)
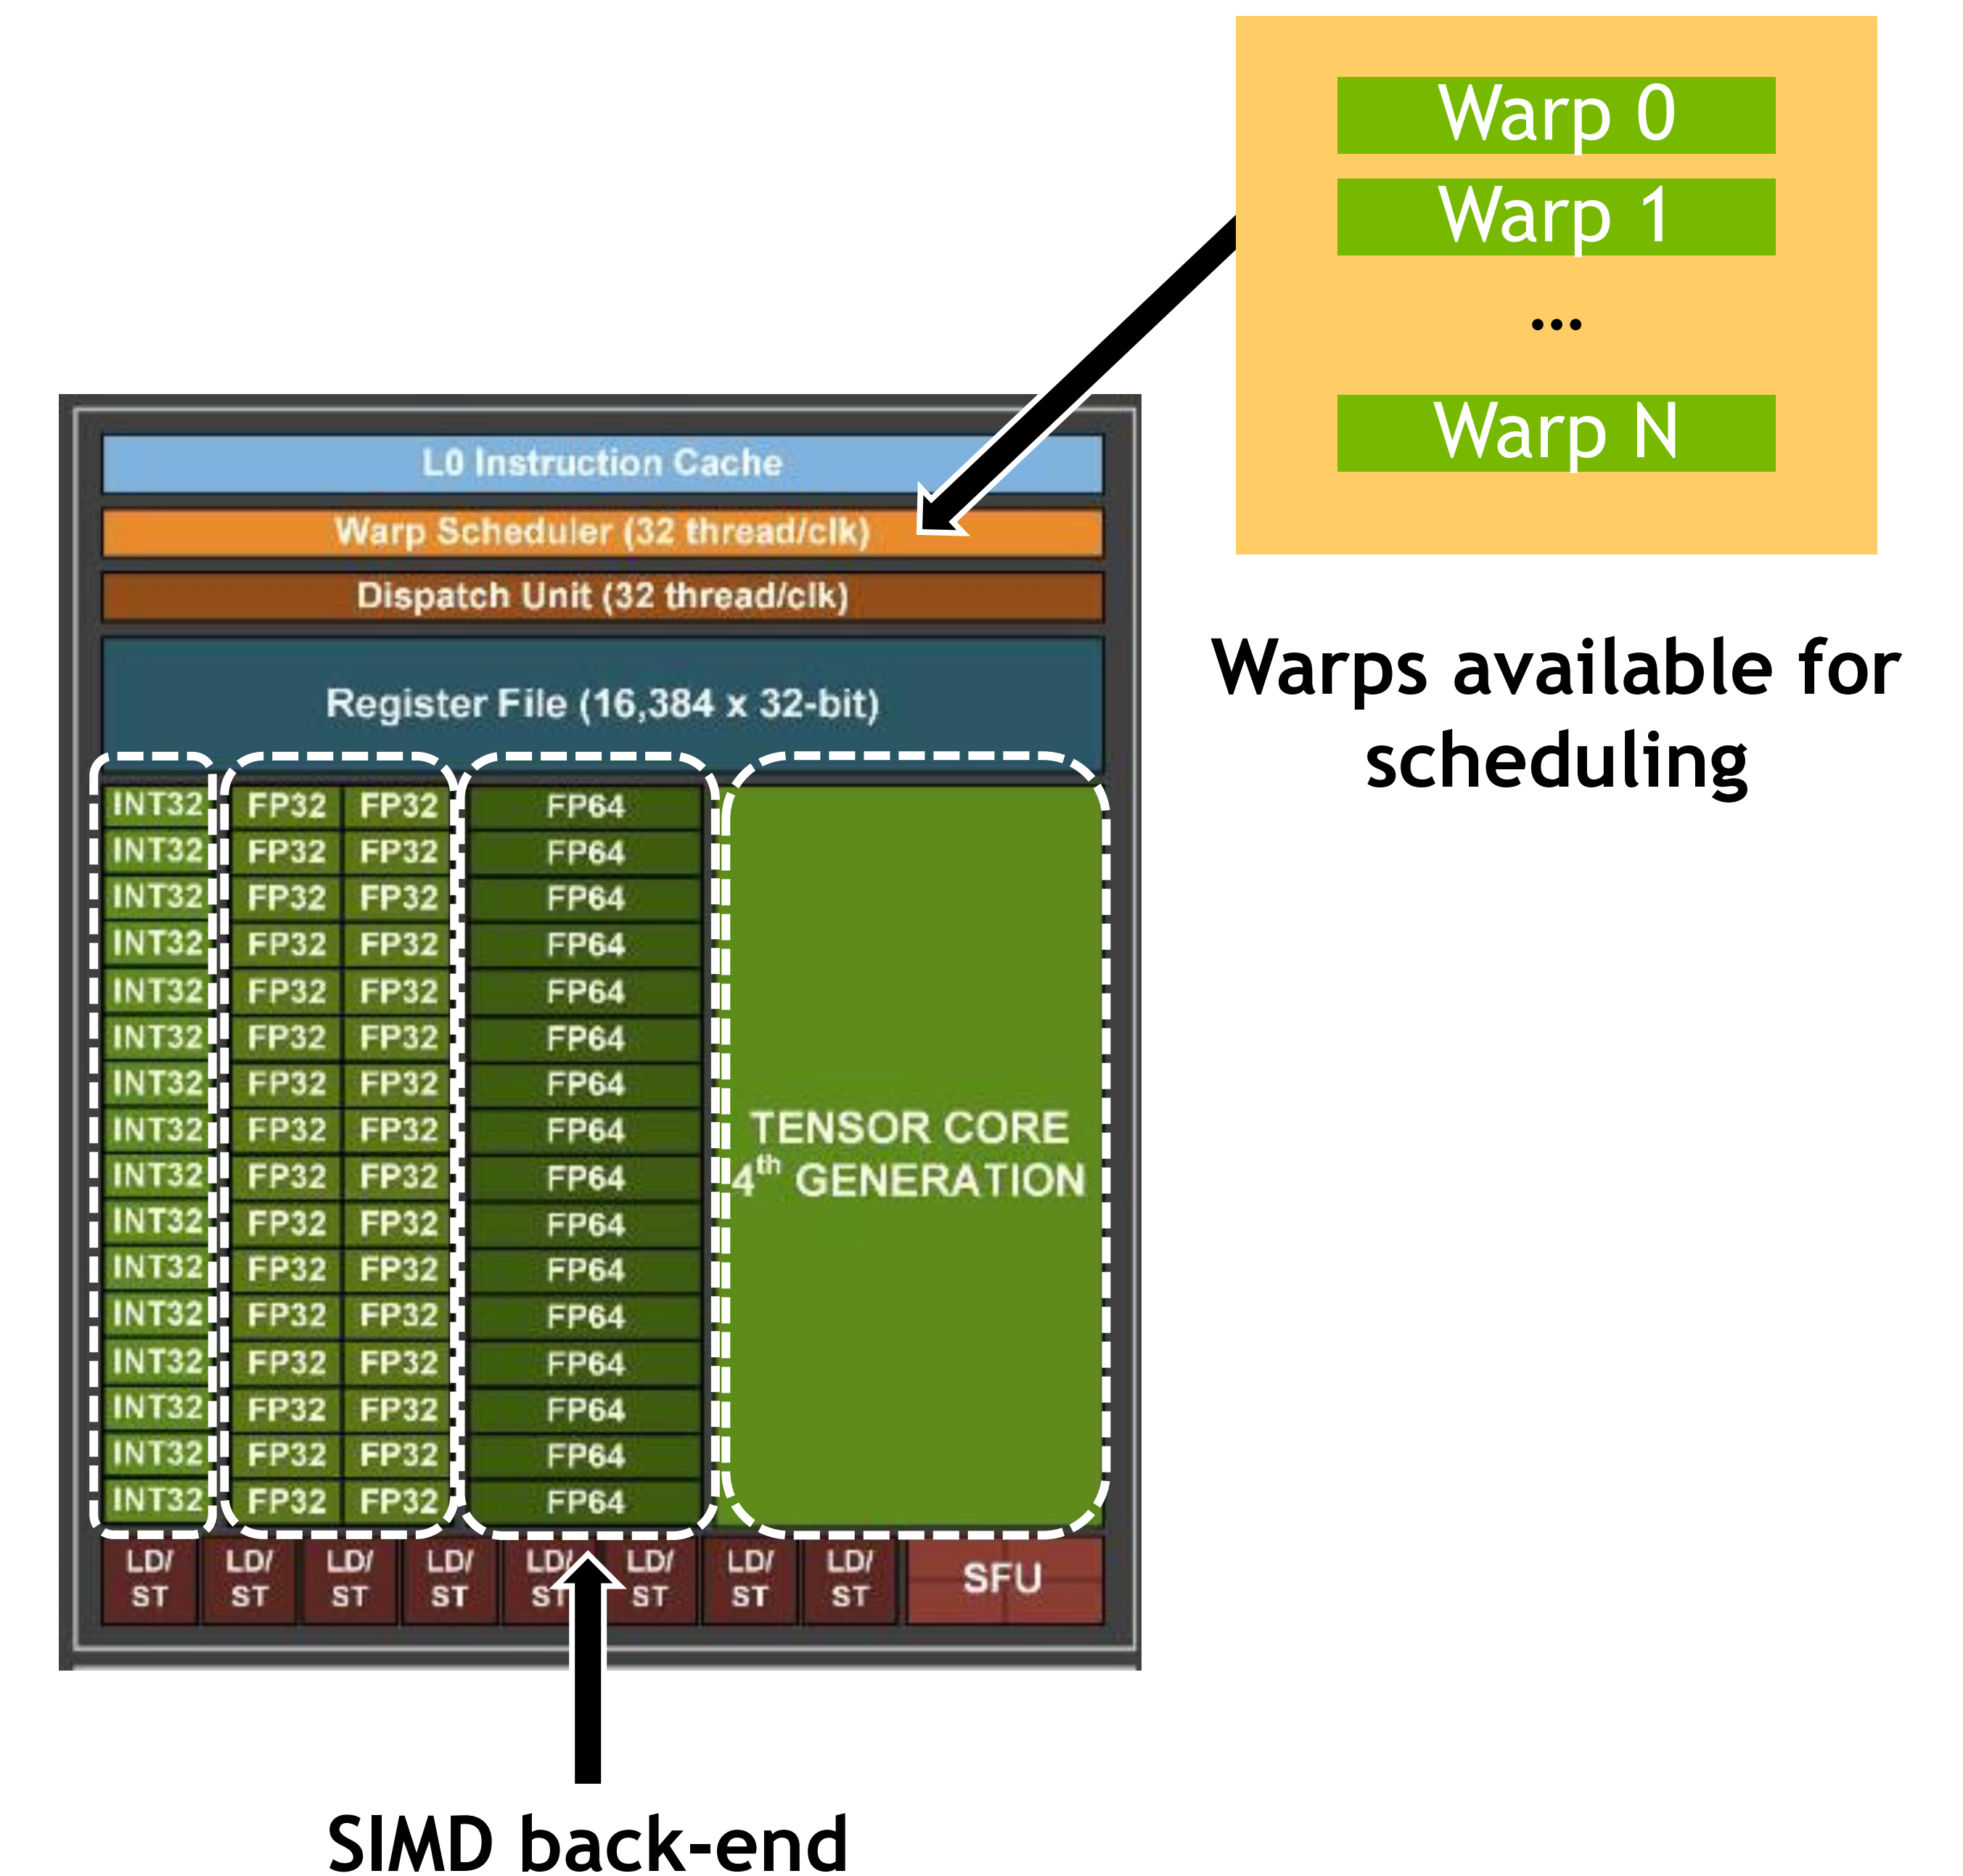
Hopper architecture

- 128 FP32 cores
- 64 FP64 cores
- 64 INT32 cores
- 4 mixed-precision Tensor Cores
- 16 special function units (transcendentals)
- 4 warp schedulers

- 32 LD/ST units
- 64K 32-bit registers
- 256 KiB unified L1 data cache and shared memory
- Tensor Memory Accelerator (TMA)

SIMT Architecture
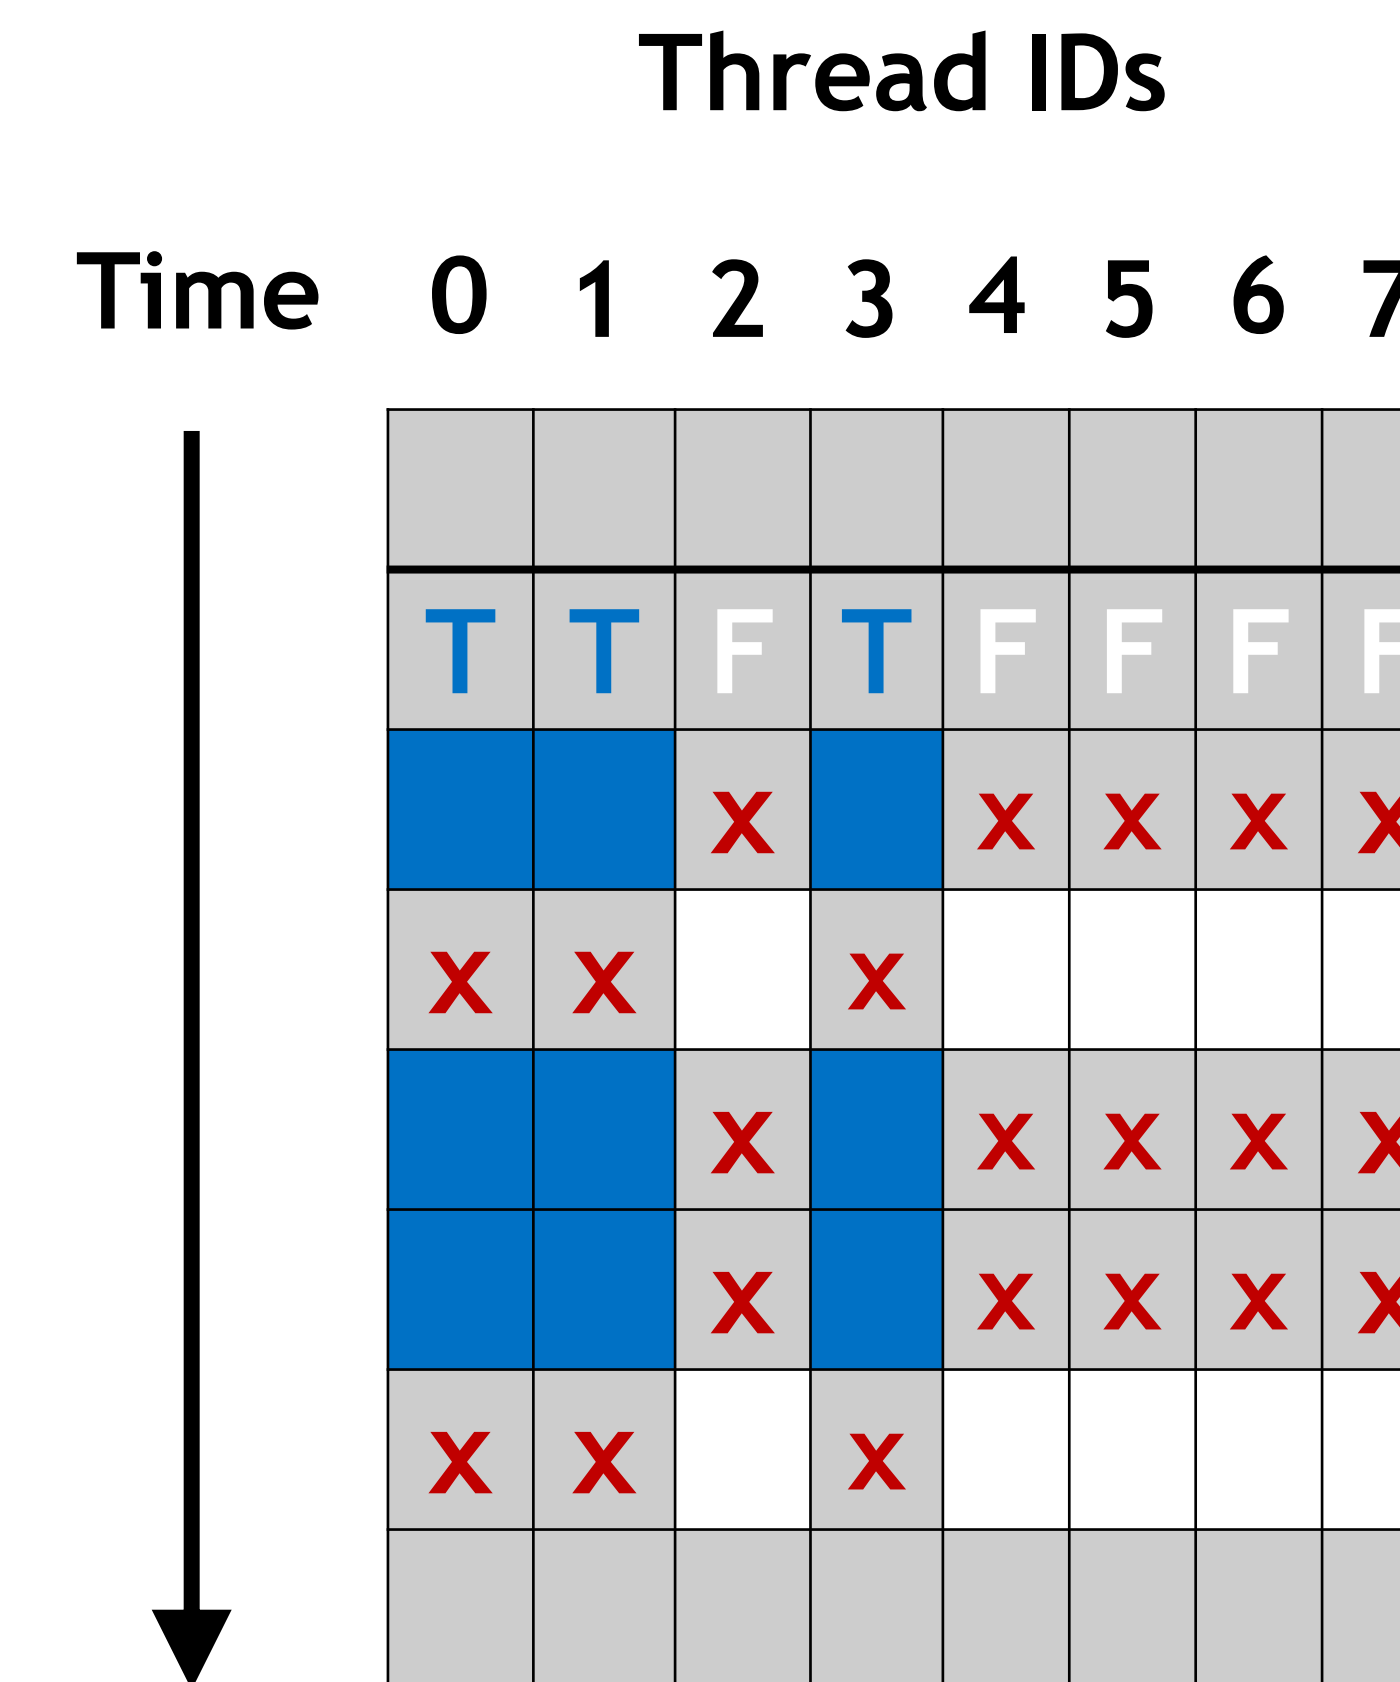## Single-Instruction, Multiple-Thread

- Akin to a **single-instruction multiple-data (SIMD) array processor** per Flynn's taxonomy combined with **fine-grained multithreading**.

- SIMT architectures expose a large set of hardware threads, which is partitioned into groups called **warps.**
  - Interleave warp execution to hide latencies.
  - Execution context for each warp is kept on-chip for fast interleaving.

- When scheduled, each thread of a warp executes on a given lane of a SIMD functional unit.

- Each SM sub-partition can be thought of as a SIMT engine that creates, manages, schedules, and executes warps of 32 parallel threads.

Warps available for scheduling

SIMD back-end

SIMT Architecture
Warp divergence

- If threads in a warp **diverge** via a conditional branch, the warp **executes every branch** path.

- **Full efficiency** is realized when all **32 threads** of a warp agree on their execution path.
 - Aka they are **converged**.

Thread IDs

Time 0 1 2 3 4 5 6 7

	0	1	2	3	4	5	6	7
T	T	F	T	F	F	F	F
			x		x	x	x	x
	x	x		x				
			x		x	x	x	x
			x		x	x	x	x
	x	x		x				

```
if (true) {
    instruction 1
    instruction 2
    instruction 3
} else {
    instruction 4
    instruction 5
}
```

SIMT Architecture
Independent Thread Scheduling

- Individual **threads** in a warp have their own program counter and call stack and are therefore **free to execute independently**.

```
if (thread_id < 4) {
    A;
    B;
} else {
    X;
    Y;
}
Z;
```


Time →

Do **not** assume threads in a warp are automatically re-converged after a conditional or at any point!

The compiler **might** sync to enforce re-convergence for better performance.

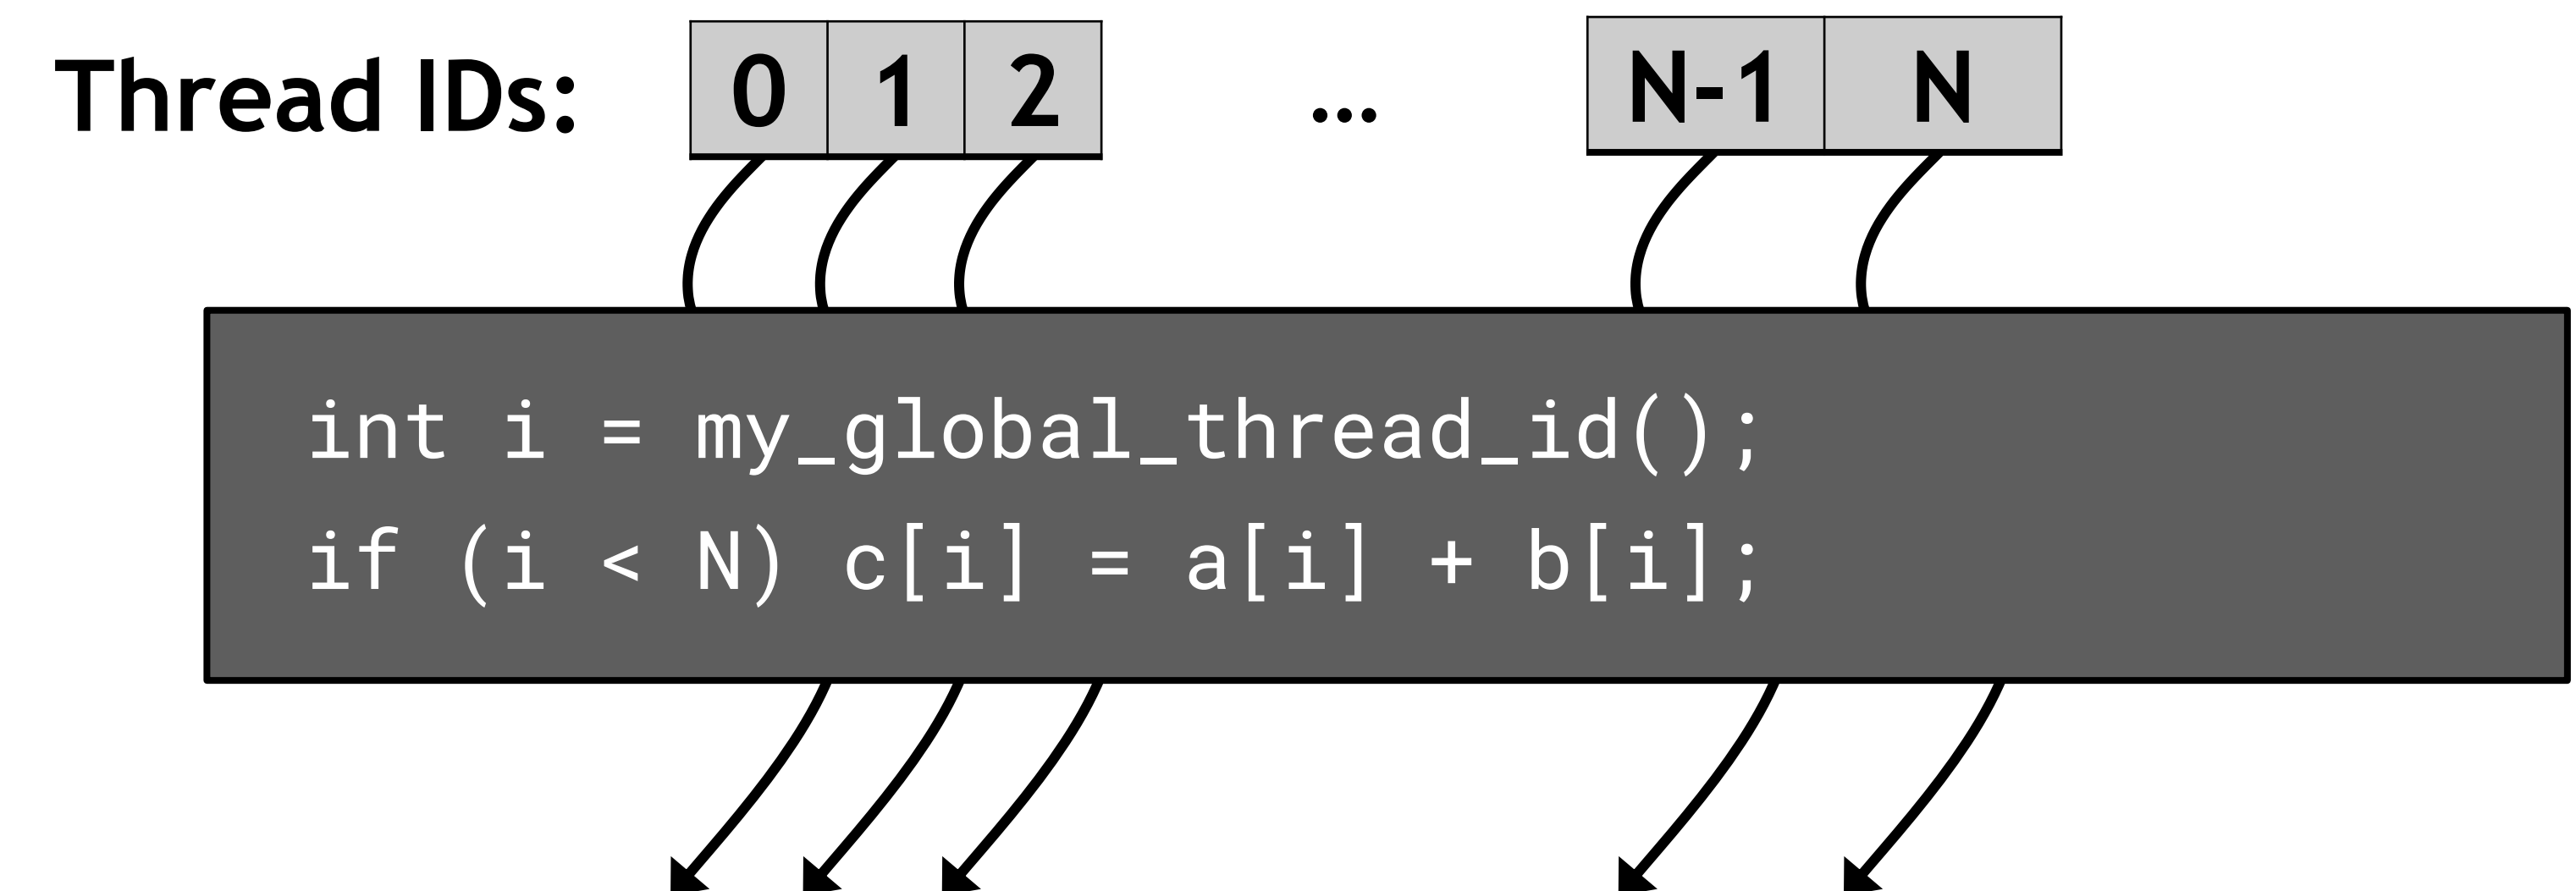NVIDIA.

CUDA Programming Model
Single-Program Multiple-Data

- SIMT instructions specify the execution of a **single** thread.

- A SIMT kernel is launched on many threads that execute in parallel.

- Threads use their thread index to work on disjoint data or to enable different execution paths.

- Three key software abstractions enable efficient programming through the CUDA programming model:
  - a hierarchy of **thread groups**,
  - **memory spaces**, and
  - **synchronization**.
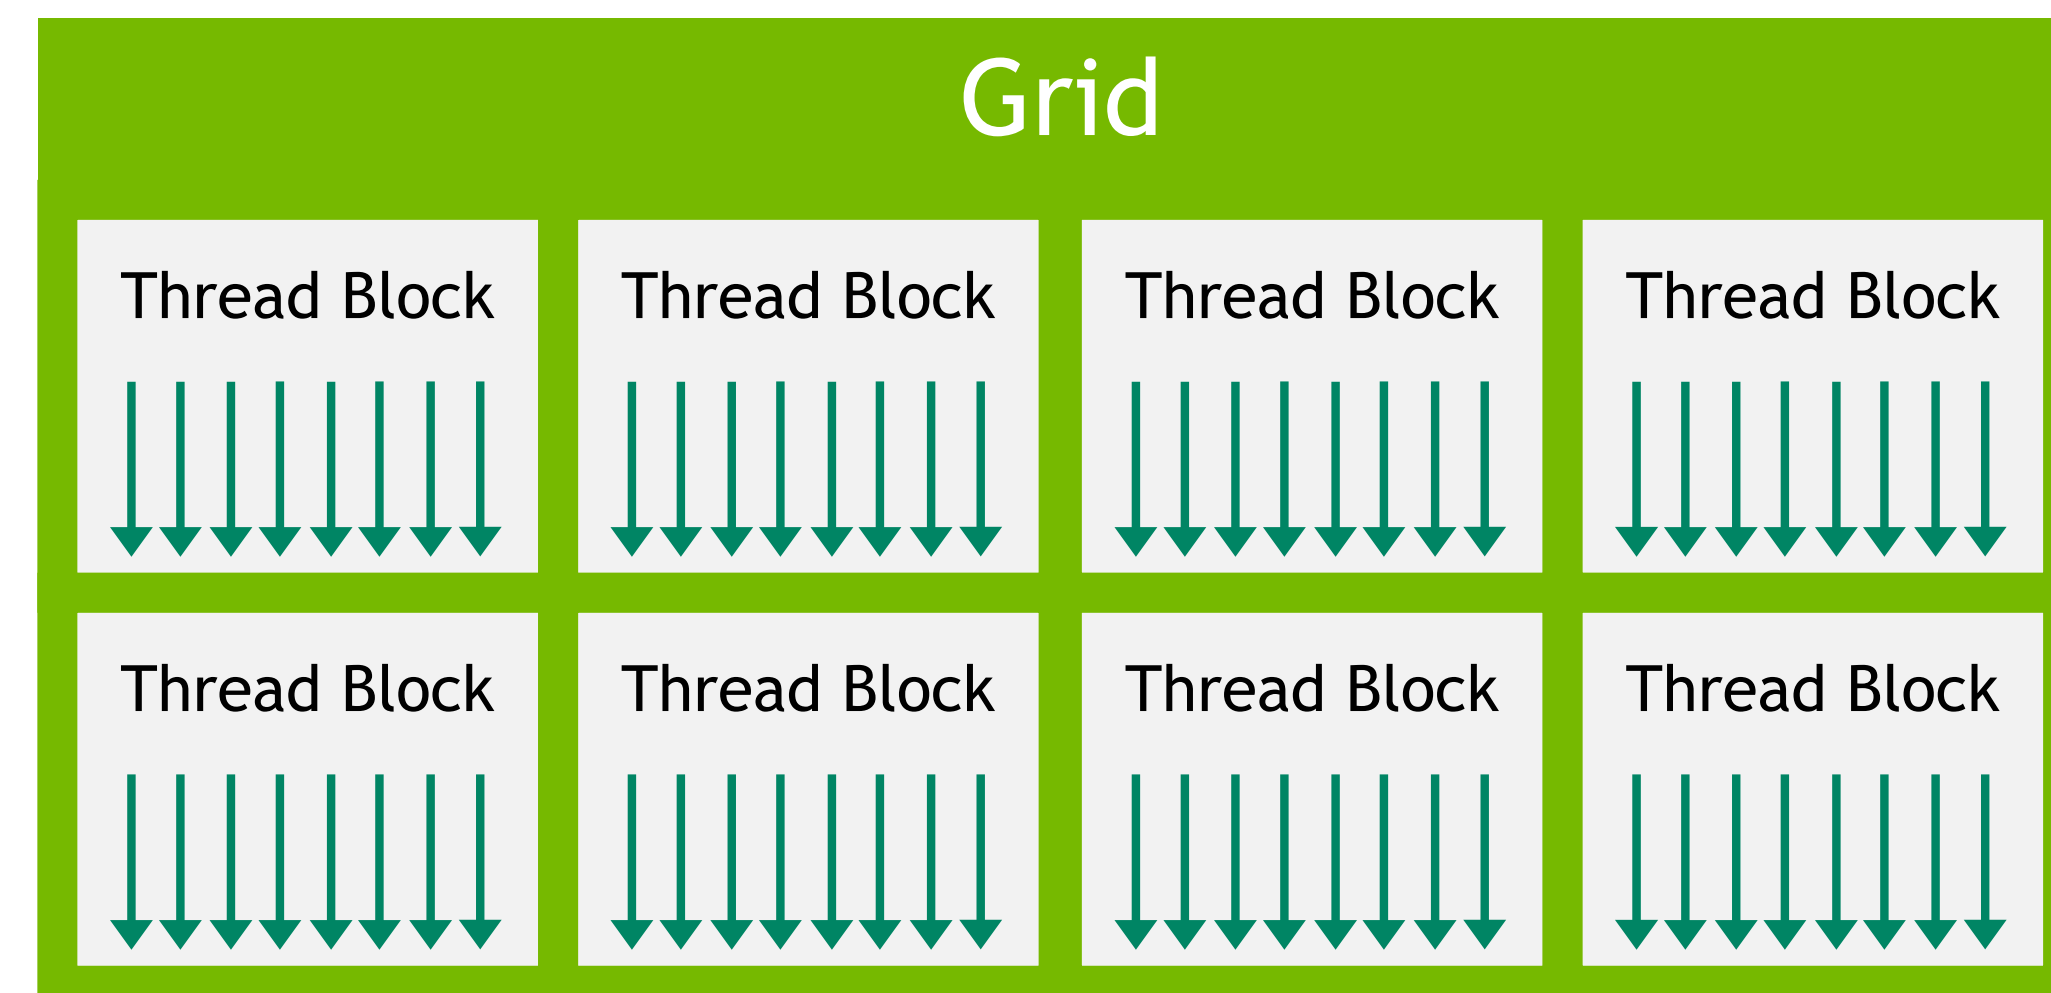
Single-threaded CPU vector addition

```
for (int i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```

GPU vector addition

Thread IDs: | 0 | 1 | 2 | … | N-1 | N |

```
int i = my_global_thread_id();
if (i < N) c[i] = a[i] + b[i];
```

9

Thread Hierarchy

CUDA/Software

Grid

| Thread Block	Thread Block	Thread Block	Thread Block
Thread Block	Thread Block	Thread Block	Thread Block

Thread Block

Thread

Hardware

Device

SM

Scalar CUDA core

- A CUDA kernel is launched on a grid of thread blocks, which are completely independent.

- Thread blocks are executed on SMs.
  - Several concurrent thread blocks can reside on an SM.
  - Thread blocks do not migrate.
  - Each block can be scheduled on any of the available SMs, in any order, concurrently or in series.

- Individual threads execute on scalar CUDA cores.

# Thread Block Clusters

- For Hopper GPUs, CUDA introduced an optional level in the thread hierarchy called **Thread Block Clusters**.

- Thread blocks in a cluster are guaranteed to be concurrently scheduled and enable efficient cooperation and data sharing for threads across multiple SMs.

- For more information on this topic visit GTC session **[S62192]: "Advanced Performance Optimization in CUDA"**.
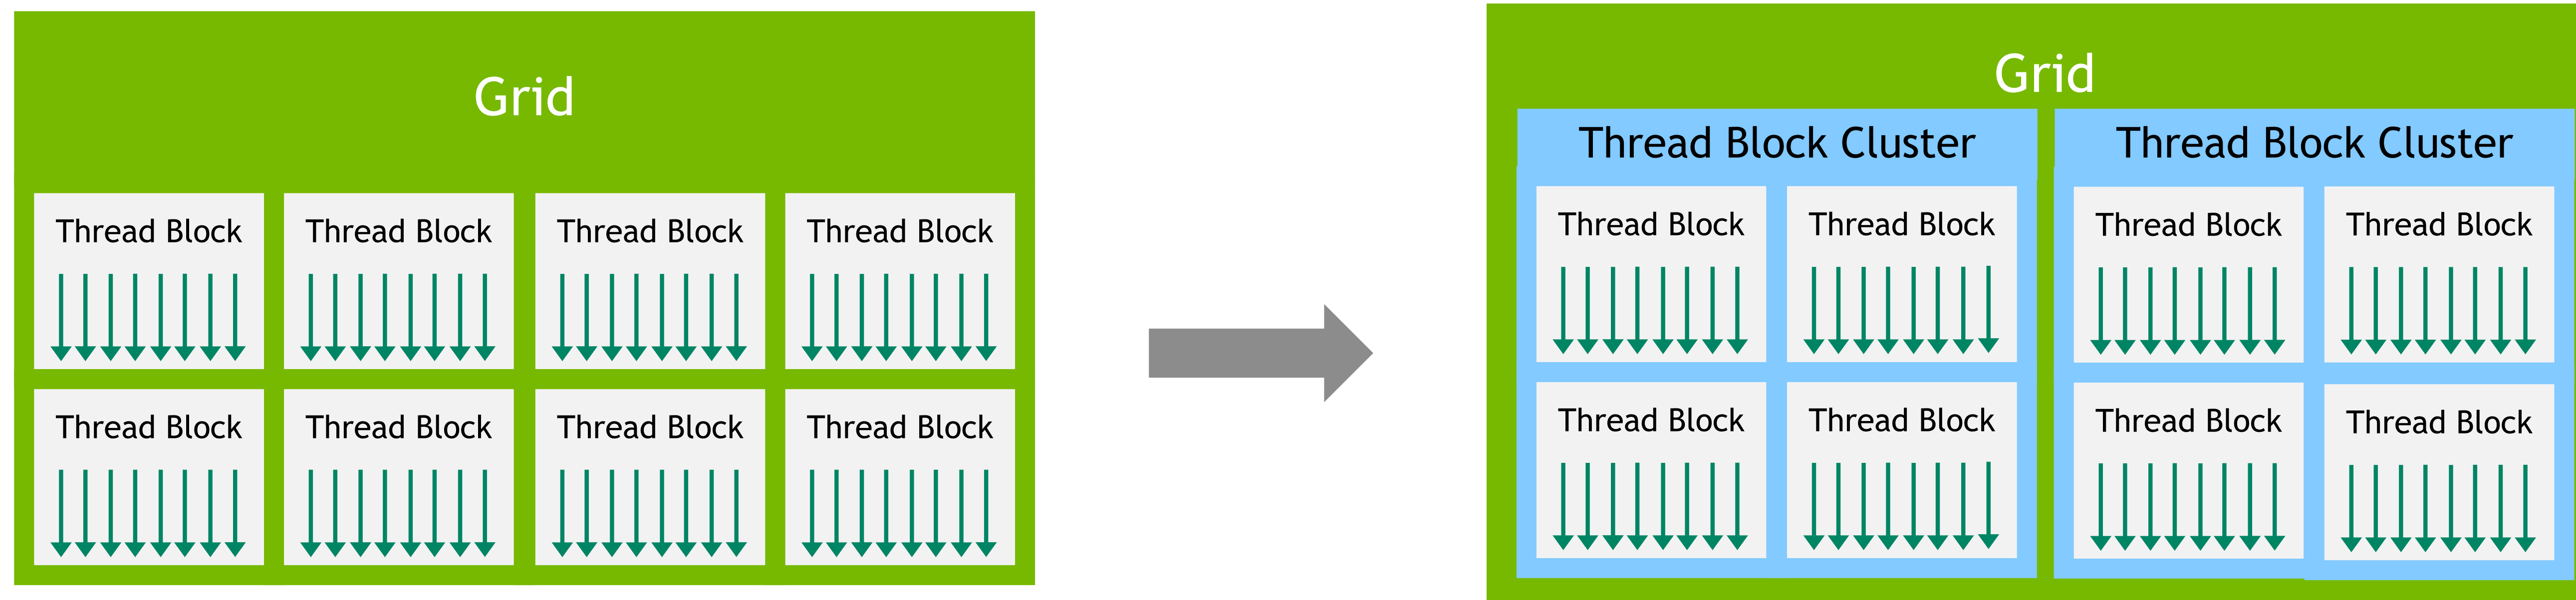
Thread Hierarchy
What about warps?

- At runtime, a block of threads is divided into warps for SIMT execution.
  - The way a block is partitioned into warps is always the same.
    - Each warp contains threads of consecutive, increasing thread IDs with the first warp containing thread 0.

- The total number of warps in a block is defined as:
  - $ceil\left(\frac{threads\ per\ block}{warp\ size}, 1\right)$
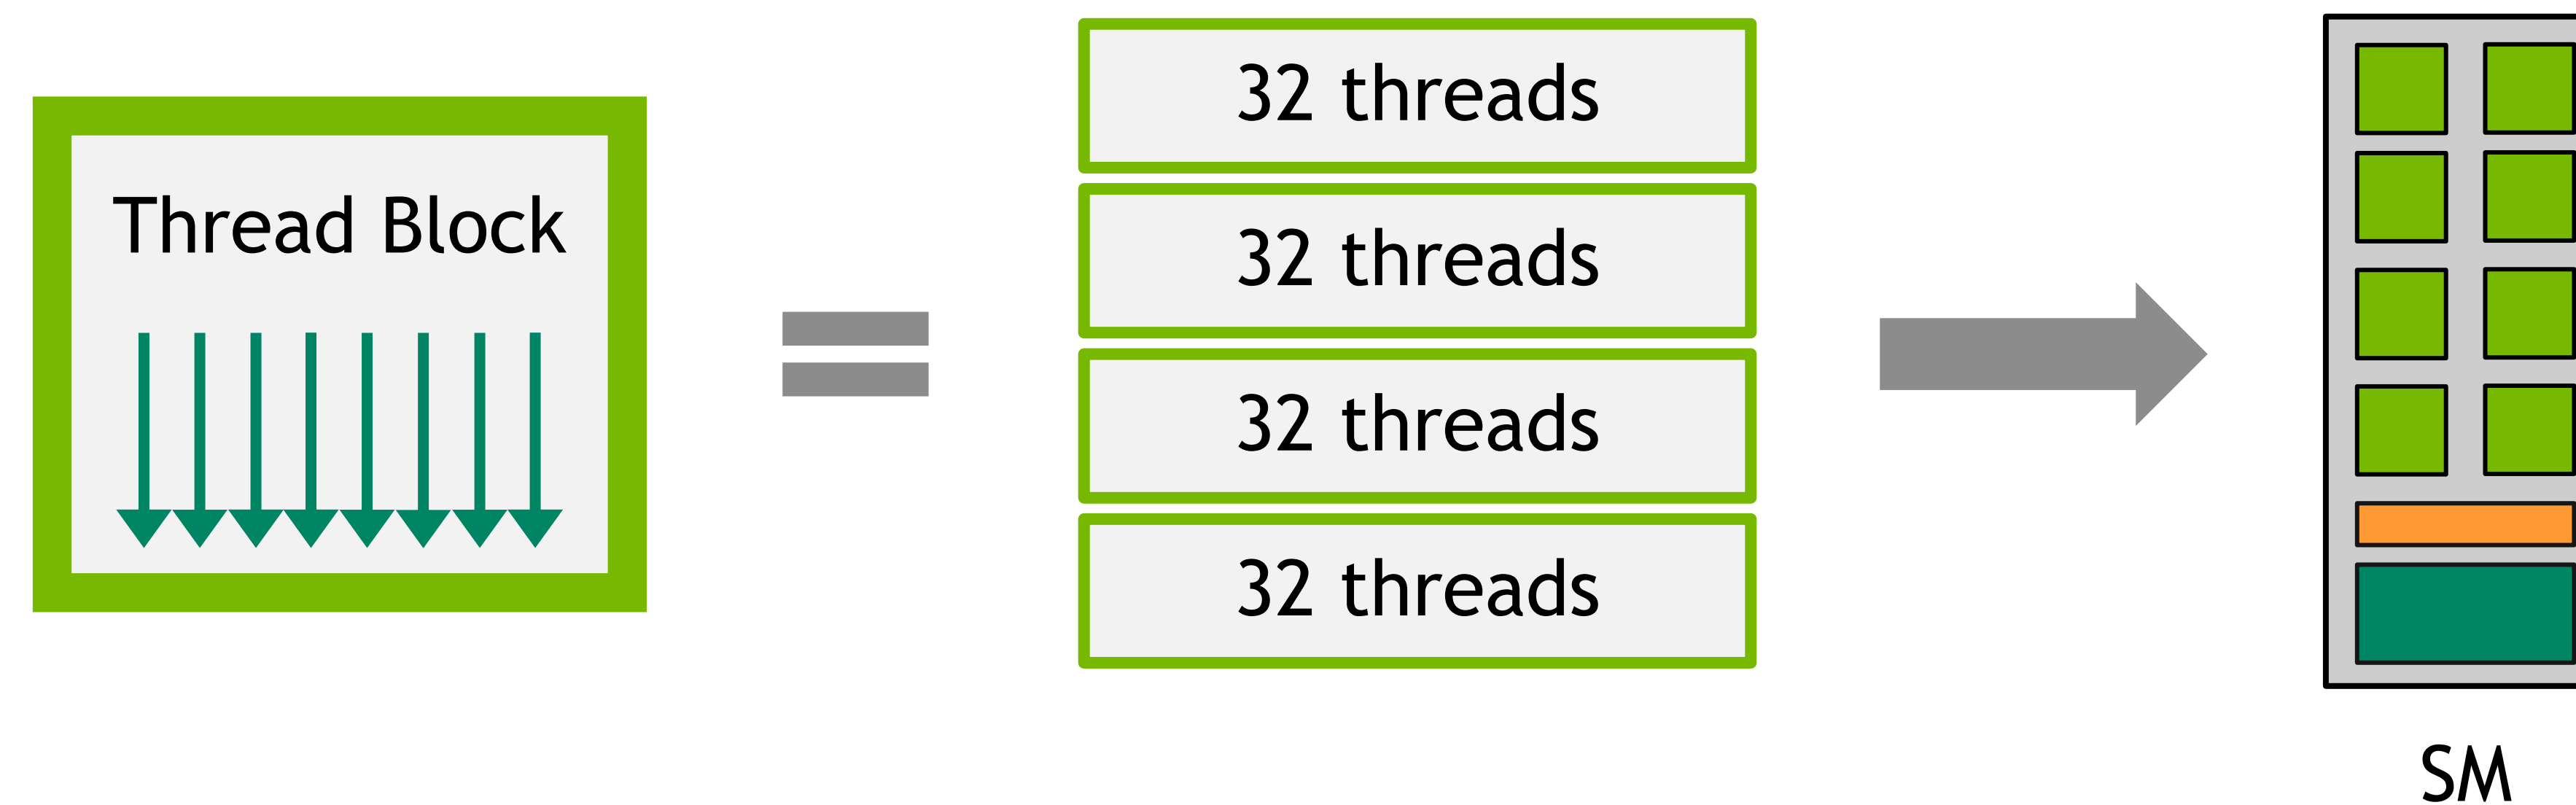
| Thread Block | | 32 threads | | SM |

Thread Hierarchy
Thread block sizing

- Let's say we want to add two vectors of size N = 1000.
  - **Scenario #1:** 1-D grid of 10 1-D blocks of size 100.
  - **Scenario #2:** 1-D grid of 8 1-D blocks of size 128.

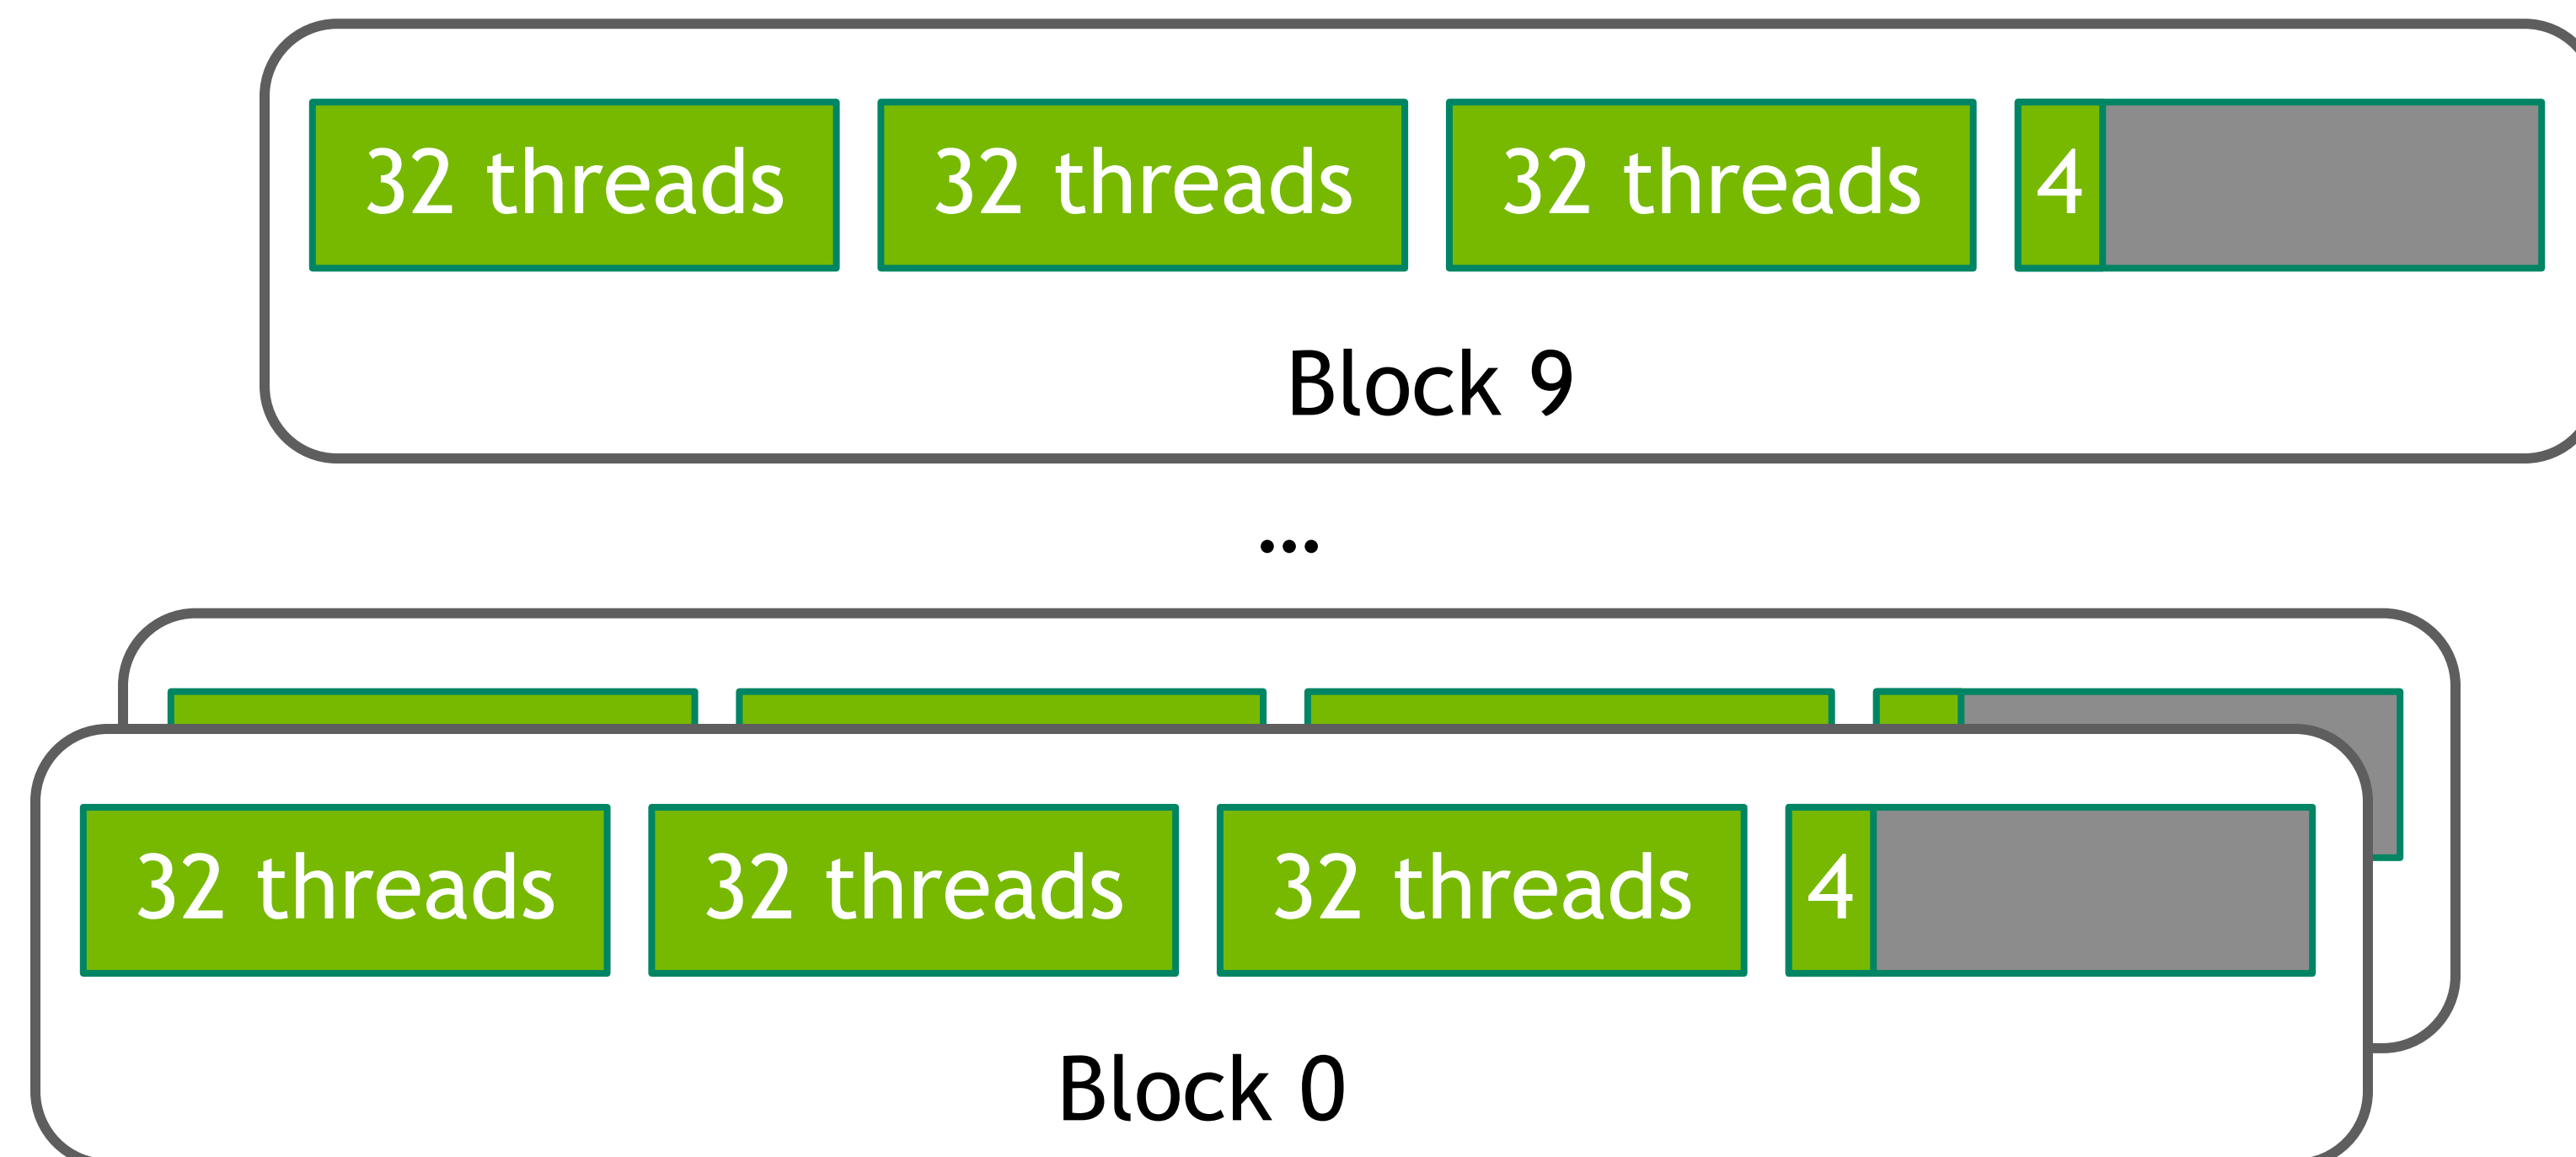- Which option is better in terms of thread resource utilization?

Thread Hierarchy
Thread block sizing

- Let's say we want to add two vectors of size N = 1000.
  - **Scenario #1:** 1-D grid of 10 1-D blocks of size 100.
  - **Scenario #2:** 1-D grid of 8 1-D blocks of size 128.

- Which option is better in terms of thread resource utilization?

| 32 threads | 32 threads | 32 threads | 4 |
Block 9

...

| 32 threads | 32 threads | 32 threads | 4 |
Block 0

Scenario #1:
3 full warps and 1 warp with 4 active threads per block
Average thread utilization = 78.125%

| 32 threads | 32 threads | 32 threads | 8 |
Block 7

...

| 32 threads | 32 threads | 32 threads | 32 threads |
Block 0
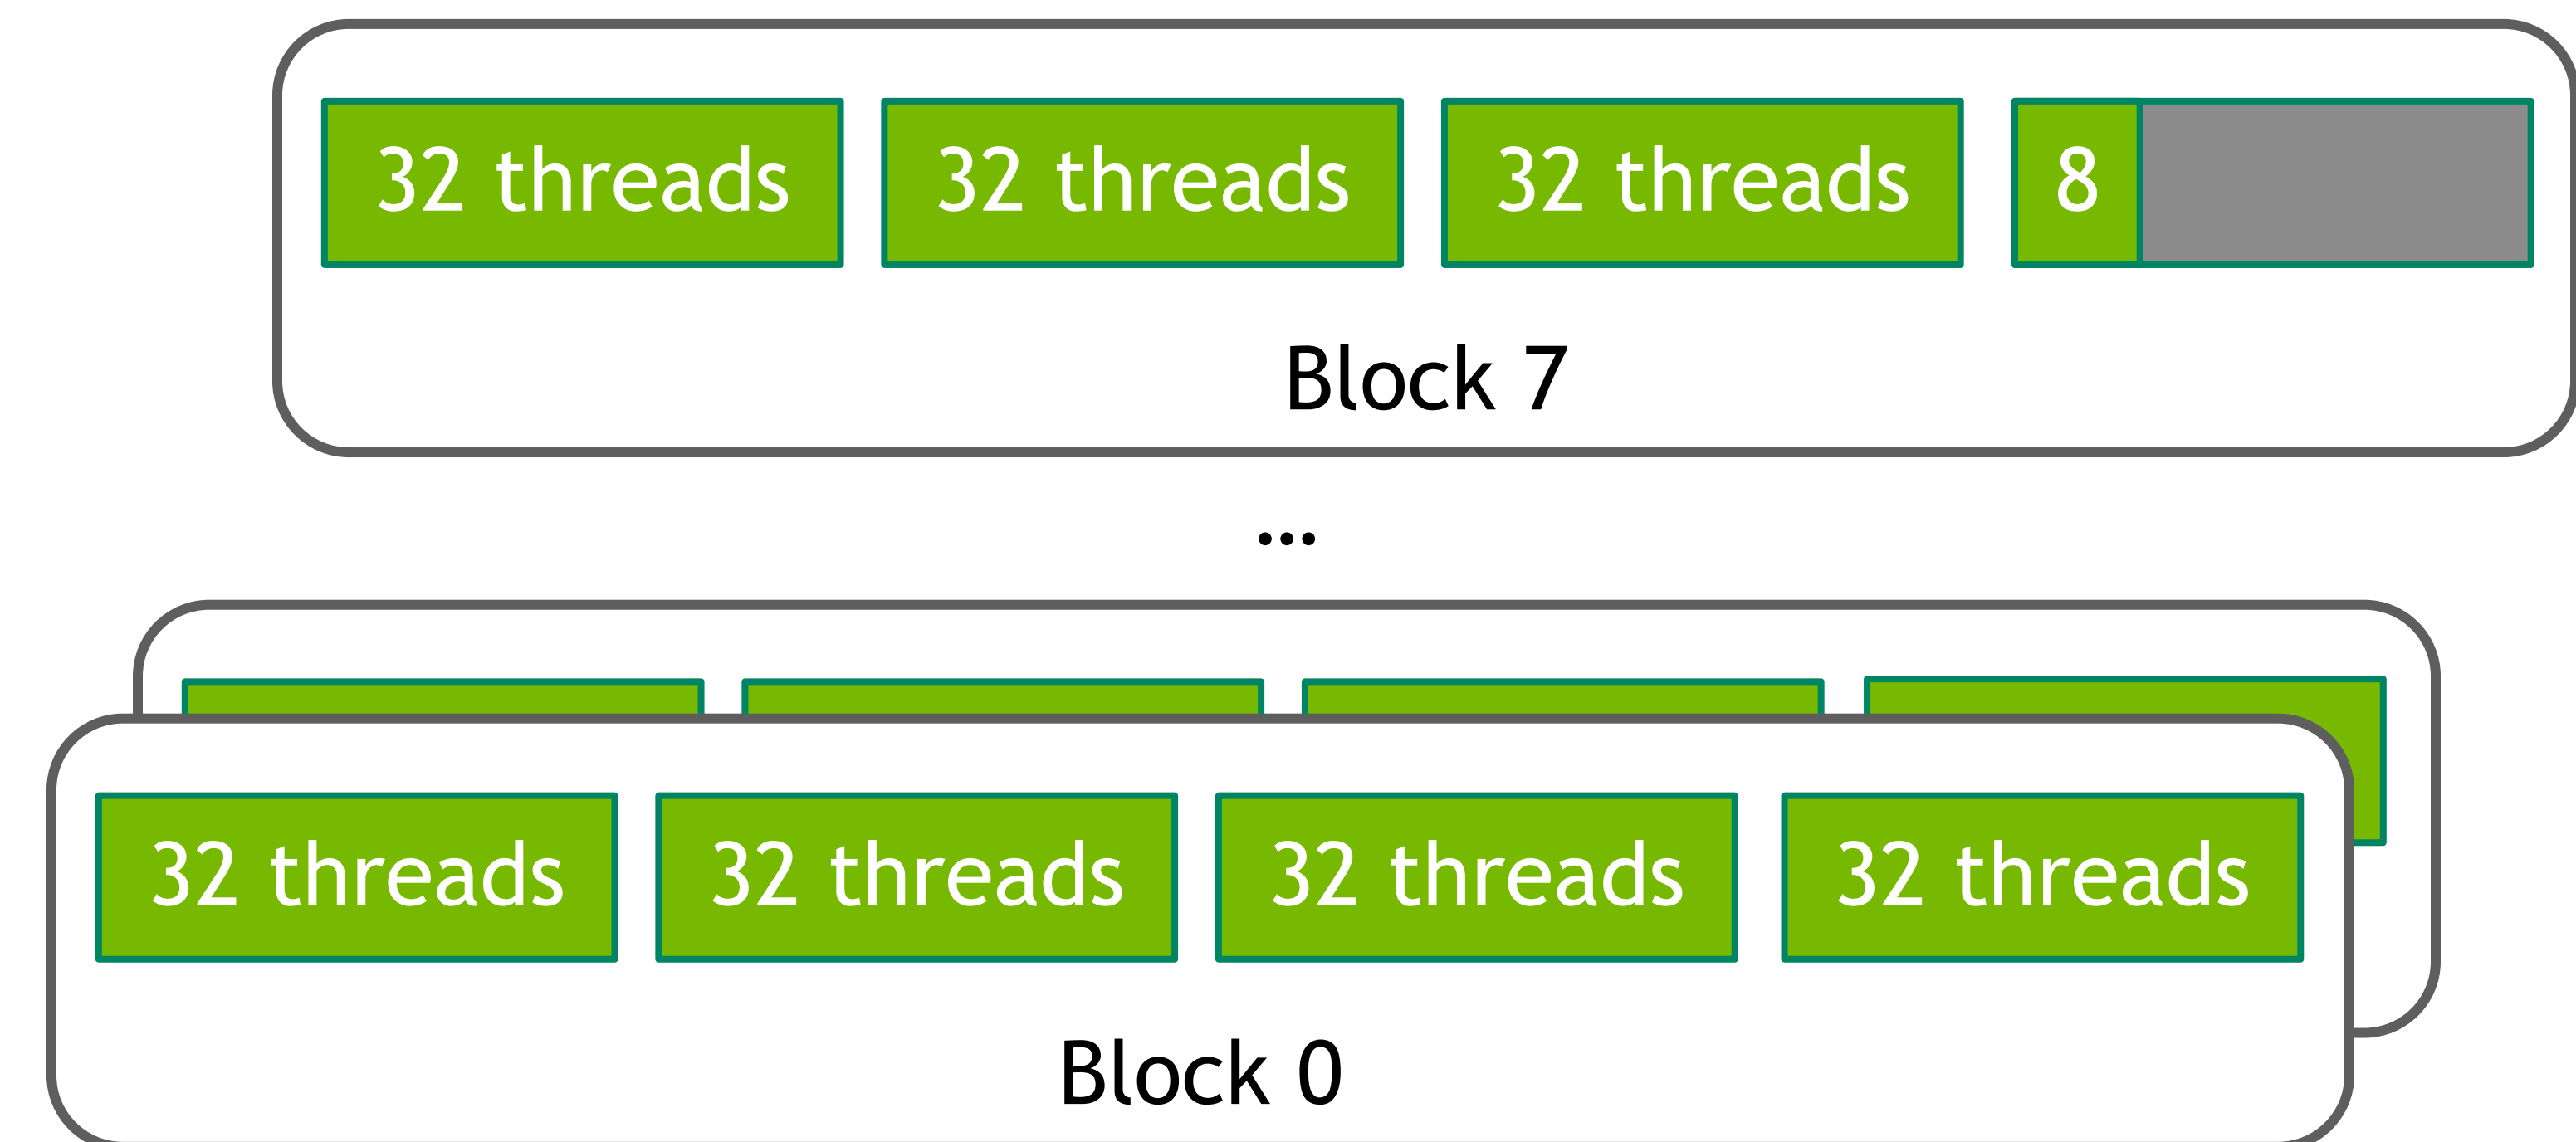
Scenario #2:
4 full warps per block, except last block
Average thread utilization = 97.656%

NVIDIA.

Memory Hierarchy

Hardware

```
SM
  Registers
  Shared/L1
L2
DRAM
```

CUDA/Software

Thread

```
Registers    Local Memory
```

```
Thread Block
  Shared
  Memory
```

```
Grid
  Thread Block  Thread Block  Thread Block  Thread Block
  Thread Block  Thread Block  Thread Block  Thread Block
  Global Memory
```
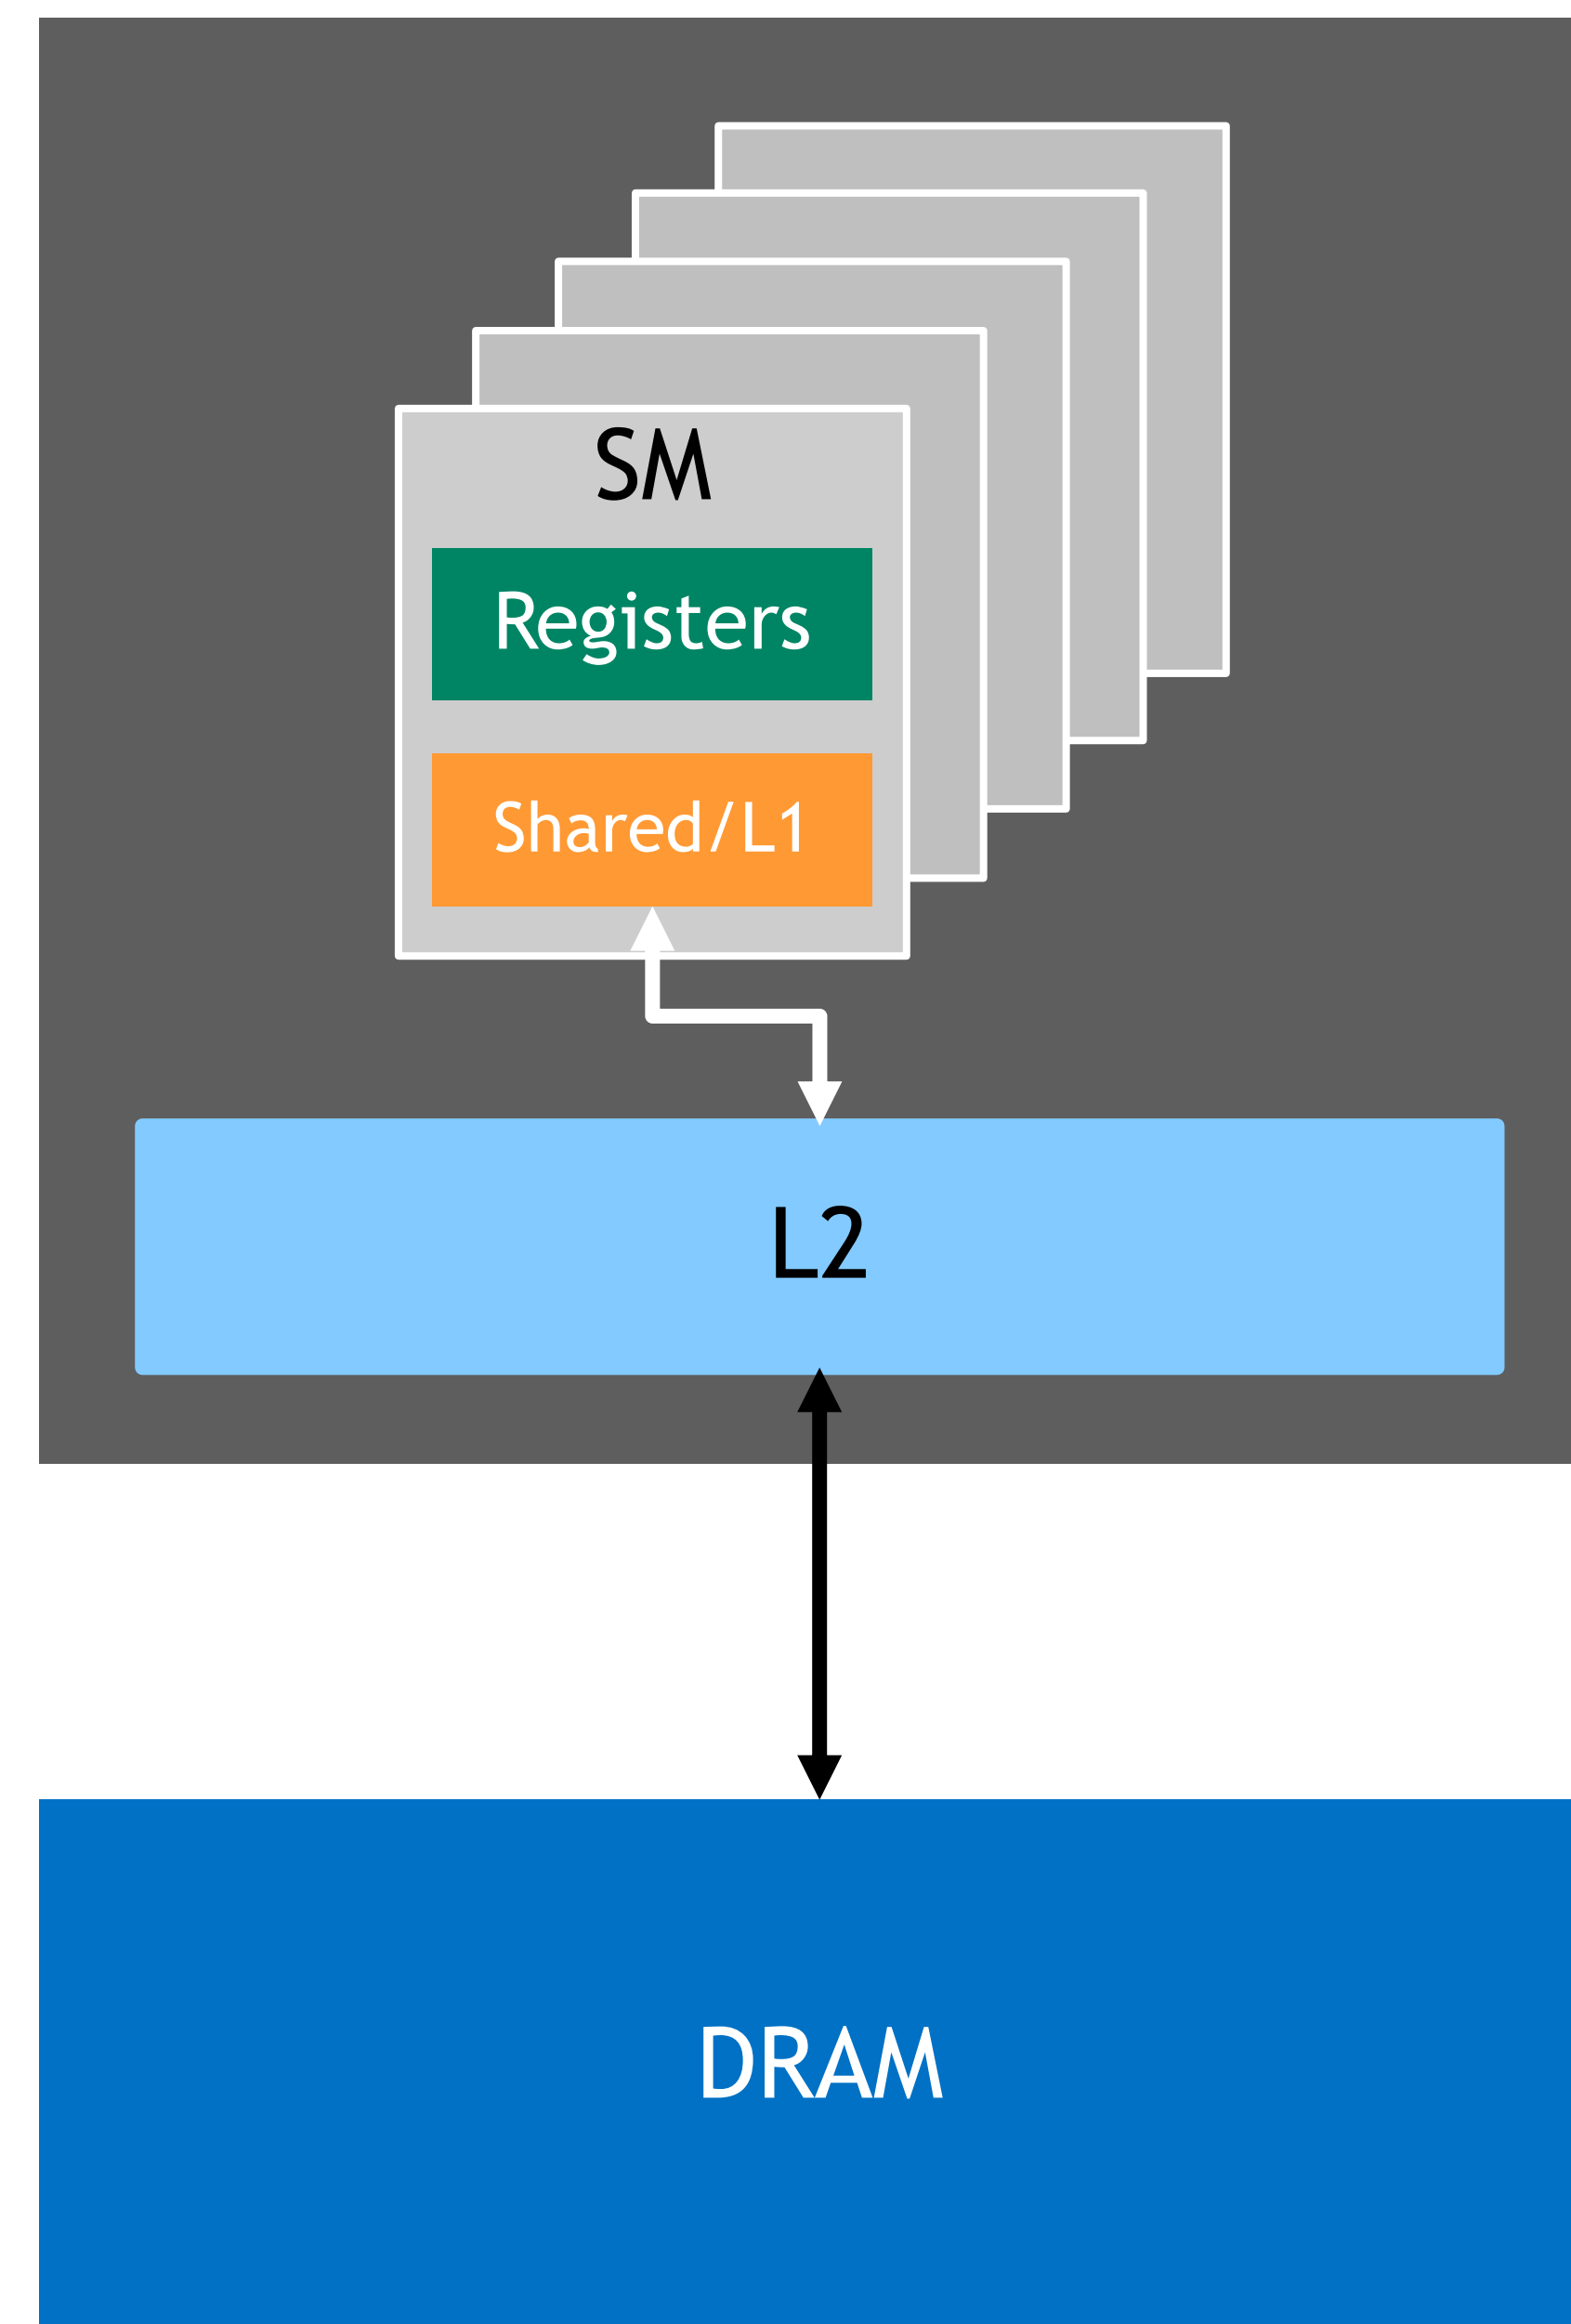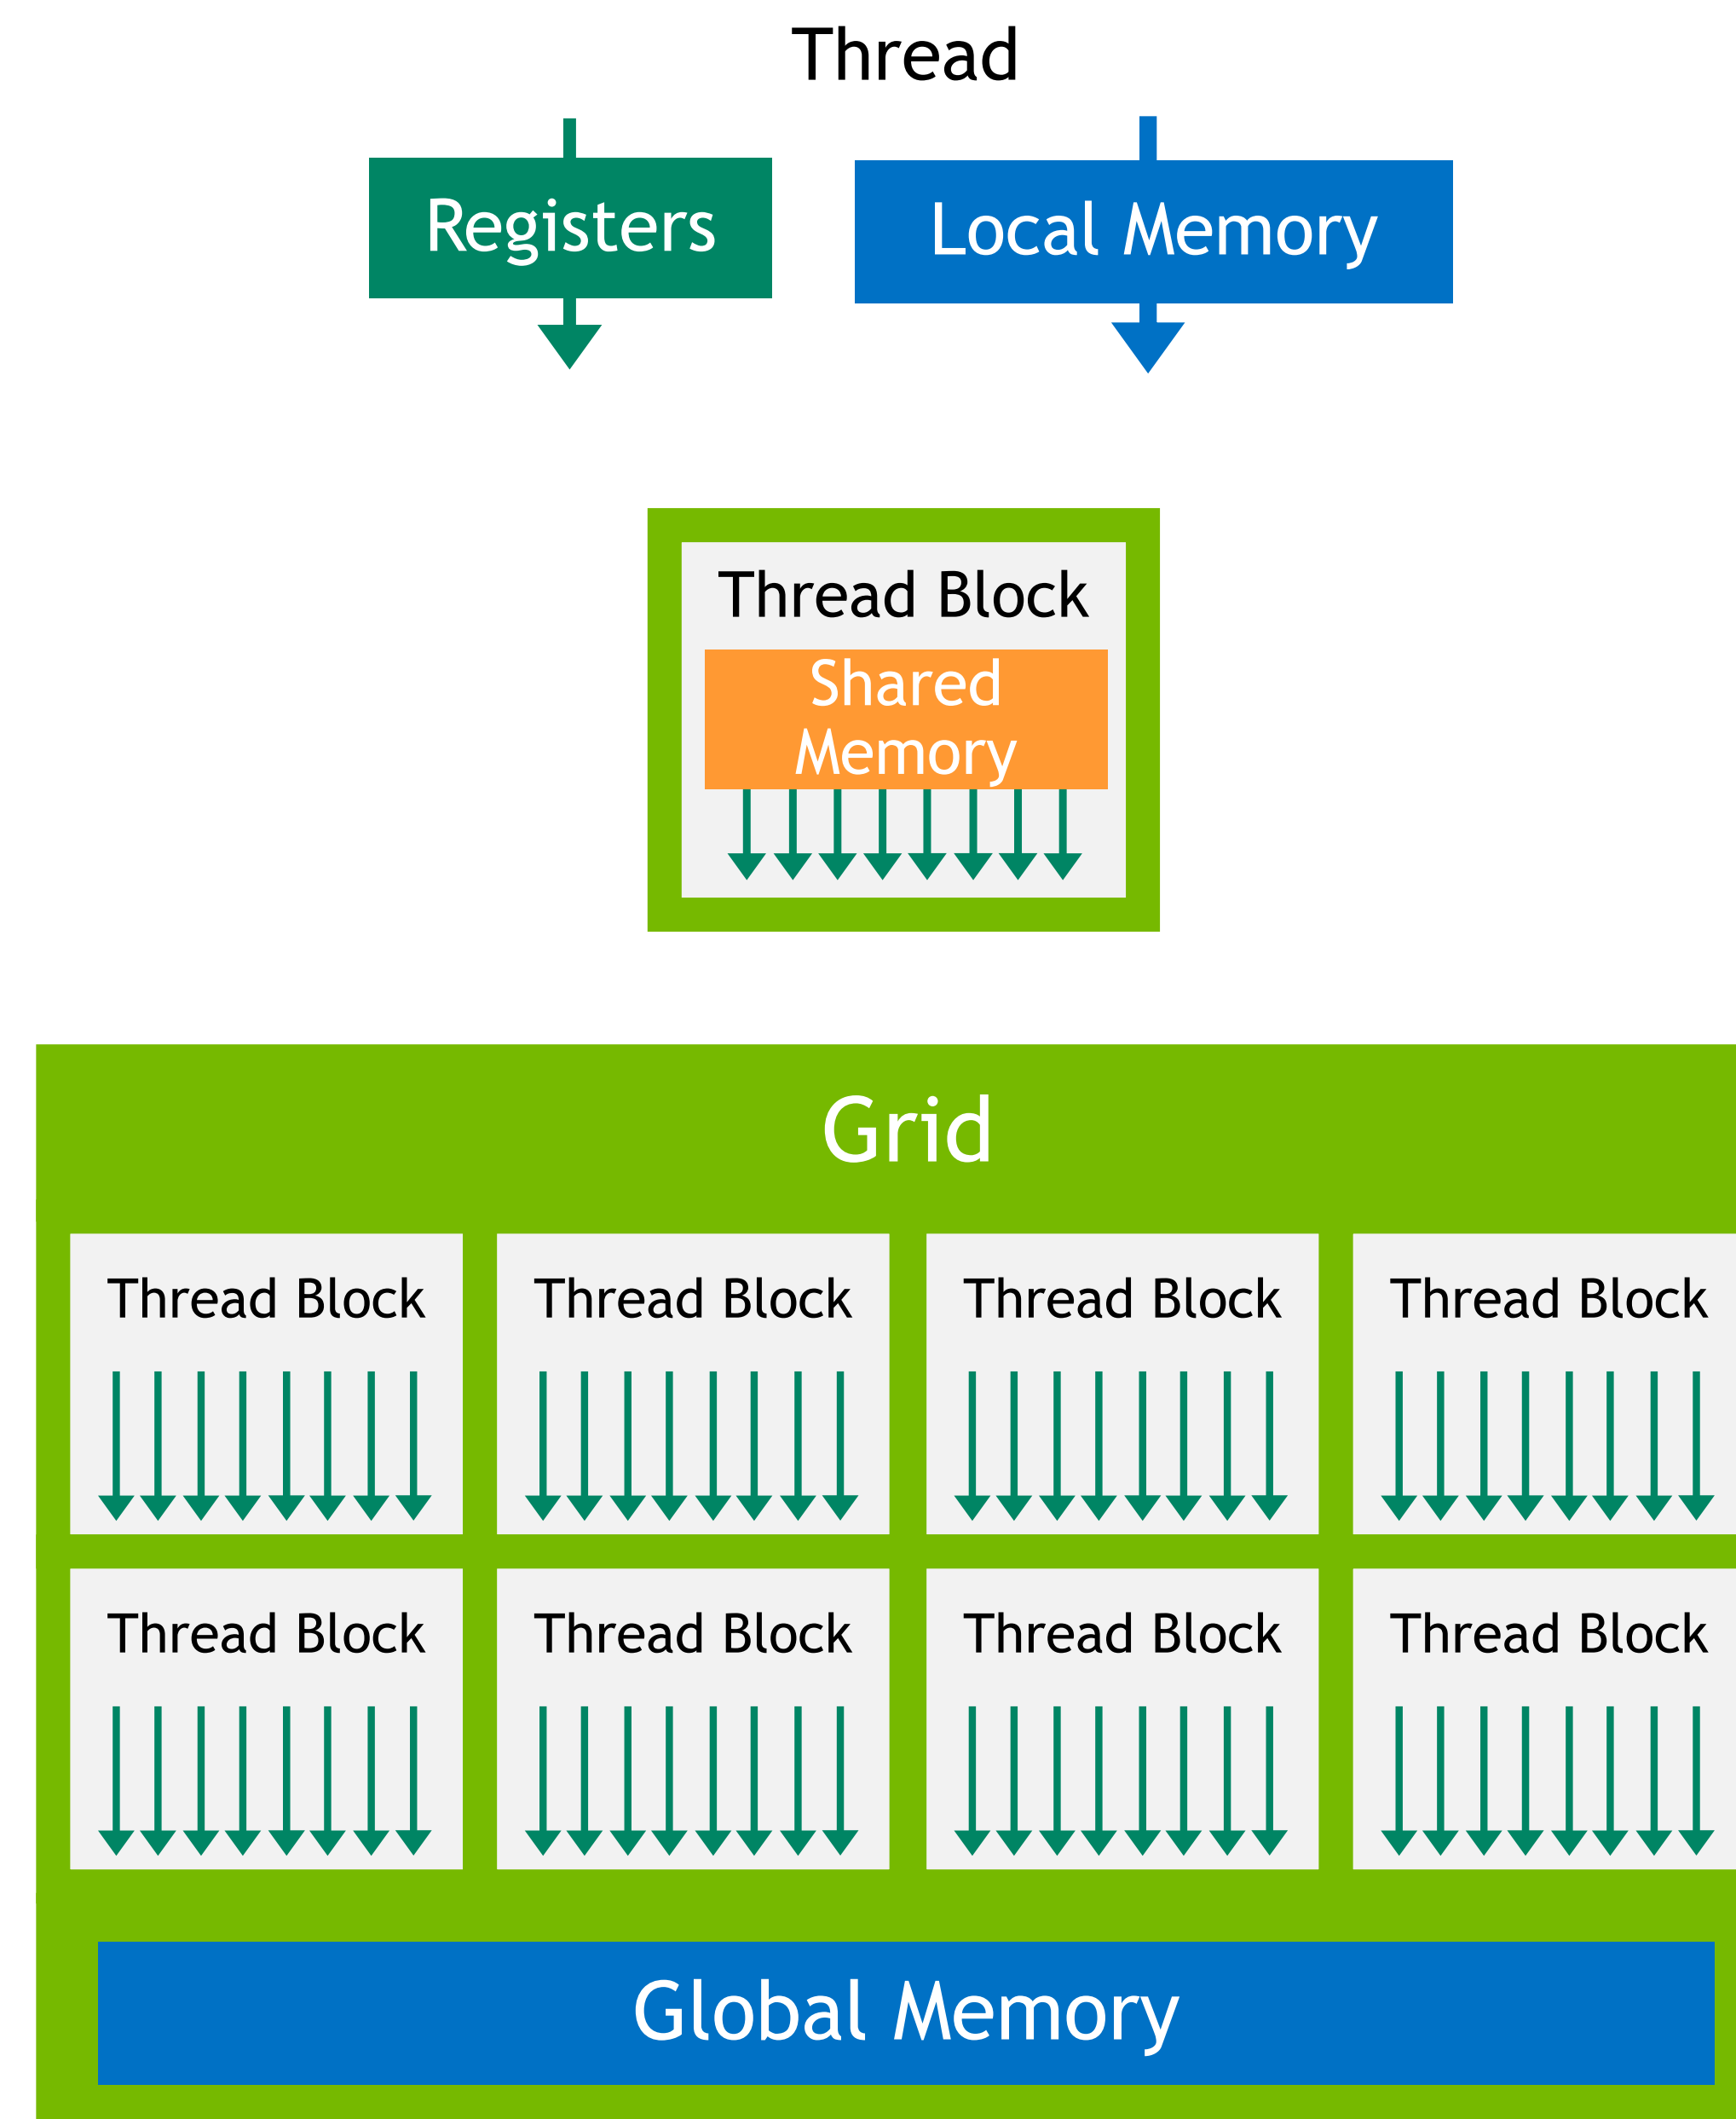
- Per-thread **registers**.
  - Lowest possible latency.
- Per-thread **local** memory.
  - Private storage.
  - Slowest access.

- Per-block **shared** memory.
  - Visible by all threads in a block.
  - Can be used to exchange data between threads in a thread block.
  - Very fast access.

- **Global** memory.
  - Visible by all threads in a grid.
  - Slowest access.

NVIDIA.

Synthronization

## Barriers

CUDA/Software

- **Grid** boundary.
  - Kernel completion.
  - `grid_group::sync()` via Cooperative Groups API
    - Requires the kernel to be launched via the `cudaLaunchCooperativeKernel()` API
    - **Slow!** Avoid unless necessary.

- **Thread-block** boundary.
  - `__syncthreads()`
  - `thread_block::sync()` via Cooperative Groups API
  - **Fast!** The most common synchronization level.

- **Warp** or **sub-warp** boundary.
  - `__syncwarp()`
  - `coalesced_group::sync()` via Cooperative Groups API
  - **Very fast!**

Atomics

Memory spaces

CUDA/Software
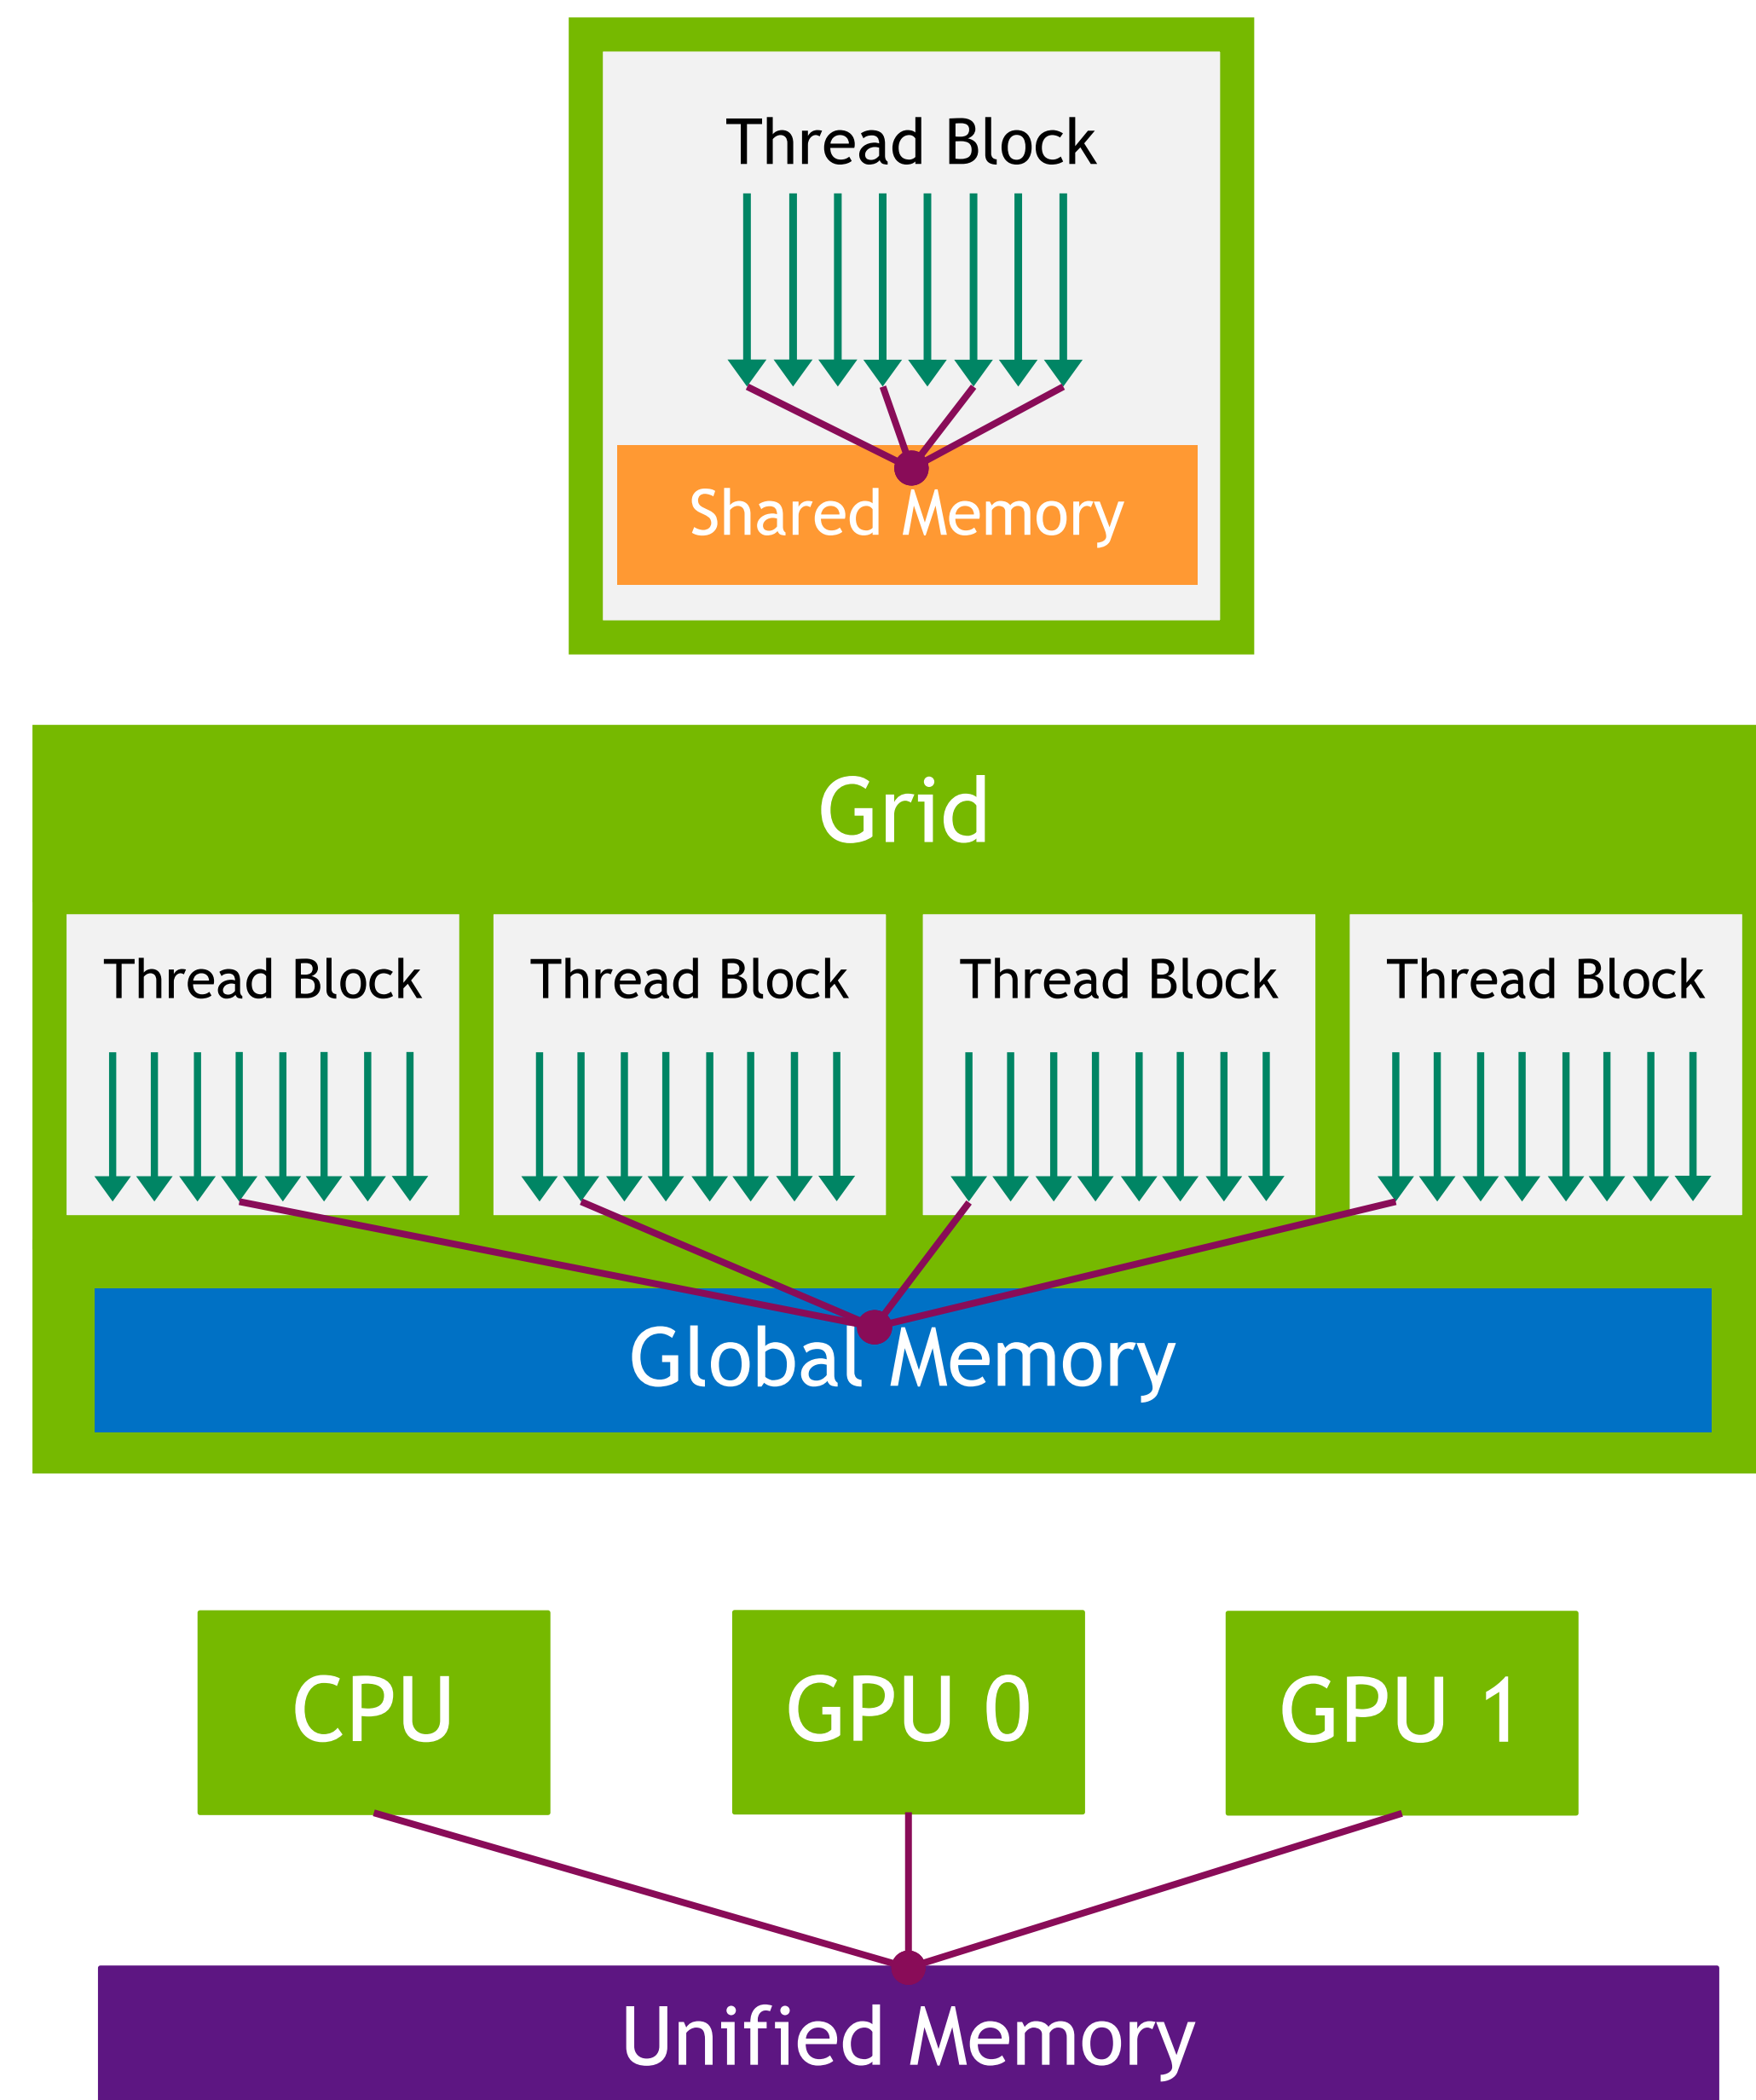
- Read-modify-write operations on 16–, 32-, 64- or 128-bit words.
    - Available as CUDA primitives or C++ atomics through `libcu++` extended API.

- **Shared memory** atomics.

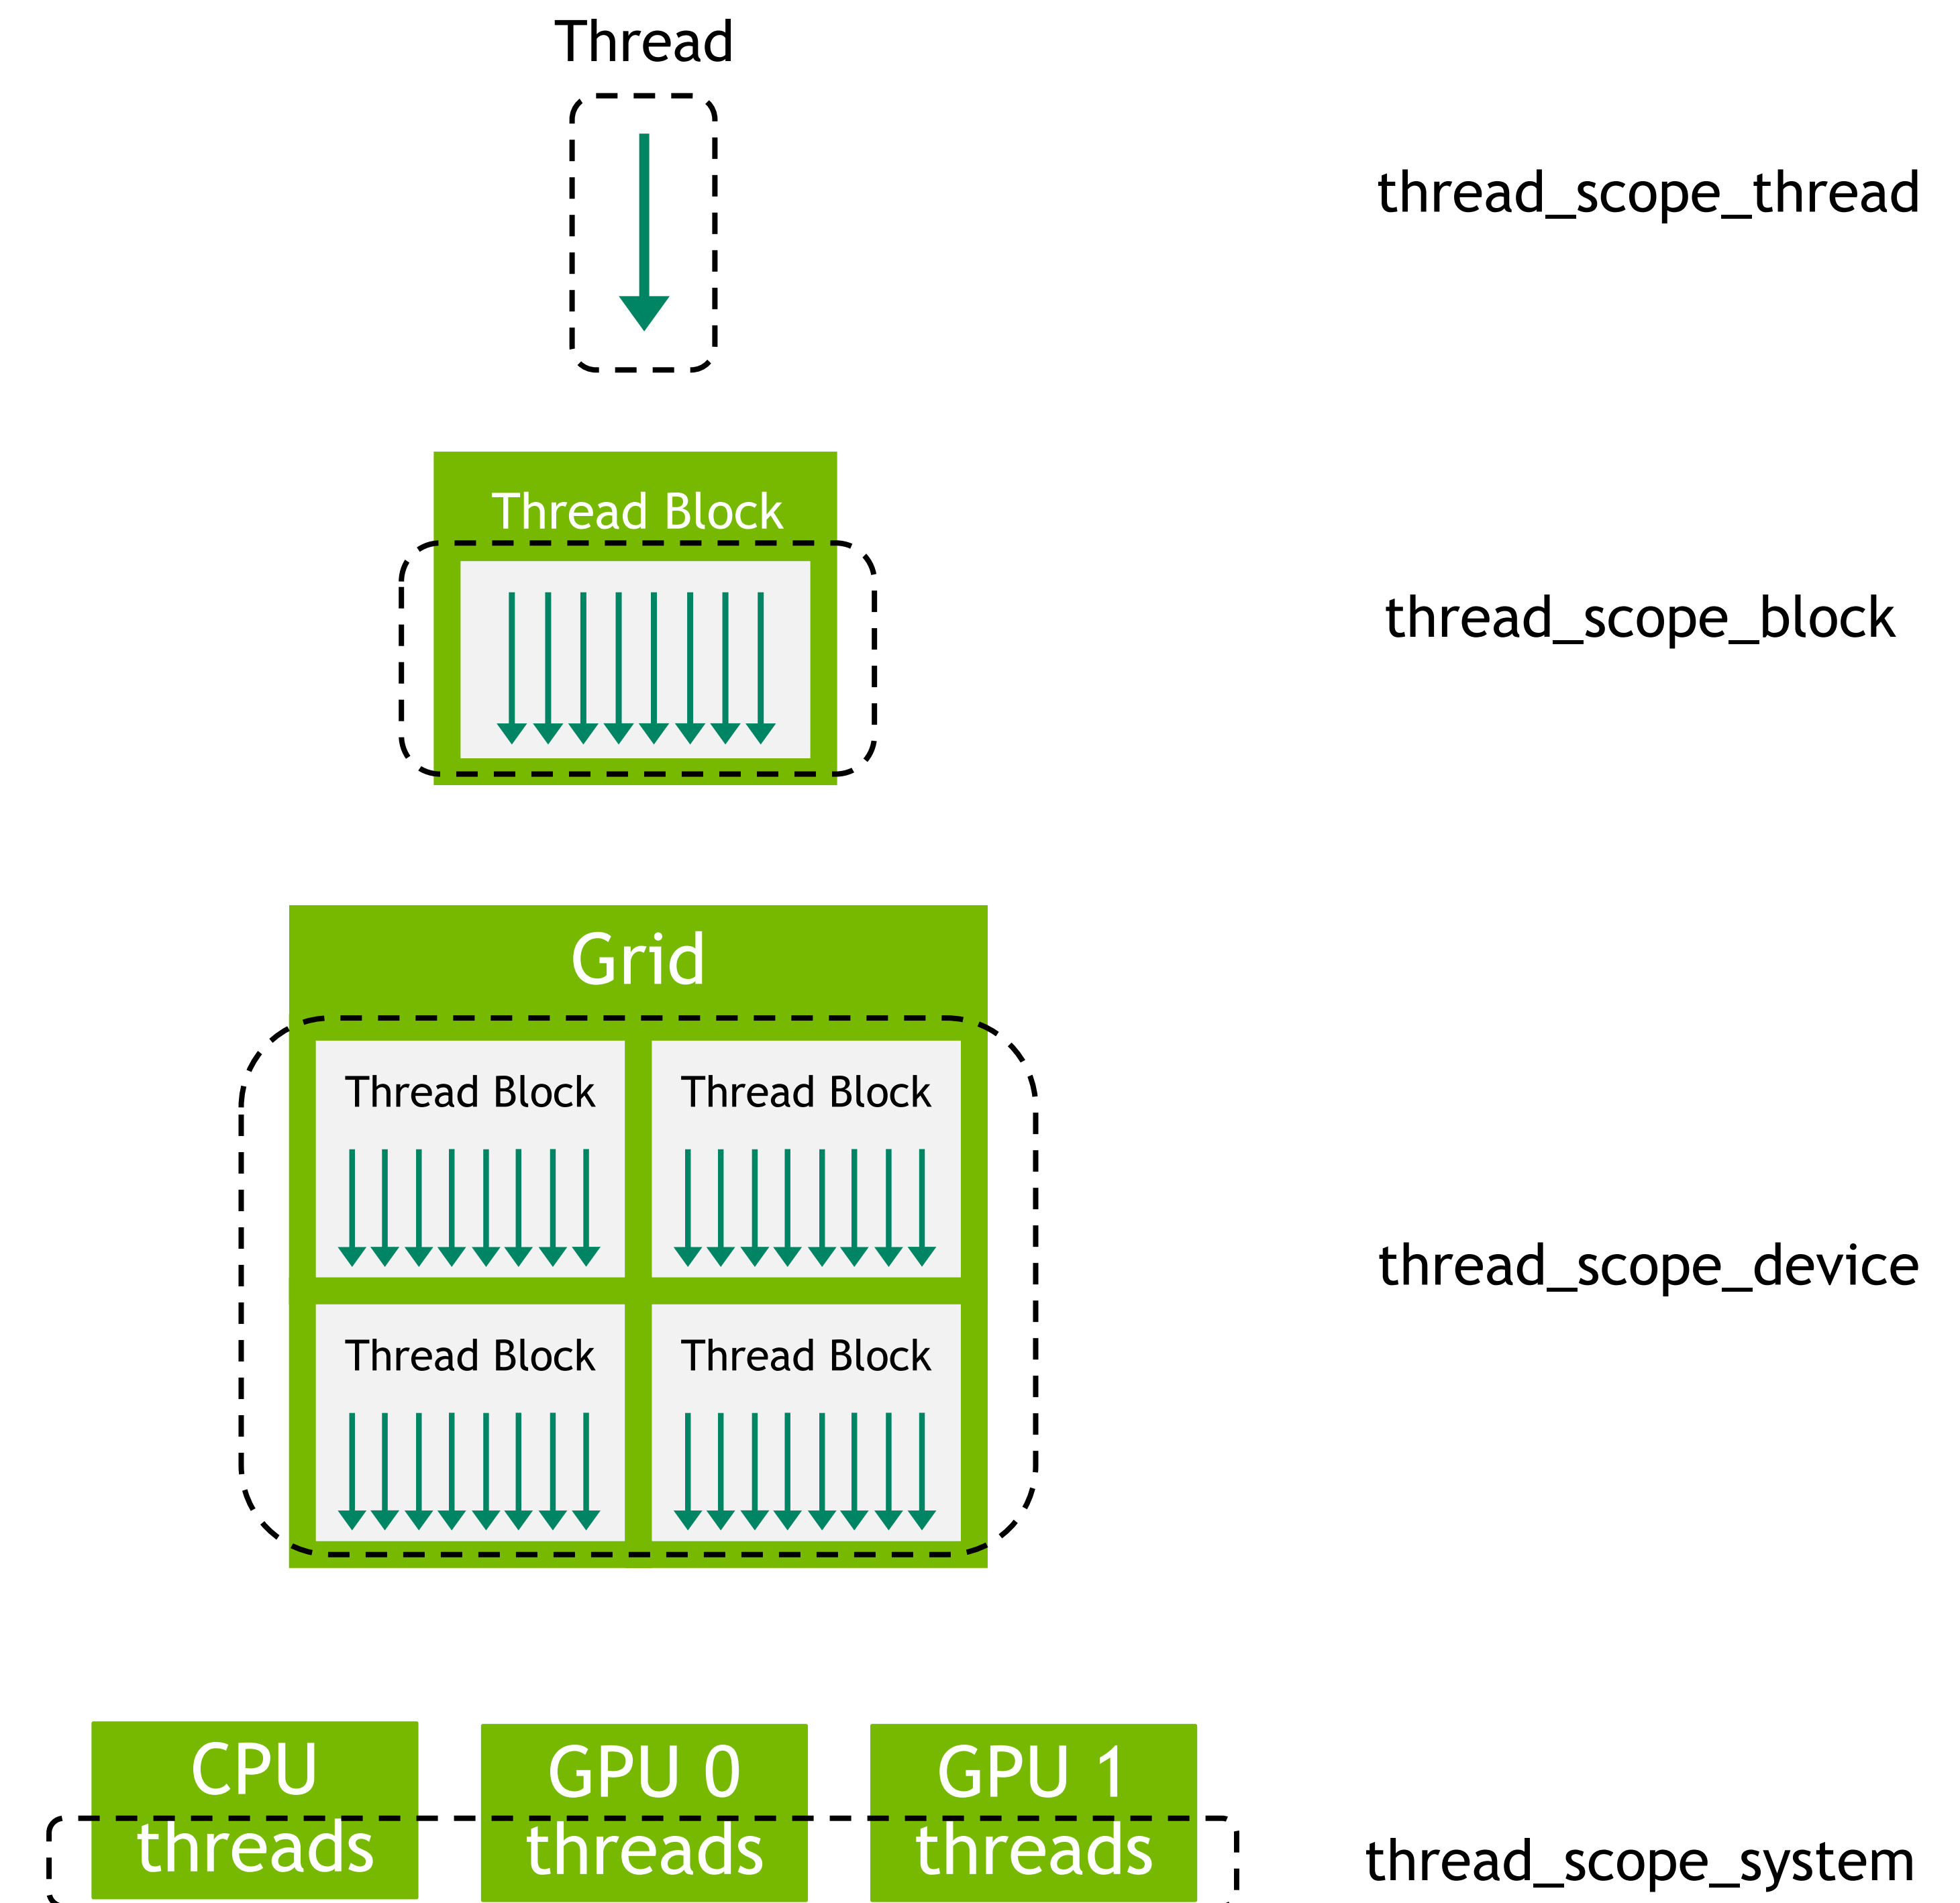- **Global memory** atomics.
    - Facilitated by special hardware in the L2 cache.

- **Unified memory** atomics.

Thread Scopes

- To account for **non-uniform** thread synchronization costs, CUDA has introduced the notion of **thread scopes**.

- A thread scope specifies **which threads can communicate** with each other using a primitive such as an atomic or a barrier.

- Thread scopes are exposed to the programmer in 3 ways:
  - PTX
  - CUDA Math API
  - CUDA C++

- Always use the narrowest scope that ensures correctness of your application.

- More on thread scopes in the GTC session **[S62192]: "Advanced Performance Optimization in CUDA"**.

Thread

thread_scope_thread

Thread Block

thread_scope_block

Grid

Thread Block   Thread Block

Thread Block   Thread Block

thread_scope_device

CPU threads   GPU 0 threads   GPU 1 threads

thread_scope_system

Fundamental Performance Optimizations

Little's Law
For escalators

Our escalator parameters:

- 1 person per step

- A step arrives every 2 seconds
  - **Bandwidth**: 0.5 person/s

- 20 steps tall
  - **Latency** = 40 seconds

One person in flight?

Throughput = 0.025 person/s

Little's Law
For escalators

Our escalator parameters:

- 1 person per step

- A step arrives every 2 seconds
  - **Bandwidth**: 0.5 person/s

- 20 steps tall
  - **Latency** = 40 seconds

How many persons do we need in-flight to saturate bandwidth?

Concurrency = Bandwidth x Latency
$\qquad$ = 0.5 persons/s x 40 s
$\qquad$ = 20 persons

Little's Law
For GPUs

- How to maximize performance?

  1. **Saturate compute units.**
  2. **Saturate memory bandwidth.**

- Need to hide the corresponding latencies to achieve this.

  - Compute latencies.
  - Memory access latencies.

- Latencies can be hidden by having more instructions in flight.

FP32	FP32
FP32	FP32
FP32	FP32
FP32	FP32

FP32 Latency = 24 cycles
8 FP32 ops per cycle

Concurrency = Bandwidth x Latency =
8 x 24 operations in-flight

Hiding Latencies

Increasing in-flight instructions

- Two ways to increase in-flight instructions:

1. Improve **Instruction-Level Parallelism** (ILP).
   - Higher ILP -> more independent instructions per thread.
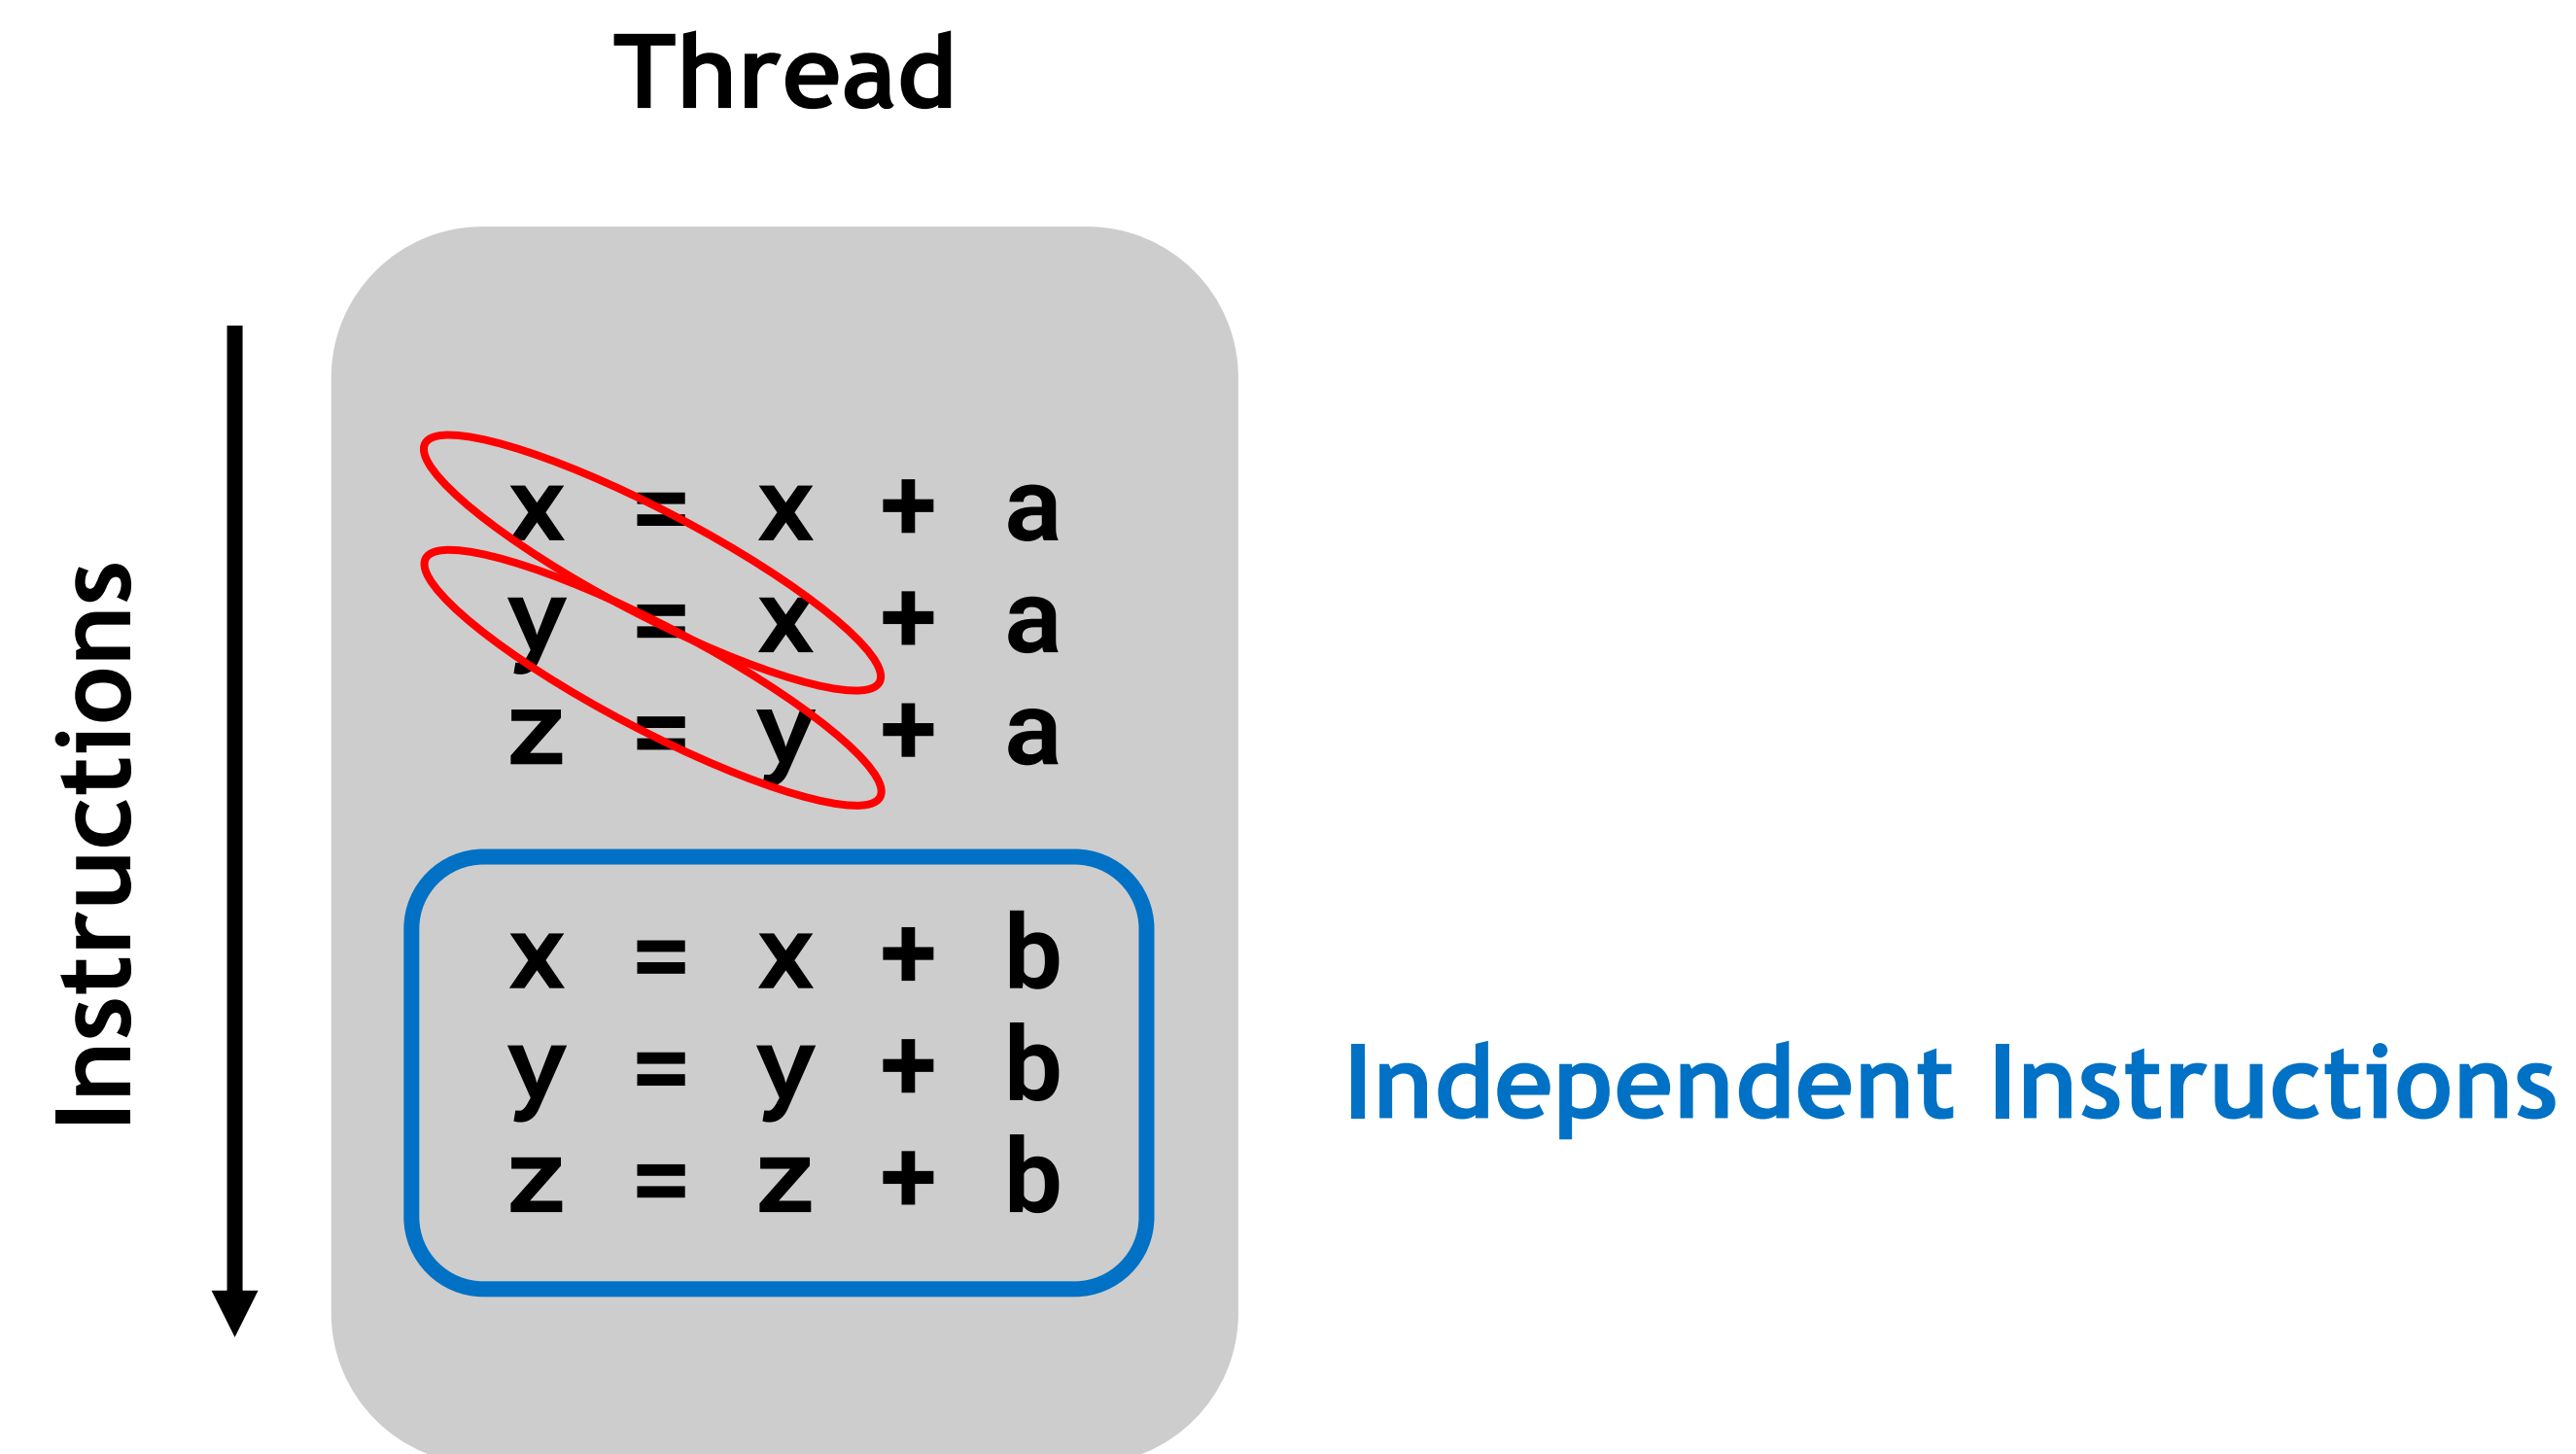
2. Improve **Thread-Level Parallelism** (TLP).
   - Higher TLP -> more threads -> more independent instructions per kernel.

Thread

Instructions

```
x = x + a
y = x + a
z = y + a

x = x + b
y = y + b
z = z + b
```

Independent Instructions

Thread 0

```
x = x + a

x = x + b
```

Thread 1

```
y = y + a

y = y + b
```
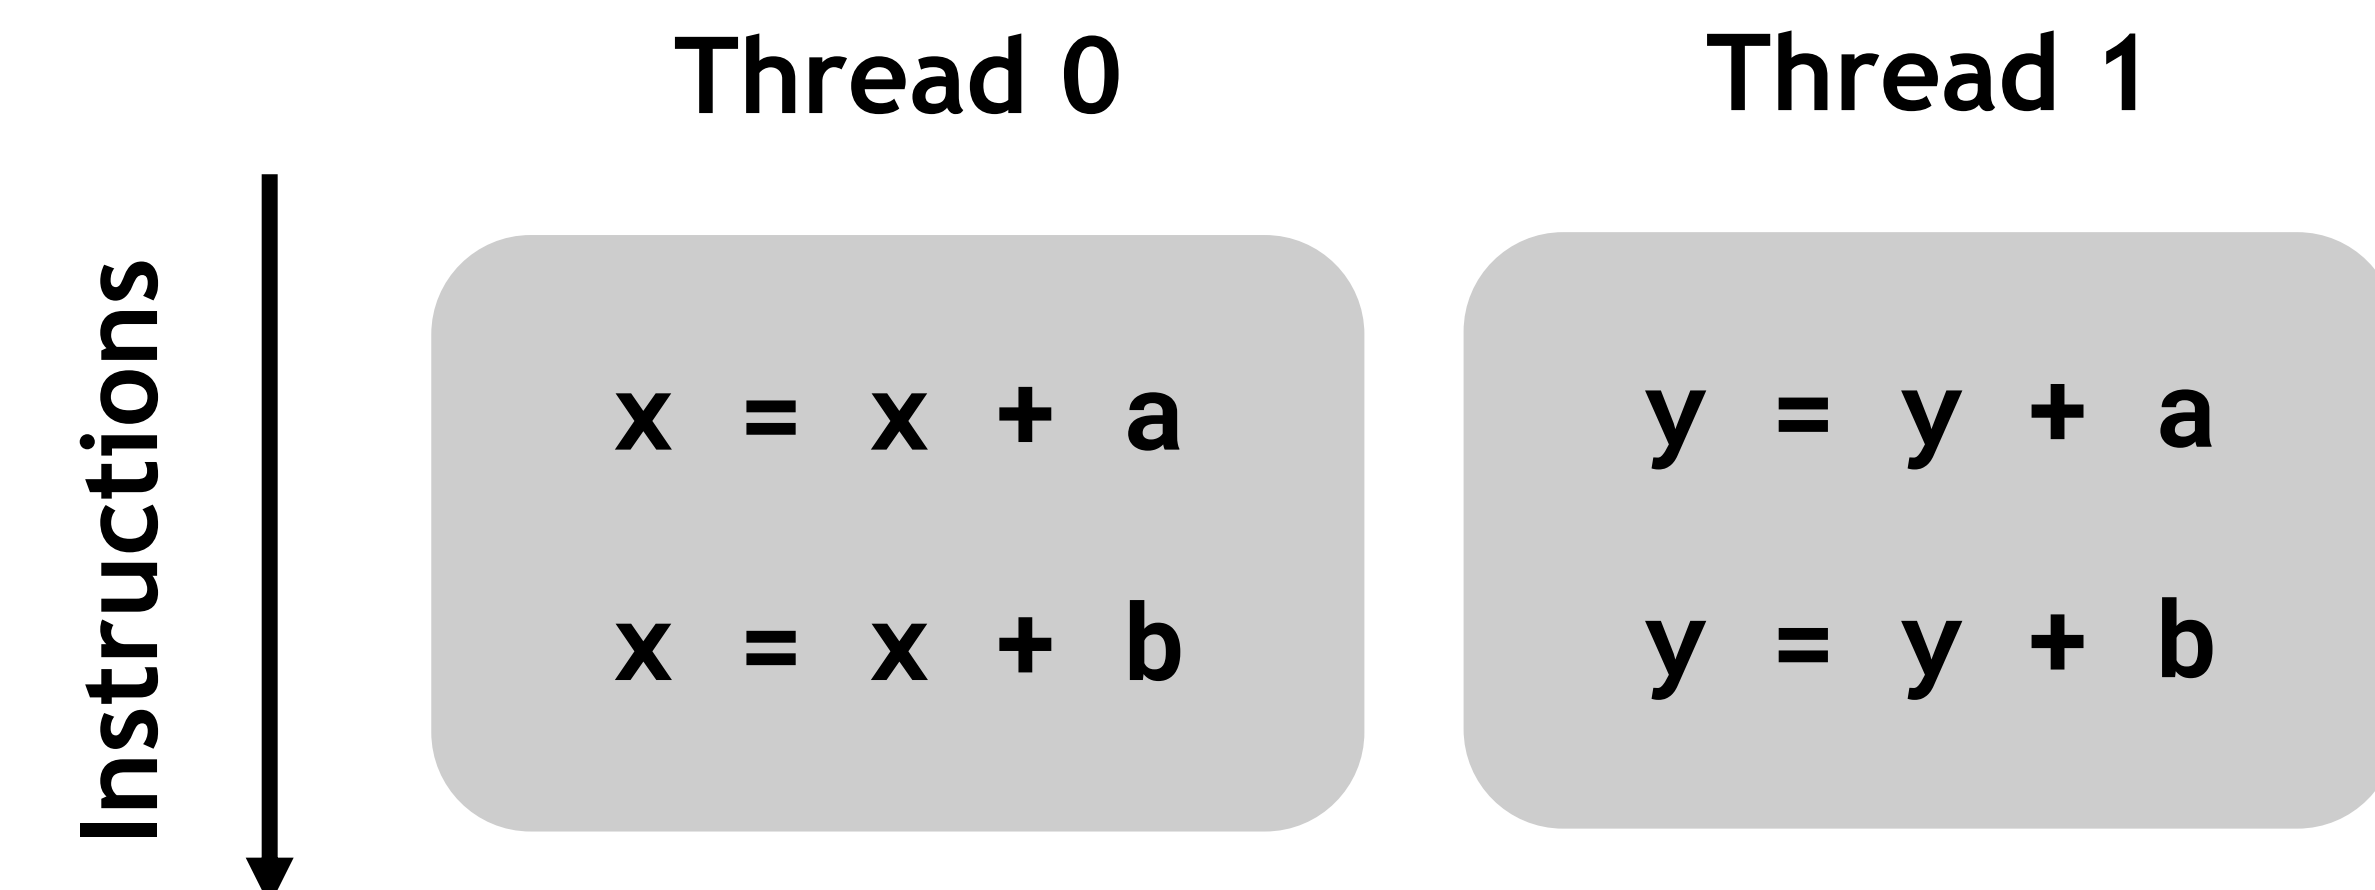
Instructions

Instruction Issue

- Assumptions
  - LDG/STG
    - Dependent Issue Rate: 1000 cycles
    - Issue Rate: 1 cycle
  - FP32 pipeline
    - Dependent Issue Rate: 4 cycles
    - Issue Rate: 2 cycles
  - 1 available warp per scheduler

Cycle	N	N+1	N+2		N+1002		N+1006
	LDG	LDG	LDG	(stall)	FFMA	(stall)	STG

Total cycles = 1006

```
__global__
void kernel(const float * __restrict__ a,
            const float * __restrict__ b,
            float * __restrict__ c)
{
  int idx = blockIdx.x * blockDim.x + threadIdx.x;

  c[idx] += a[idx] * b[idx];
}
```

load a
load b
load c   } 12 bytes in-flight

fma c, a, b

store c   } 4 bytes in-flight
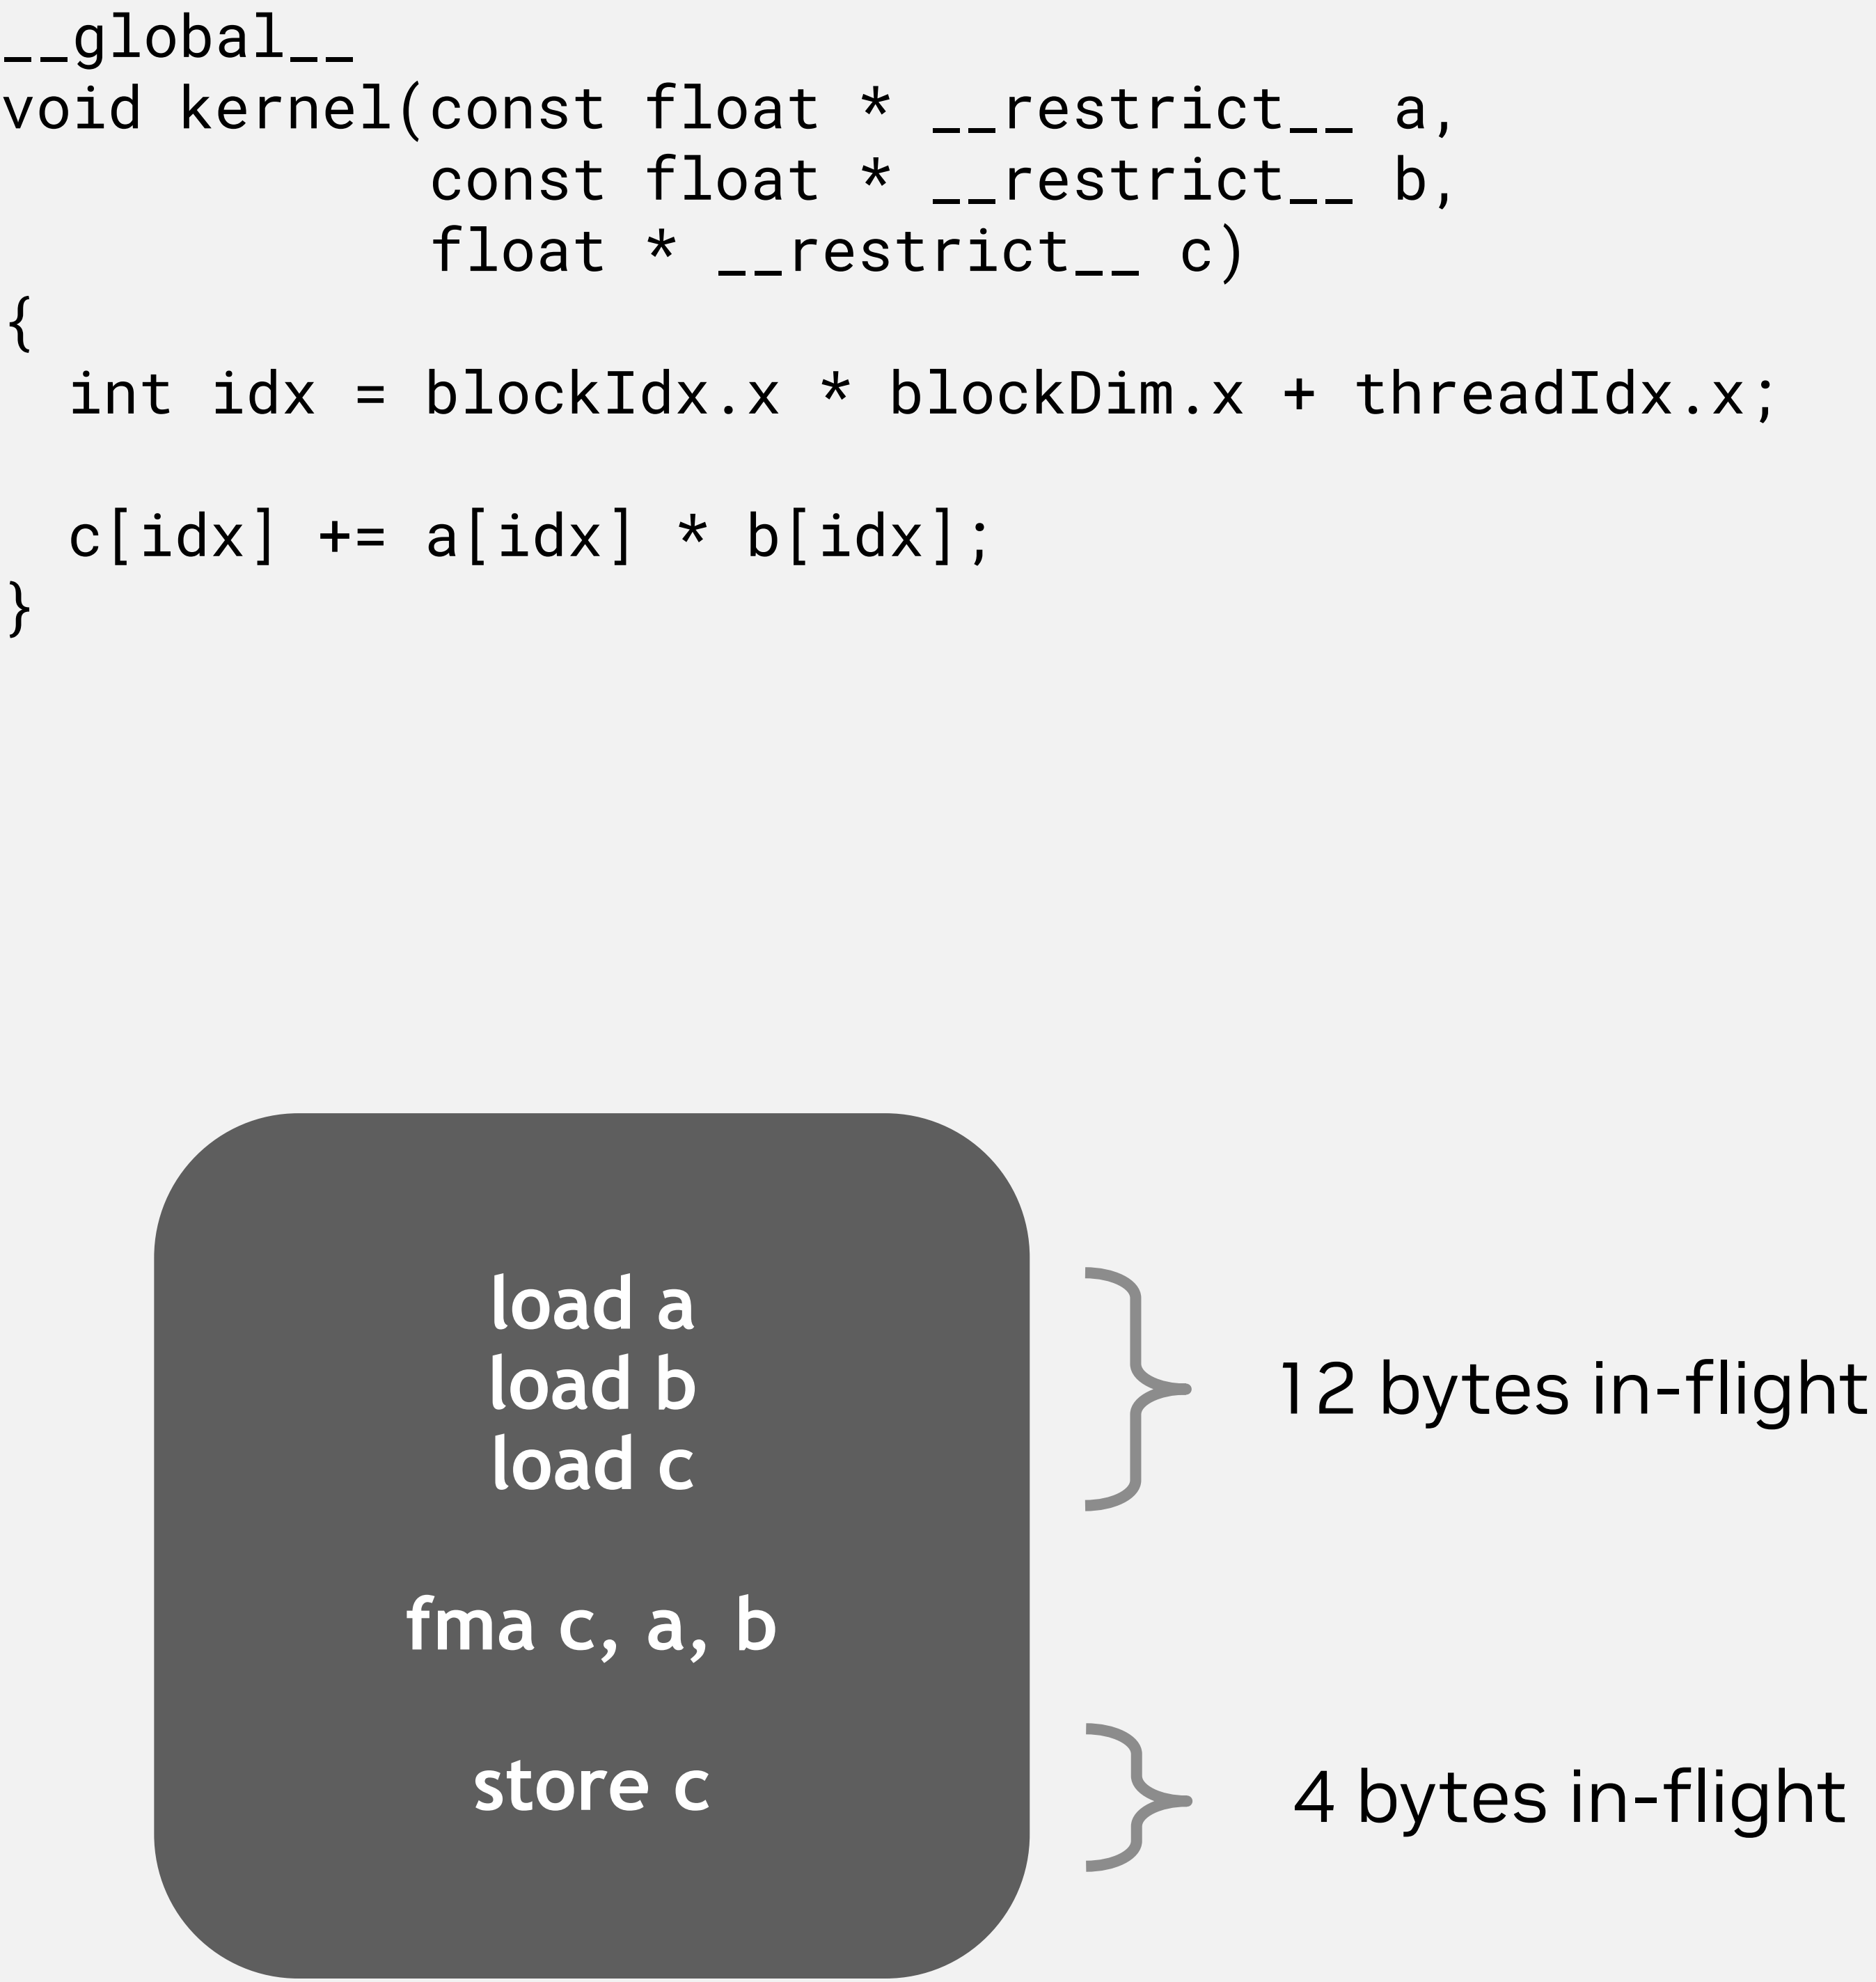
Increasing ILP

Computing 2 elements per thread – version #1

- Every thread computes 2 elements using a **grid stride.**

Cycle

N	N+1	N+2	N+3	N+4		N+1002		N+1006
LDG	LDG	LDG	LDG	LDG	(stall)	FFMA	(stall)	STG

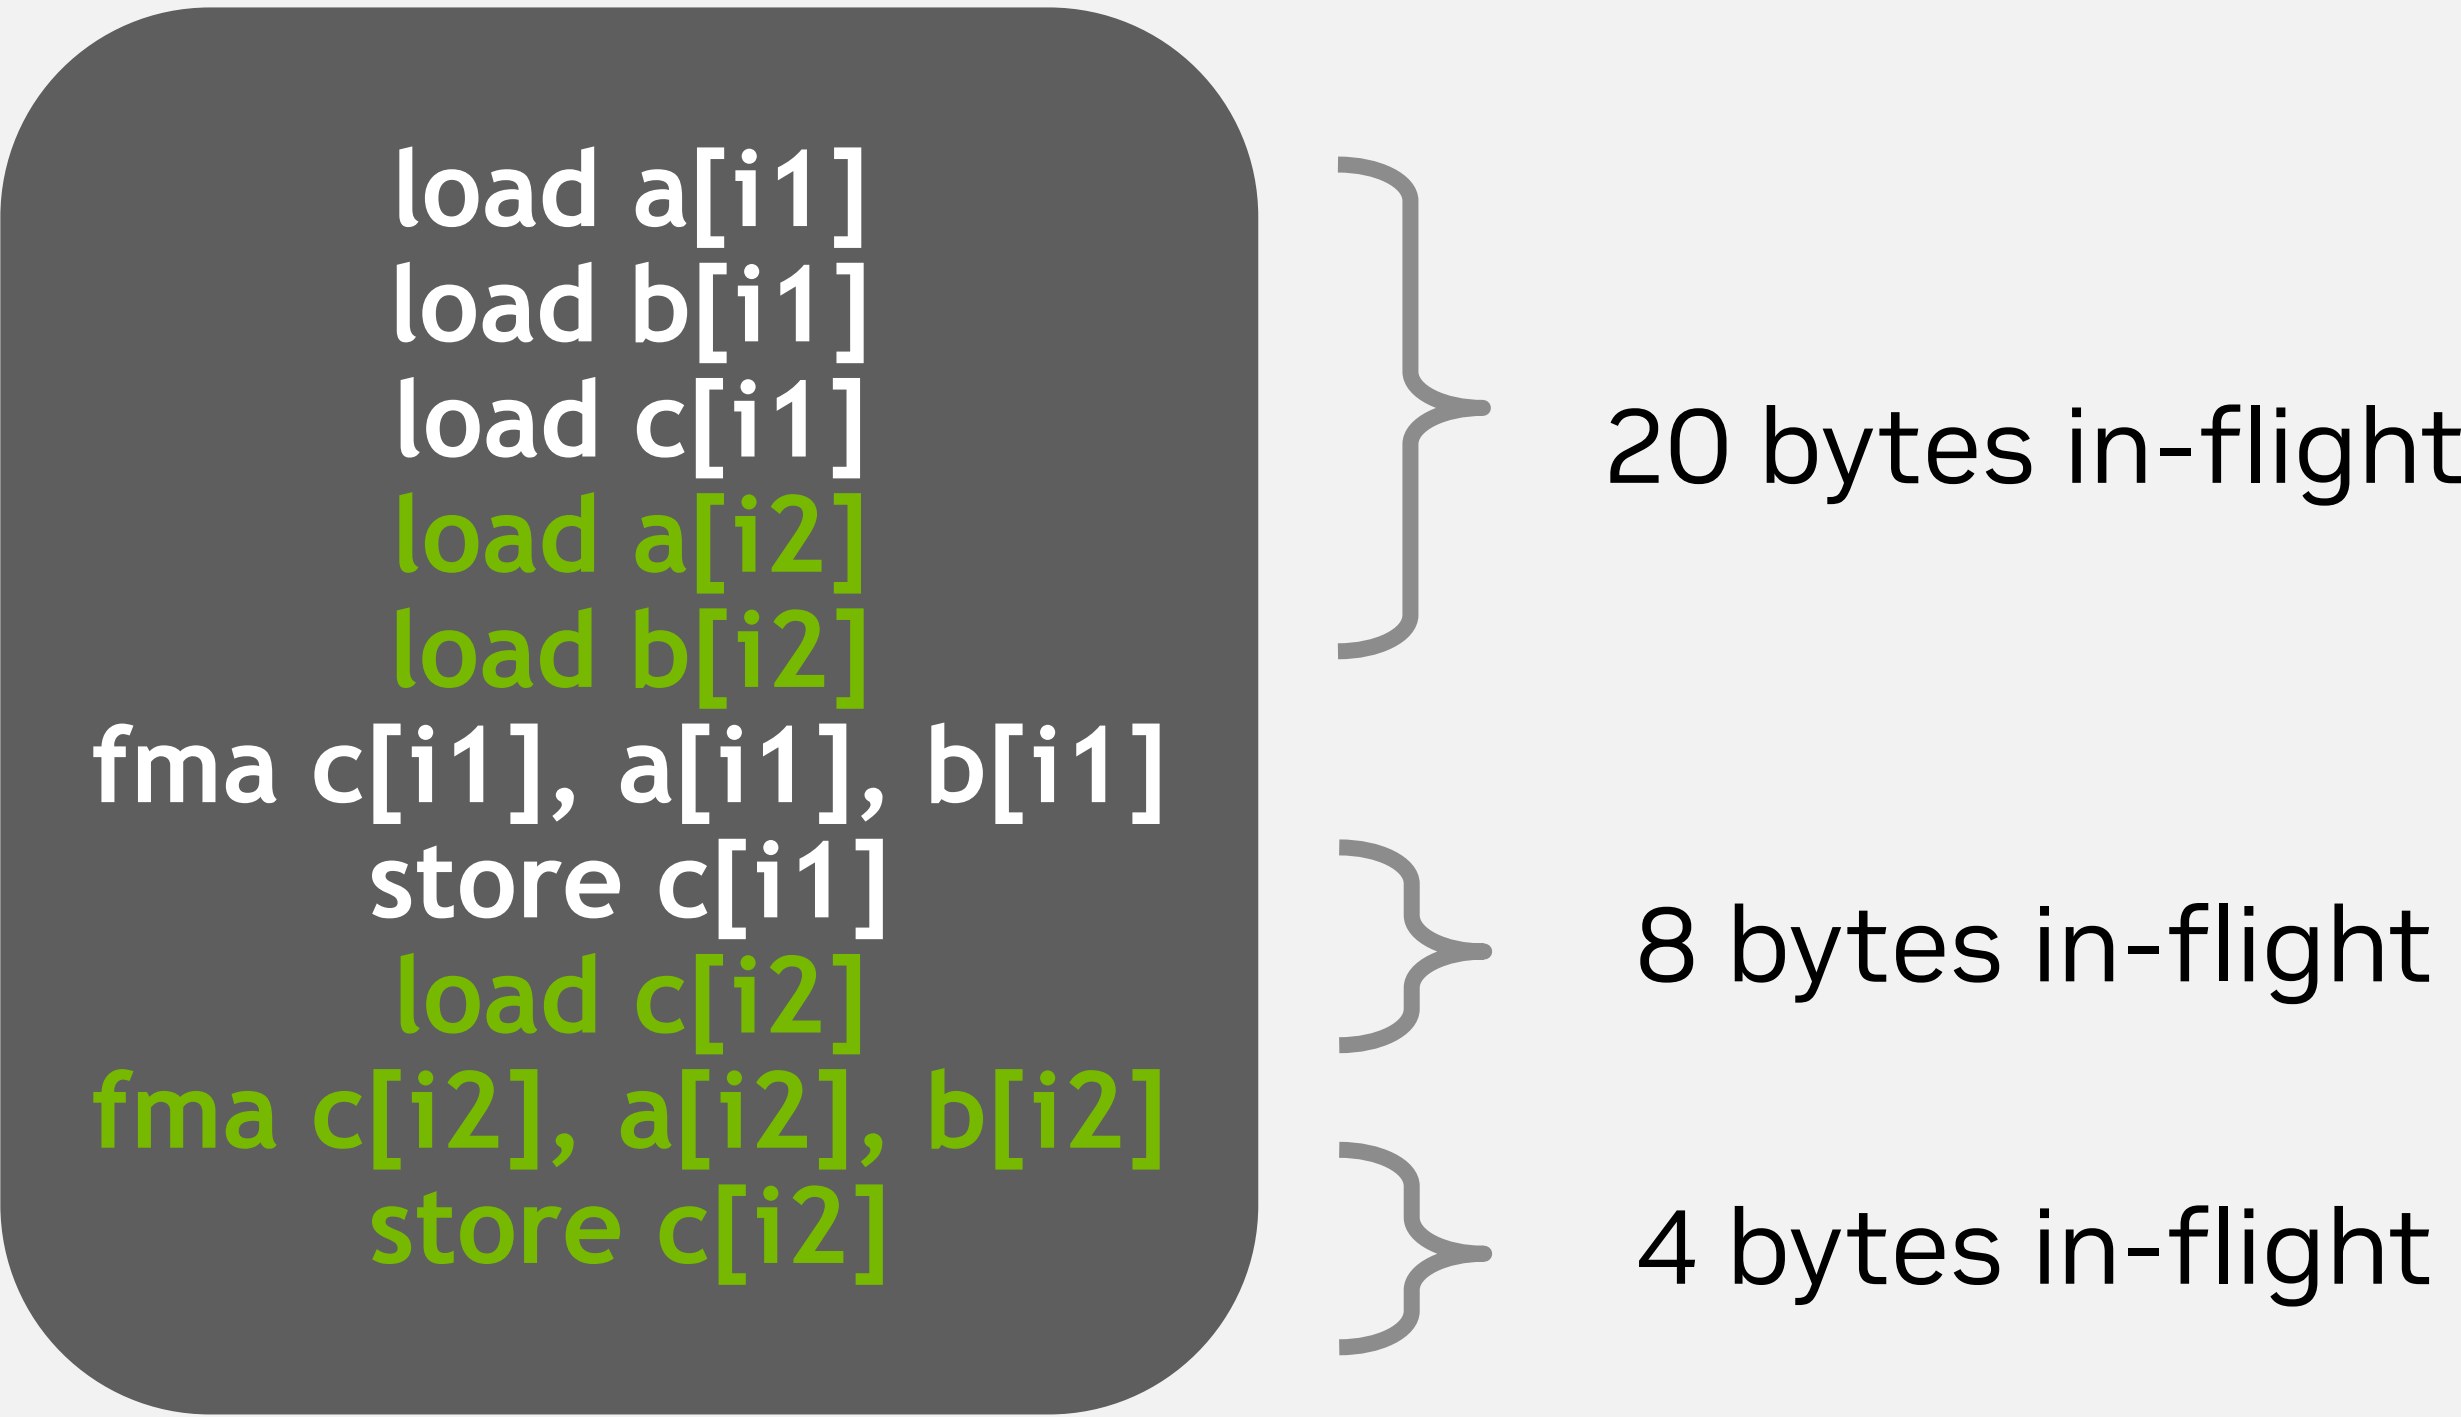| N+1007		N+2007		N+2011
LDG	(stall)	FFMA	(stall)	STG

Total cycles = 2011

2x the amount of work in **2x** more cycles!

```cpp
__global__
void kernel(const float * __restrict__ a,
            const float * __restrict__ b,
            float * __restrict__ c)
{
  int tid = blockIdx.x * blockDim.x + threadIdx.x;
  int stride = blockDim.x * gridDim.x;

  #pragma unroll 2
  for (int i = 0; i < 2; i++) {
    const int idx = tid + i * stride;
    c[idx] += a[idx] * b[idx];
  }
}
```

load a[i1]
load b[i1]
load c[i1]
load a[i2]
load b[i2]   } 20 bytes in-flight
fma c[i1], a[i1], b[i1]
store c[i1]
load c[i2]   } 8 bytes in-flight
fma c[i2], a[i2], b[i2]
store c[i2]   } 4 bytes in-flight

Increasing ILP

Computing 2 elements per thread – version #2

- Every thread computes 2 elements using a **constant block stride.**

Cycle

N	N+1	N+2	N+3	N+4	N+5		N+1002	
LDG	LDG	LDG	LDG	LDG	LDG	(stall)	FFMA	(stall)

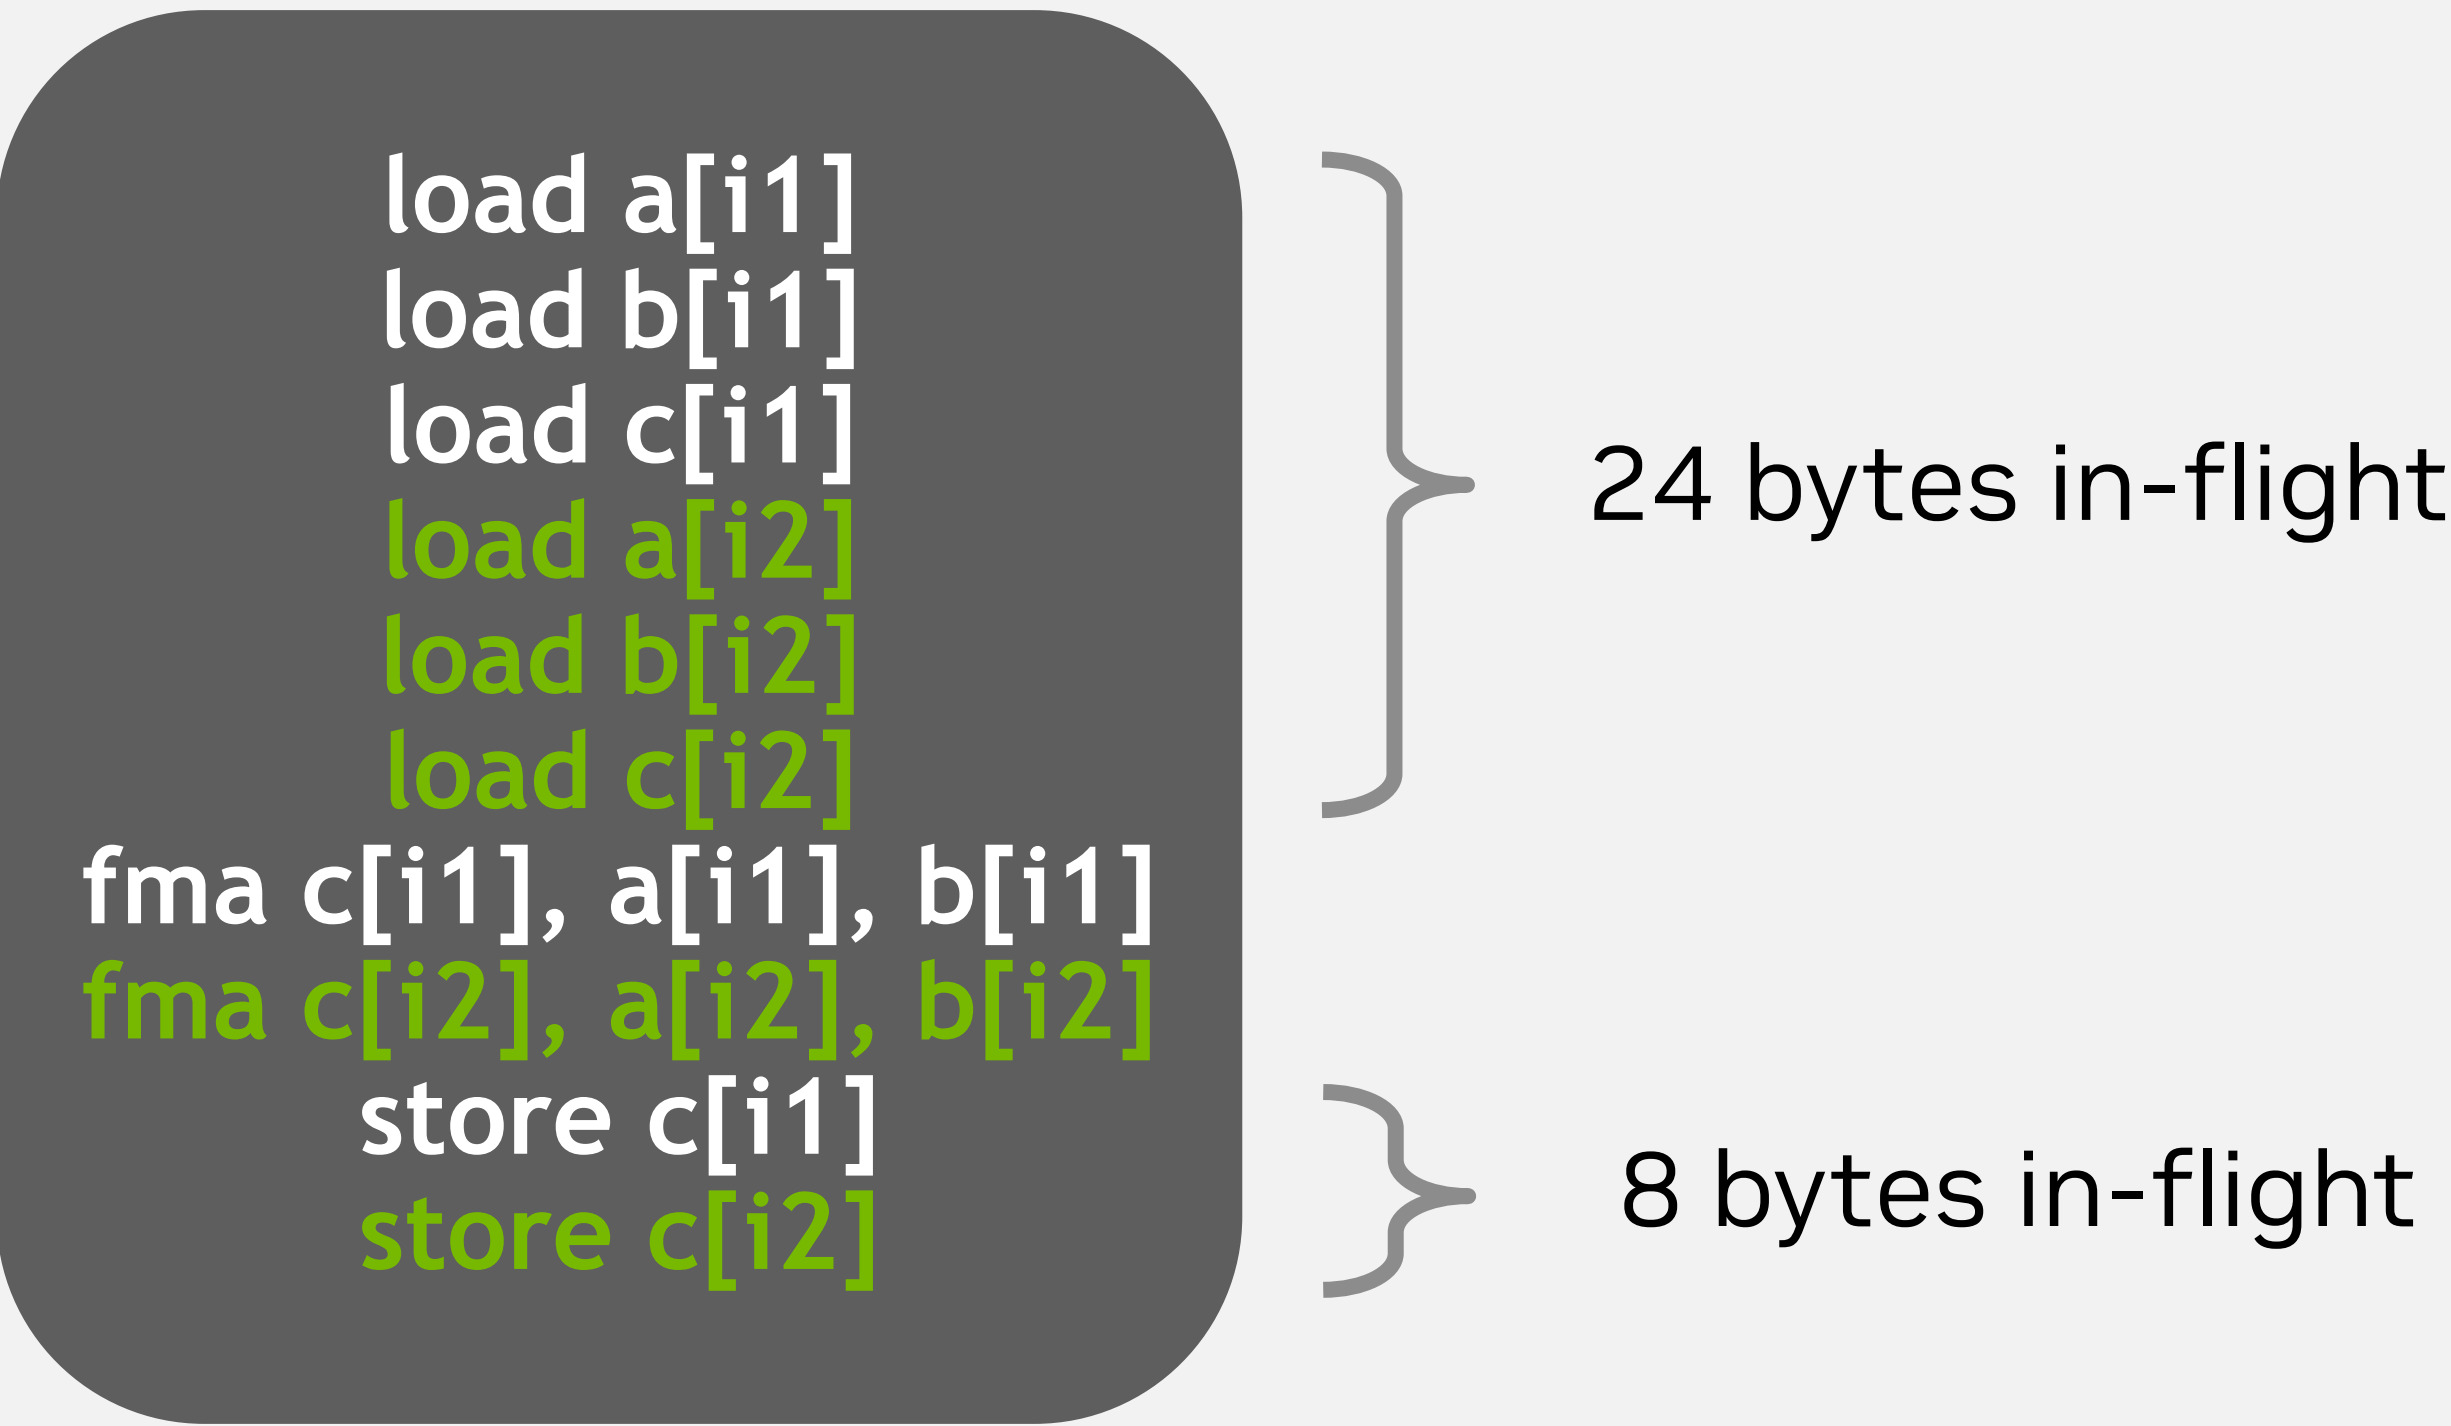| N+1004		N+1006		N+1008
FFMA	(stall)	STG	(stall)	STG

Total cycles = 1008

2x the amount of work in the **~same** number of cycles!
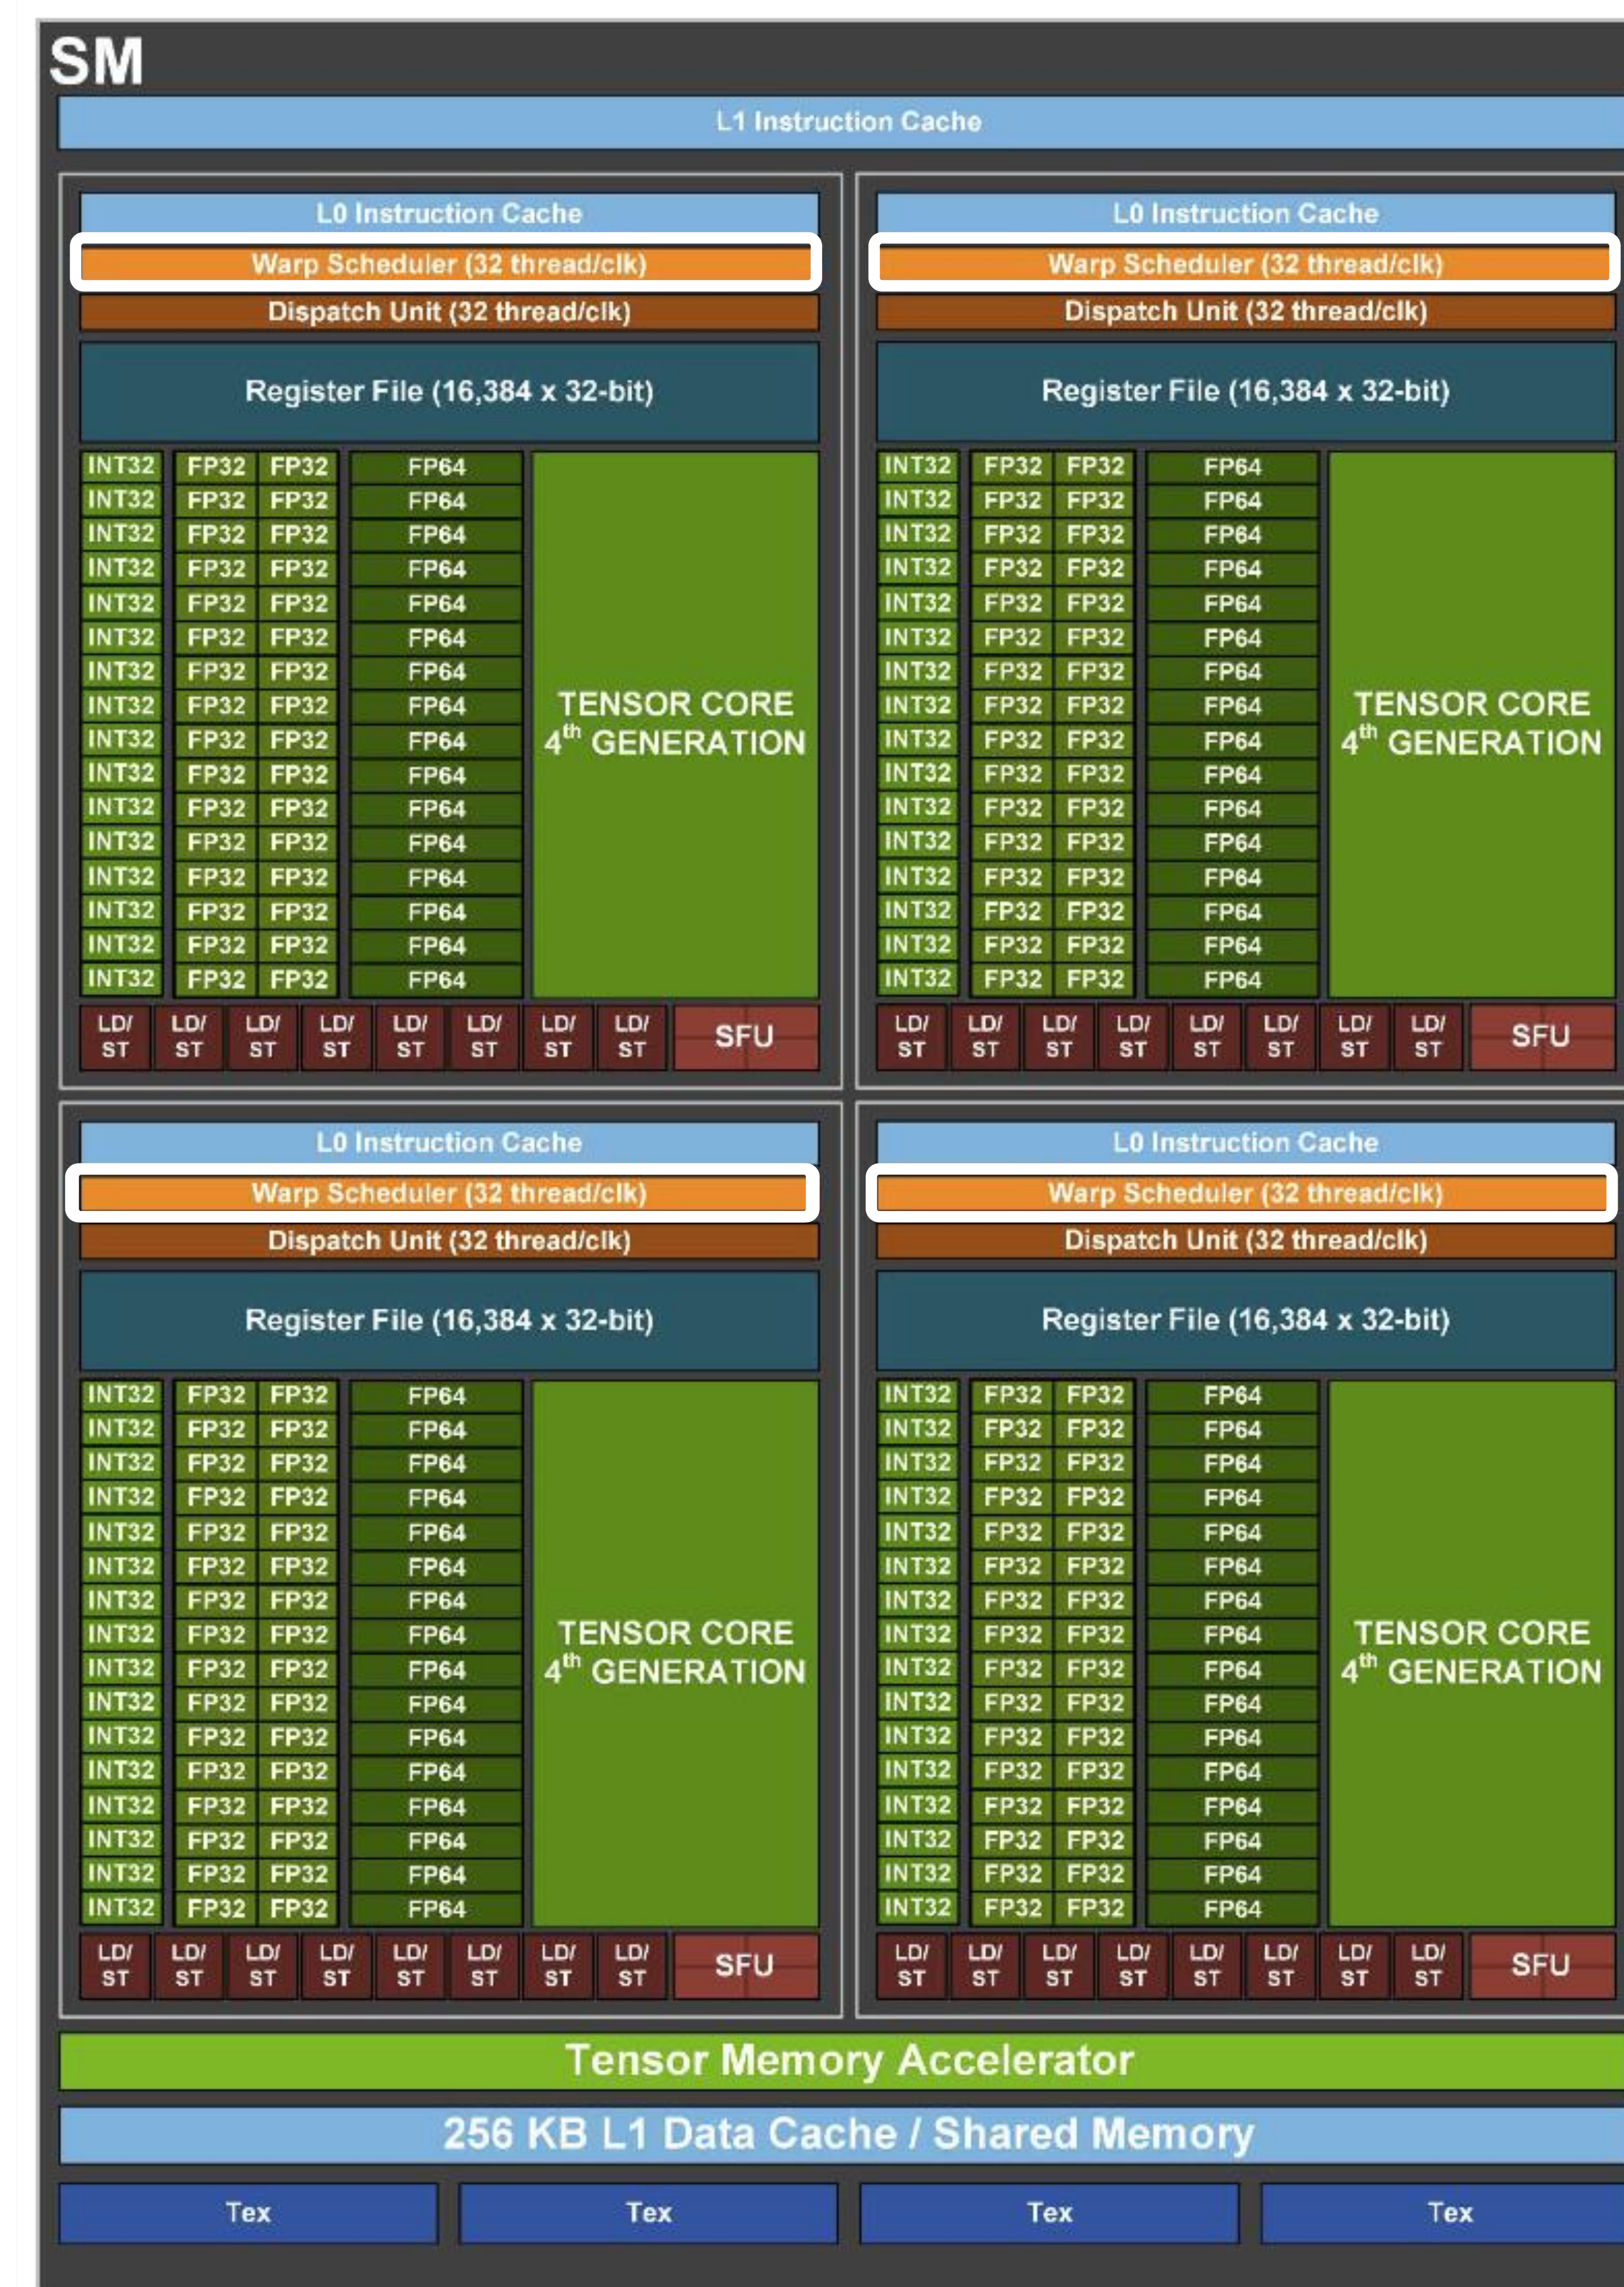
```
#define THREAD_BLOCK_DIM 128

__global__
void kernel(const float * __restrict__ a,
            const float * __restrict__ b,
            float * __restrict__ c)
{
  int tid = blockIdx.x * blockDim.x + threadIdx.x;
  int off = 2 * THREAD_BLOCK_DIM * blockIdx.x + threadIdx.x;

  #pragma unroll 2
  for (int i = 0; i < 2; i++) {
    const int idx = off + i * THREAD_BLOCK_DIM;
    c[idx] += a[idx] * b[idx];
  }
}
```

```
        load a[i1]
        load b[i1]
        load c[i1]
        load a[i2]                    24 bytes in-flight
        load b[i2]
        load c[i2]
    fma c[i1], a[i1], b[i1]
    fma c[i2], a[i2], b[i2]
        store c[i1]                   8 bytes in-flight
        store c[i2]
```

Warp Scheduling

Hopper SM

- 4 warp schedulers per SM.

- Each scheduler manages a pool of warps.
  - Hopper: 16 warp slots per scheduler.

- Each scheduler can issue 1 warp per cycle.

Warp Scheduling
Mental model

Active Warp States:

Stalled
Waiting on:
    an instruction fetch,
    a memory dependency,
    an execution dependency, or
    a synchronization barrier.

Eligible
Ready to issue an instruction.

Selected
Eligible that is selected to issue
an instruction.

Warp Slots

N
...
5
4
3
2
1
0

Warps occupying scheduler slots are considered active

Warp Scheduling
Mental model

Active Warp States:

■ **Stalled**
Waiting on:
    an instruction fetch,
    a memory dependency,
    an execution dependency, or
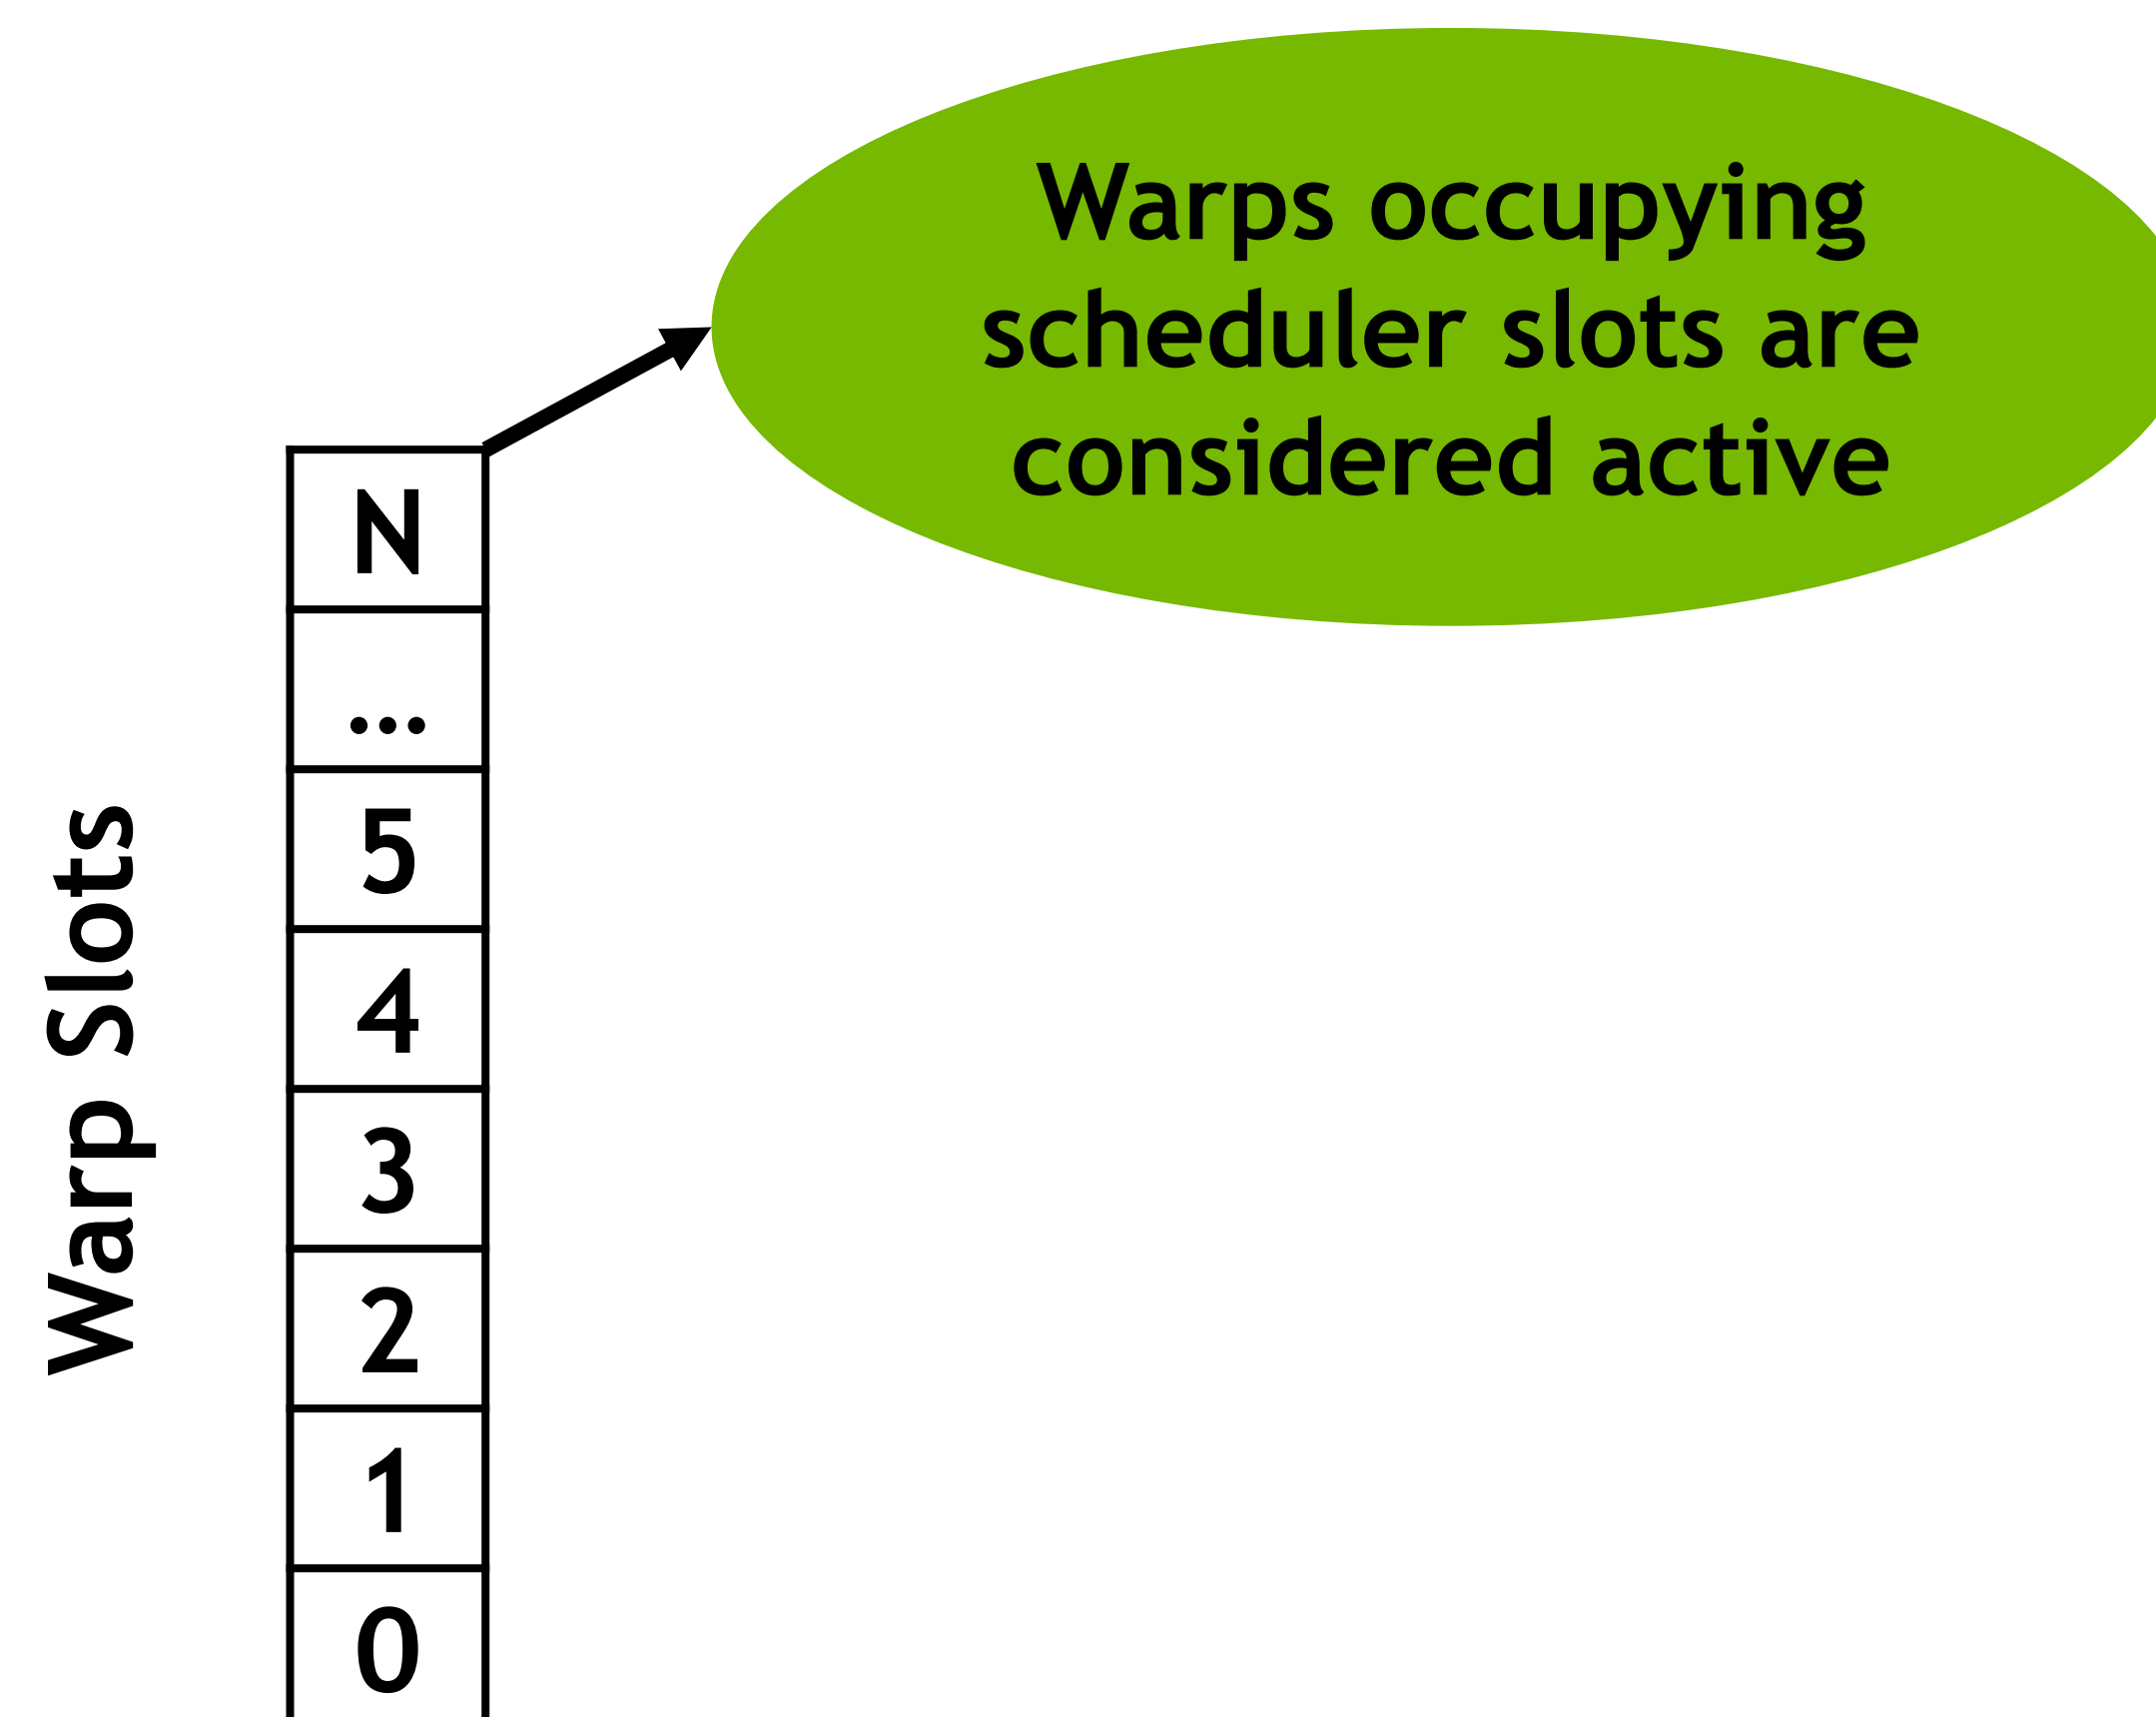   a synchronization barrier.
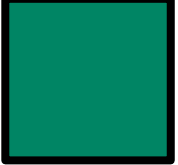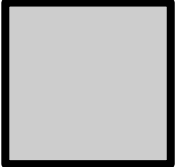
□ **Eligible**
Ready to issue an instruction.

■ **Selected**
Eligible that is selected to issue
an instruction.

Cycle: N

Warp Slots

N
...
5
4
3
2
1
0

Issue slot:

Warp Scheduling
Mental model

Active Warp States:

■ **Stalled**
Waiting on:
   an instruction fetch,
   a memory dependency,
   an execution dependency, or
   a synchronization barrier.

□ **Eligible**
Ready to issue an instruction.

■ **Selected**
Eligible that is selected to issue
an instruction

Cycle: N

Warp Slots

N
...
5
4
3
2
1
0

Issue slot: 2

Each cycle: out of all eligible warps, select one to issue on that cycle

Warp Scheduling
Mental model

Active Warp States:

Stalled
Waiting on:
    an instruction fetch,
    a memory dependency,
    an execution dependency, or
    a synchronization barrier.

Eligible
Ready to issue an instruction.

Selected
Eligible that is selected to issue
an instruction

Cycle: N   N+1

Warp Slots

N
...
5
4
3
2
1
0

Issue slot: 2

Warp selected at cycle N is not eligible in cycle N+1.
E.g., instructions with longer latencies.

Warp Scheduling
Mental model

Active Warp States:

Stalled
Waiting on:
    an instruction fetch,
    a memory dependency,
   an execution dependency, or
    a synchronization barrier.

Eligible
Ready to issue an instruction.

Selected
Eligible that is selected to issue
an instruction

Cycle: N   N+1

Warp Slots

| N |
| ... |
| 5 |
| 4 |
| 3 |
| 2 |
| 1 |
| 0 |

Issue slot: 2   ✖

No eligible warps! Issue slot unused.

Warp Scheduling
Mental model

Active Warp States:

Stalled
Waiting on:
   an instruction fetch,
   a memory dependency,
   an execution dependency, or
   a synchronization barrier.

Eligible
Ready to issue an instruction.

Selected
Eligible that is selected to issue
an instruction

Cycle: N   N+1   N+2

Warp Slots

N
...
5
4
3
2
1
0

Issue slot: 2   ❌
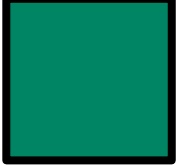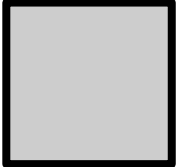
Warp at slot 0 becomes eligible.

Warp Scheduling
Mental model

Active Warp States:

■ **Stalled**
Waiting on:
  an instruction fetch,
  a memory dependency,
  an execution dependency, or
  a synchronization barrier.

□ **Eligible**
Ready to issue an instruction.

■ **Selected**
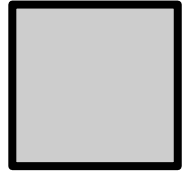Eligible that is selected to issue
an instruction

Cycle: N   N+1   N+2

Warp Slots

N
...
5
4
3
2
1
0

Issue slot: 2   ❌   0
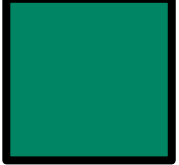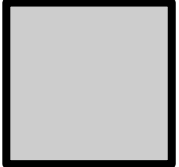
Warp at slot 0 is selected.

Warp Scheduling
Mental model

Active Warp States:

■ **Stalled**
Waiting on:
  an instruction fetch,
  a memory dependency,
  an execution dependency, or
  a synchronization barrier.

□ **Eligible**
Active warp that is not stalled.

■ **Selected**
Eligible warp that is selected to
issue an instruction.

Cycle: N   N+1   N+2   N+3

Warp Slots

| N |
| ... |
| 5 |
| 4 |
| 3 |
| 2 |
| 1 |
| 0 |

Issue slot: 2   ✖   0   ✖

No eligible warps! Issue slot unused.

Warp Scheduling
Mental model

Active Warp States:

■ **Stalled**
Waiting on:
an instruction fetch,
a memory dependency,
an execution dependency, or
a synchronization barrier.

□ **Eligible**
Active warp that is not stalled.

■ **Selected**
Eligible warp that is selected to
issue an instruction.

Cycle: N N+1 N+2 N+3 N+4 N+5 N+6

Warp Slots: N ... 5 4 3 2 1 0
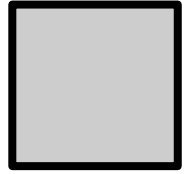
Issue slot: 2 ✖ 0 ✖ ✖ ✖ ✖

Having more **active warps** would help reduce idle issue slots and hide latencies of stalled warps.

How to Increase Active Warps?
Occupancy

- There is a maximum number of warps which can be concurrently active on an SM.
    - **Device** (depends on compute capability of the GPU)
    - **Achievable** (depends on kernel implementation + compiler)
    - **Achieved** (depends mostly on the grid size)

$$Occupancy = \frac{Achievable\ \#\ active\ warps\ per\ SM}{Device\ \#\ active\ warps\ per\ SM}$$

- Occupancy of a CUDA kernel may be limited by:
    - **Register usage**
        - SM registers are partitioned among threads.
    - **Shared memory usage**
        - SM shared memory is partitioned among thread blocks.
    - **Thread block size**
        - Threads are allocated at thread block granularity.

Analyze the occupancy of CUDA kernels with **NVIDIA Nsight Compute**!

NVIDIA

Occupancy Limiters
Registers

- Register usage: compile with `--ptxas-options=-v`
  - Reports registers per thread

- The maximum number of registers per thread can be set manually:
  - At compile time on a per-file basis using the `--maxrregcount` flag of nvcc
  - Per-kernel using the `__launch_bounds__` qualifier

- Hopper has 64K (65536) registers per SM
  - Allocated in fixed-size chunks of 256 registers

- **Example:**
  - Kernel uses 63 registers per thread
  - Registers per warp = 63 * 32 = 2016
  - Registers allocated per warp = 2048
  - Achievable active warps per SM = 65536 / 2048 = 32
  - Occupancy = 32 / 64 * 100 = 50%
    - Hopper supports up to 64 warps per SM

Occupancy Limiters
Shared memory

- Shared memory usage: compile with `--ptxas-options=-v`.
  - Reports static shared memory usage per thread block.

- Hopper has 228 KiB of shared memory.
  - 1KiB per thread block is reserved for system use.
  - With opt-in using dynamic shared memory.

- **Example:**
  - Kernel uses 17408 bytes of shared memory per 128-thread block.
  - Blocks per SM = 233472 / (17408 + 1024) = 12.66
  - Achievable active warps per SM = 12 * 128 / 32 = 48
  - Occupancy = 48 / 64 * 100 = 75%
    - Hopper supports up to 64 warps per SM.

Occupancy Limiters
Thread block size

- Thread block size is a multiple of warp size (32).
  - Even if you request fewer threads, HW rounds up.

- Each thread block can have a maximum size of 1024.

- Each SM can have up to 64 warps, 32 blocks and 2048 threads (Hopper).

Block Size	Active threads per SM	Active Warps per SM	Active Warps per Block	Active Blocks per SM	Occupancy (%)
32	1024	32	1	32	**50**
64	2048	64	2	32	100
256	2048	64	8	8	100
512	2048	64	16	4	100
768	1536	48	24	2	**50**
1024	2048	64	32	2	100

ILP vs TLP for Hiding Latencies
Computing **c = c + a * b**

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
  - Problem size = 2^28
  - Datatype = float
  - Baseline thread block size = 32 (50% occupancy)
  - **Experiment #1**: increase occupancy
    - Thread block size = 64 (100% occupancy)
  - **Experiment #2**: increase ILP by computing more elements per thread
    - Elements per thread = 2, 4

Implementation	Elements per Thread	Thread Block Size	Main Memory Bandwidth Utilization (%)	SM Occupancy (%)	GPU Time (ms)
Baseline	1	32	25	50	5.0
Experiment #1	1	64	51	100	2.5
Experiment #2	2	32	51	50	2.5
Experiment #2	4	32	82	50	**1.6**

What Occupancy Do I Need?

General guidelines

Rule of thumb: Try to maximize occupancy.

But some algorithms will run better at low occupancy.

More registers and shared memory can allow higher data reuse, higher ILP, higher performance.

Low Occupancy   **High Occupancy**

Fewer threads per SM.   **+** Registers per thread and shared memory **-**   More threads per SM.

More resources per thread.   Fewer registers per thread.

Enough instruction-level parallelism
or GPU will starve!   **-** Occupancy **+**   Rely on thread parallelism
to hide latencies!

Complex algorithms ⟵――――――――――――――――――⟶ Simple algorithms

Maximizing Memory Throughput

Memory Hierarchy

NVIDIA H200 SXM

Register File (64K 32-bit registers per SM)

Unified Shared Memory / L1 Cache (228 KiB per SM, variable split)

L2 Cache (50 MiB)

HBM3e (141 GB)

SM

Registers

Shared/L1

L2

DRAM

4.8 TB/s
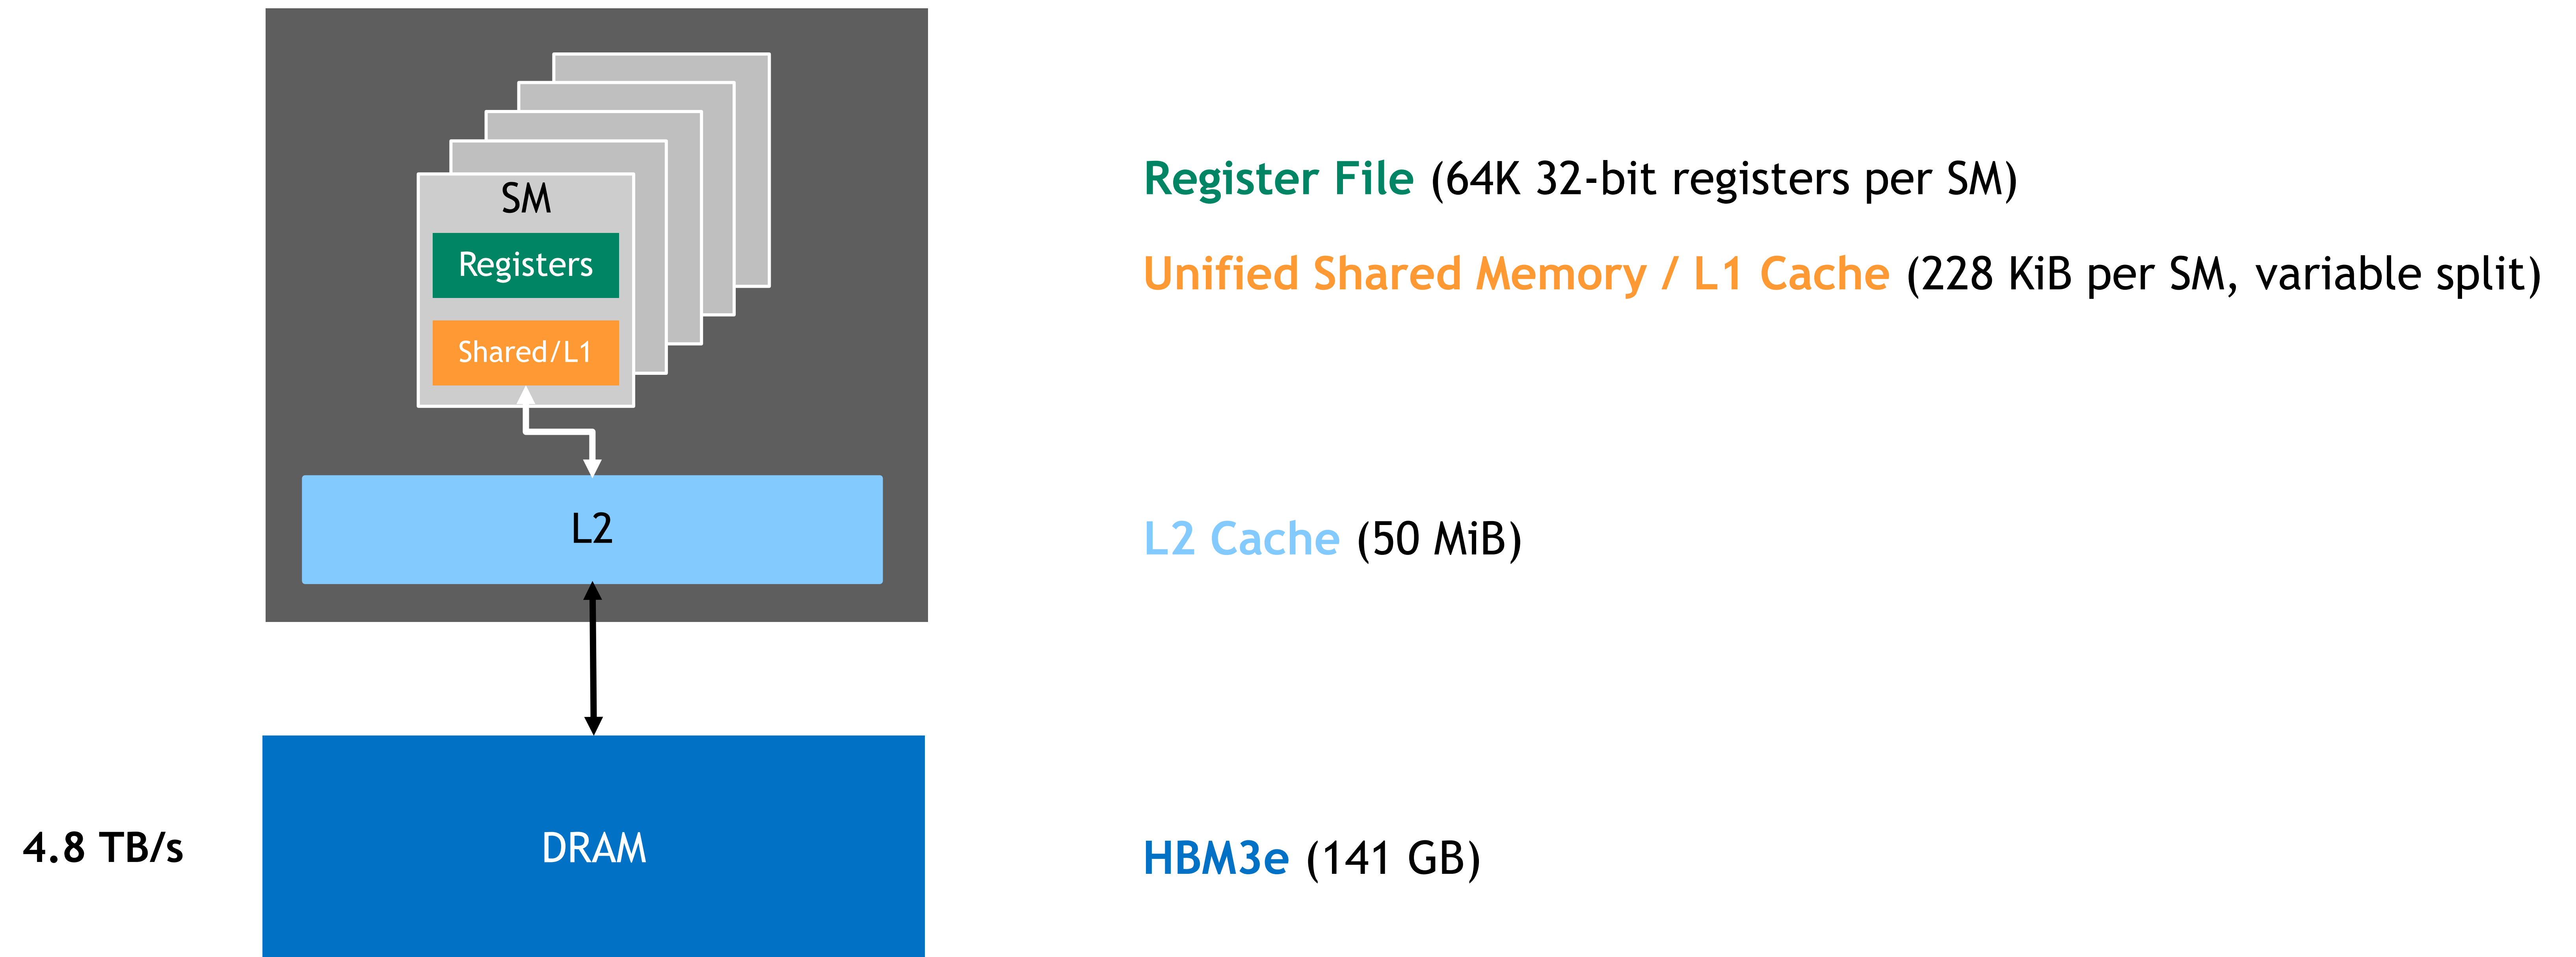
Why Do GPUs Have Caches?

- 100s ~ 1000s of threads sharing the L1 and ~100000s of threads sharing the L2.
  - L1, L2 capacity per thread is relatively small.

Caches on GPUs are mostly useful for:
  - "Smoothing" irregular, misaligned access patterns.
  - Caching common data accessed by multiple threads.
  - Faster register spills, local memory.
  - Faster atomics.

What about cache blocking?
  - L2 cache blocking may be feasible.
    - For an example of efficient use of L2 cache blocking, see **[S62192]: "Advanced Performance Optimization in CUDA"**.

Memory Transactions
Cache lines and sectors

- Minimum memory access granularity: **32 bytes** = **1 sector**
  - **L1 to L2:** 1 sector
  - **L2 to Global:** 2 sectors (default)
    - User can set a preferred granularity with `cudaDeviceSetLimit()` and `cudaLimitMaxL2FetchGranularity`.
    - Only a hint though!

- Cache line size: **128 bytes** = **4 sectors**
  - Cache "management" granularity = 1 cache line
    - Coalescing of requests.
    - Evictions.

128-Byte alignment

| Sector 0 | Sector 1 | Sector 2 | Sector 3 |

128 Byte cache line

NVIDIA.

Memory Reads & Writes

SM

Registers

Shared/L1

L2
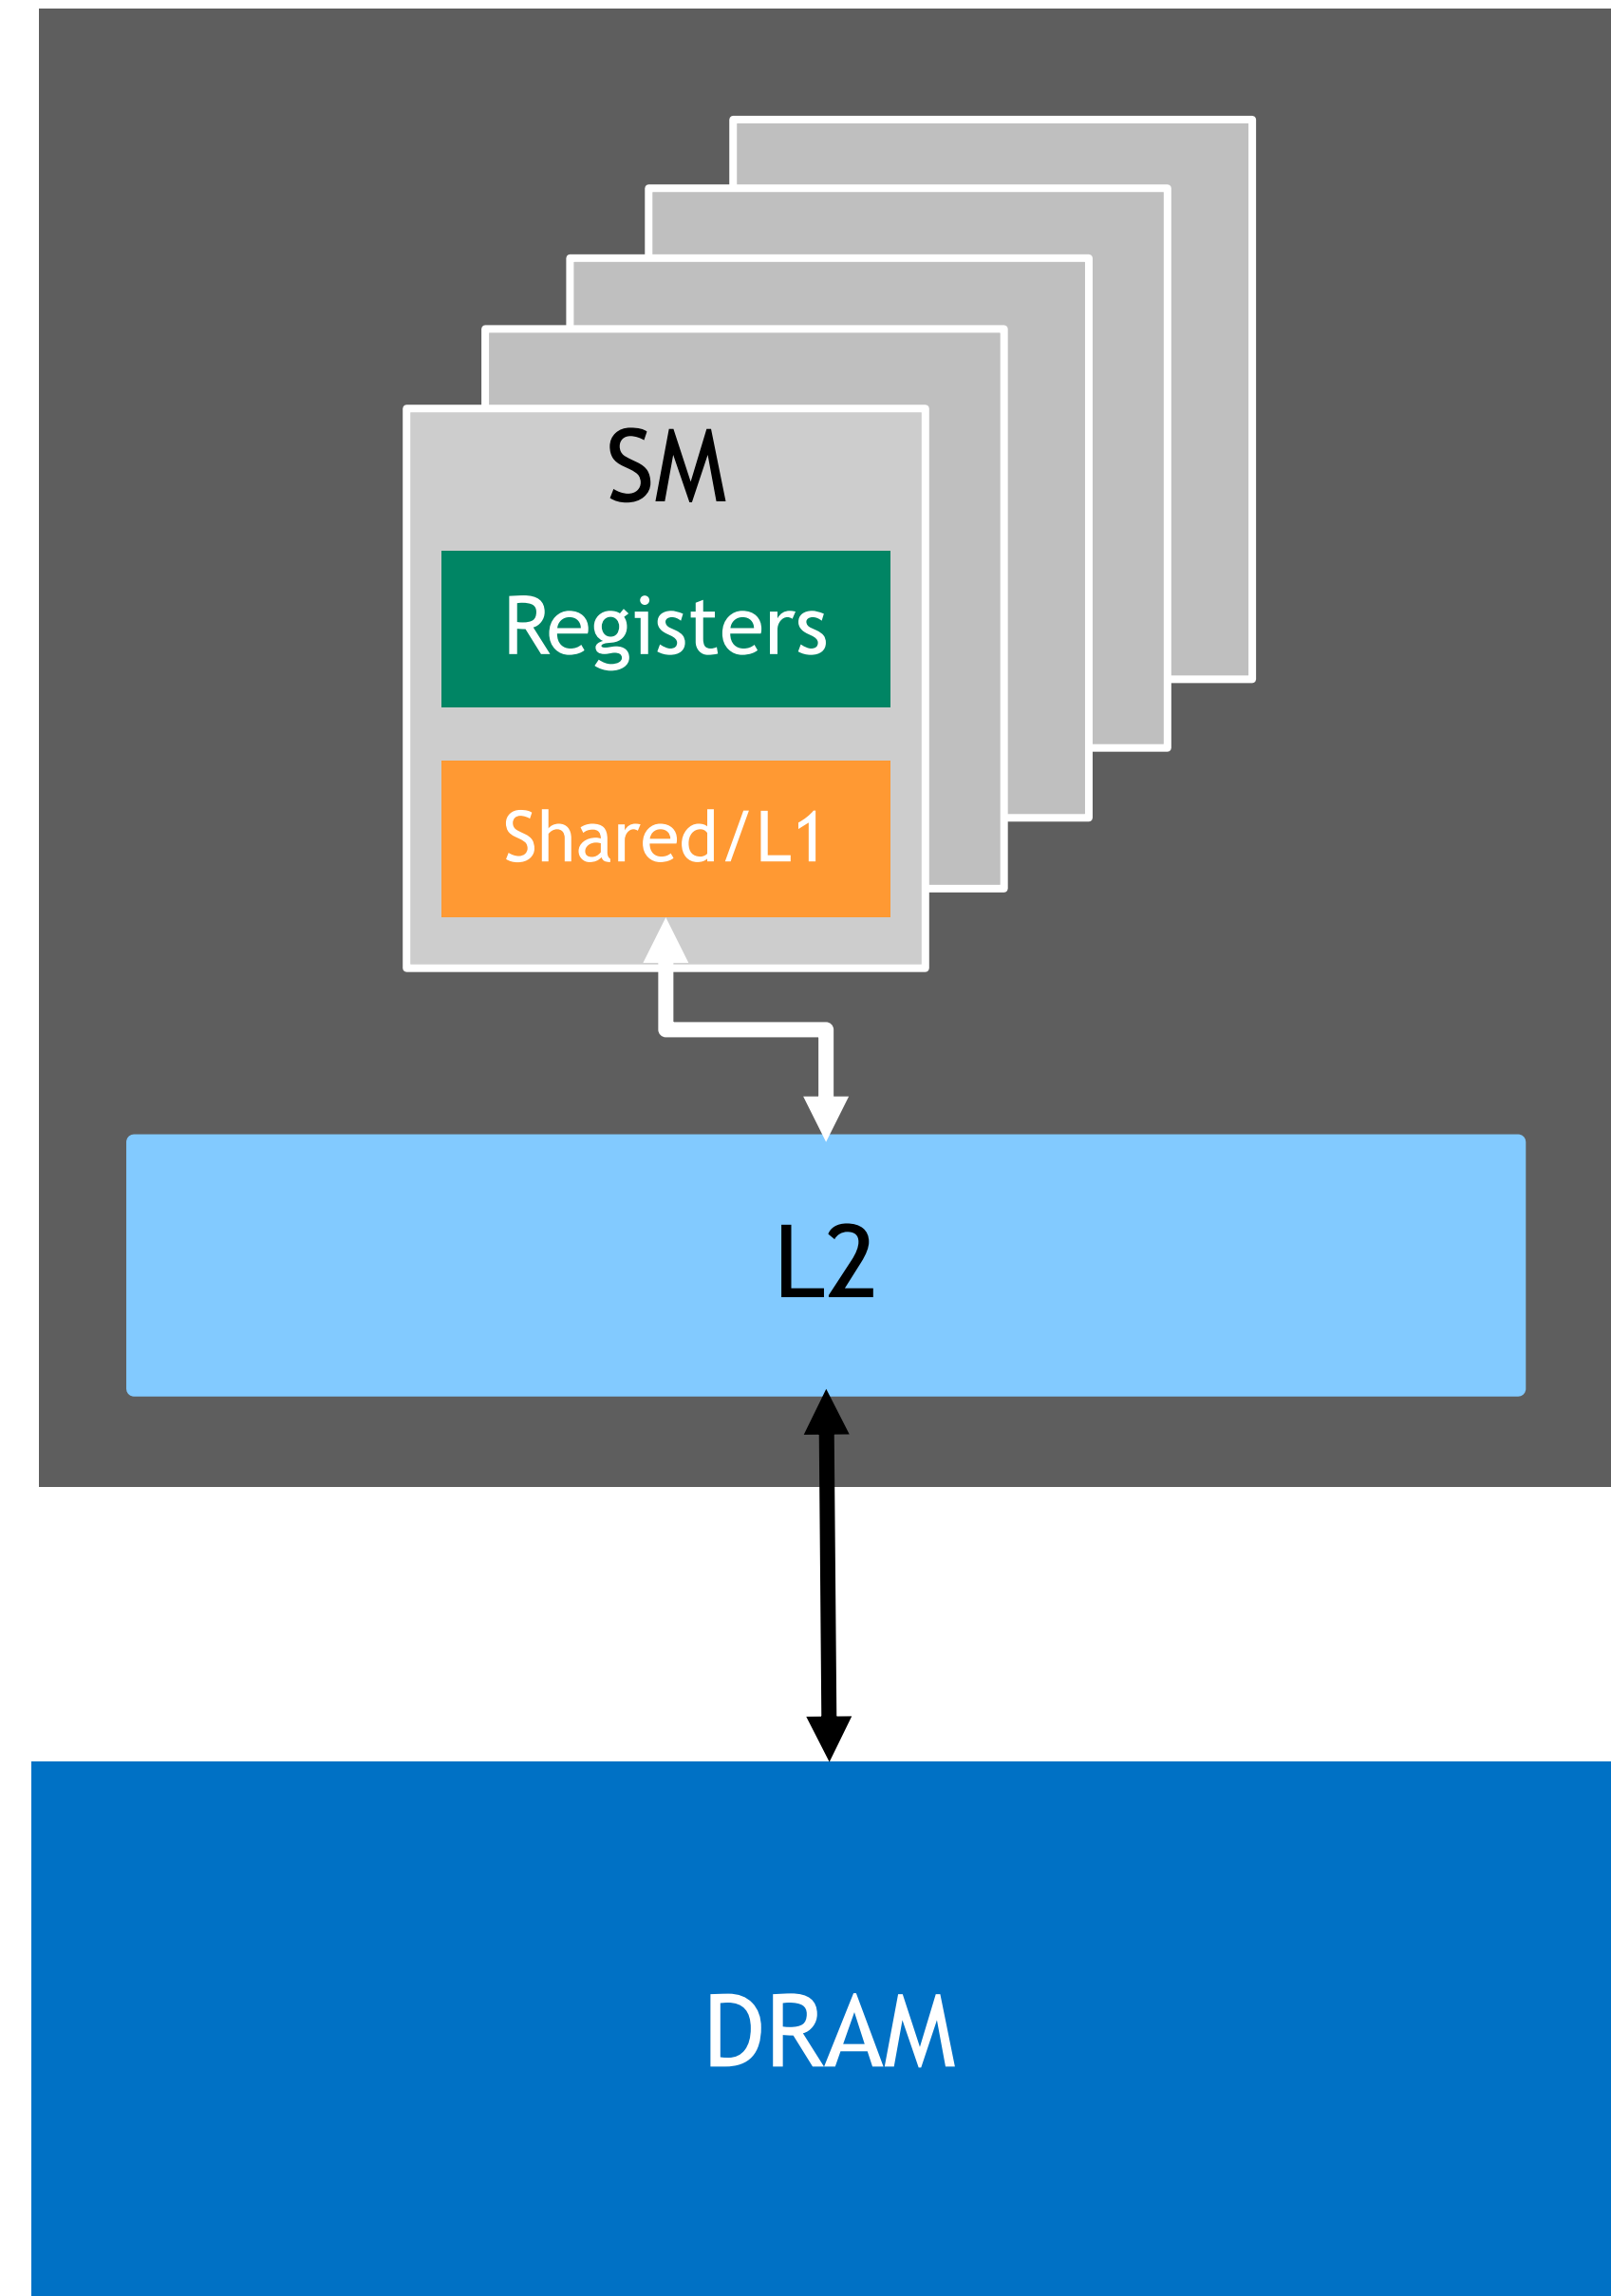
DRAM

Reads

Check if data is in L1 (if not, check L2)

Check if data is in L2 (if not, get from DRAM)

Unit of data moved: **full sector**

Writes

L1 is write-through: update both L1 and L2

L2 is write back: flush data to DRAM only when needed

Unit of data moved: **partial sector***

* Depends on whether ECC is enable/disabled.

Global Memory Access Patterns

Aligned and sequential

4-byte element access
4 sectors

0	WARP	31

| 0 | 32 | 64 | 96 | 128 | 160 | 192 | 224 | 256 | 288 | 320 | 352 |

Memory Addresses

COALESCED!

Global Memory Access Patterns

Aligned and sequential

8-byte element access
8 sectors

0   31

WARP

0   32   64   96   128   160   192   224   256   288   320   352

Memory Addresses

COALESCED!

Global Memory Access Patterns

Aligned and non-sequential

4-byte element access
4 sectors

0   31

| WARP |

0   32   64   96   128   160   192   224   256   288   320   352

Memory Addresses

COALESCED!

Global Memory Access Patterns

Mis-aligned and sequential

4-byte element access
5 sectors

0   31

| WARP |

Memory Addresses

0   32   64   96   128   160   192   224   256   288   320   352

Global Memory Access Patterns
Mis-aligned and sequential

4-byte element access
5 sectors

0   310   31

| WARP | WARP |

0   32   64   96   128   160   192   224   256   288   320   352

Memory Addresses

Global Memory Access Patterns

Same address

4-byte element access
1 sector

```
0                          31
┌──────────────────────────┐
│          WARP            │
└──────────────────────────┘
```

Memory Addresses

Global Memory Access Patterns

Aligned and strided

4-byte element access
32 sectors

0   31

| WARP |

0   32   64   96   128   160   192   224   256   288   320   352

Memory Addresses
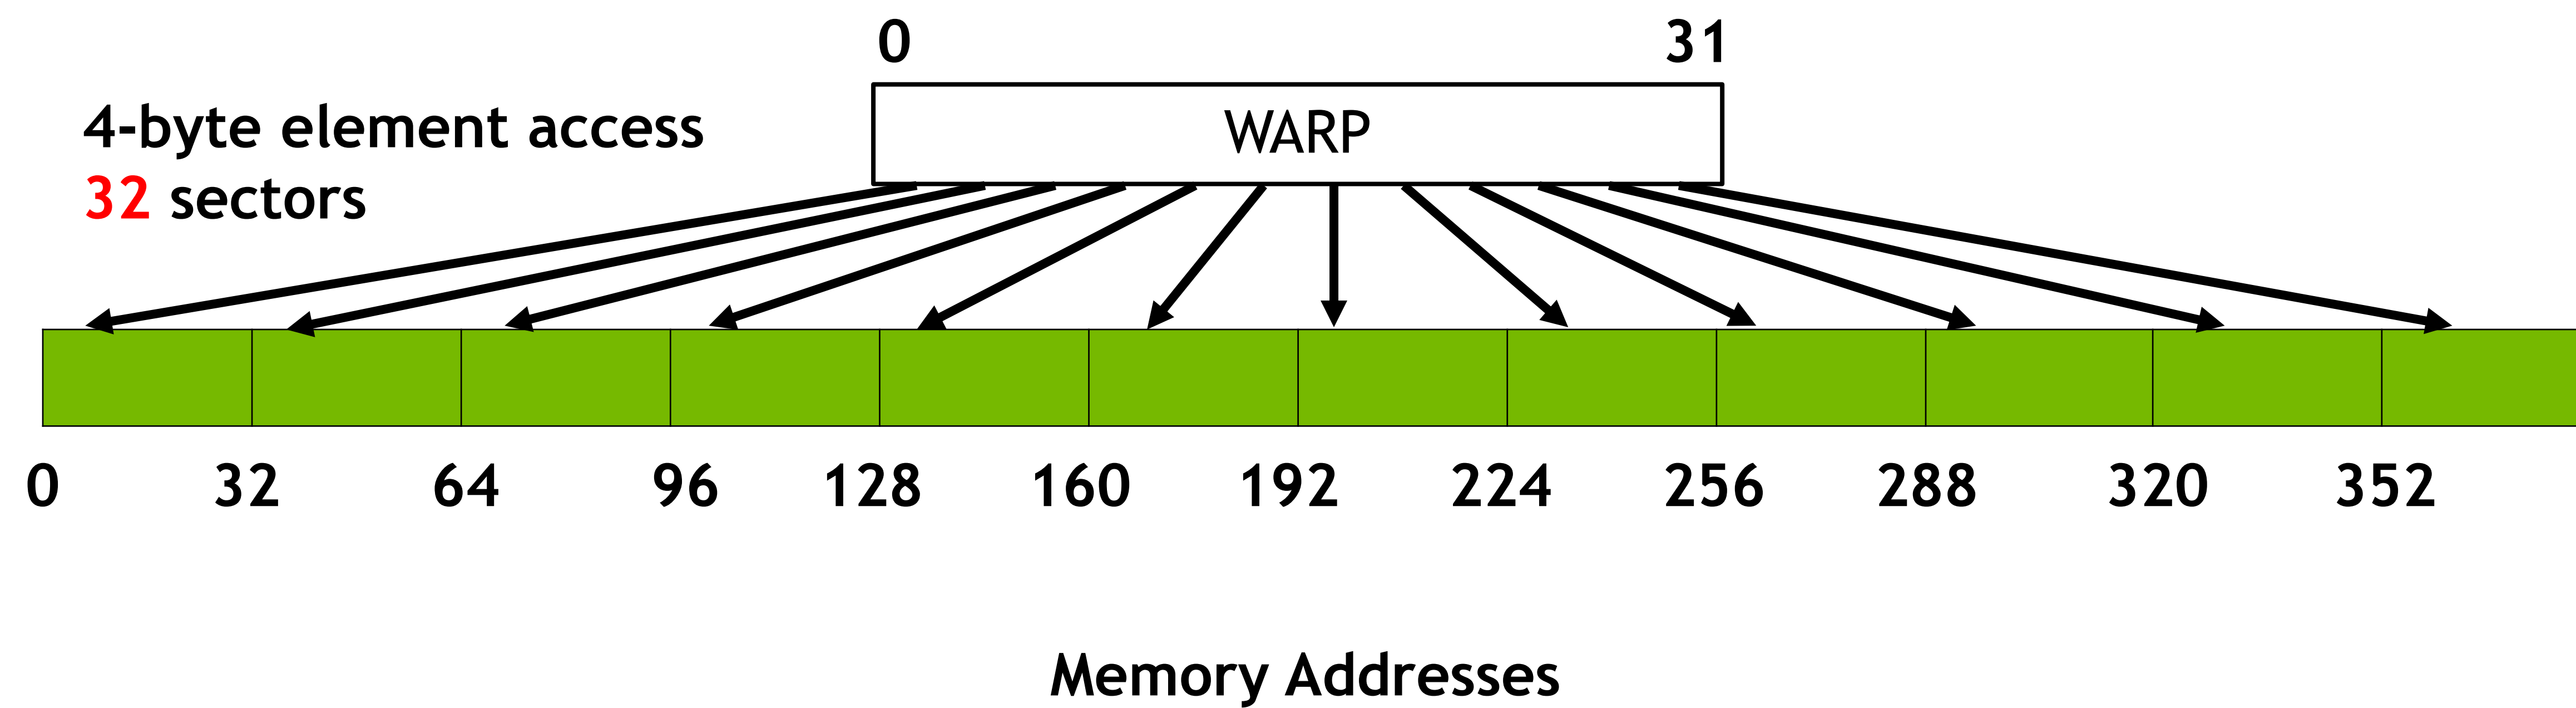
Impact of Data Layout

Array-of-Structures (AoS) vs Structure-of-Arrays (SoA)

AoS Memory Layout

| u | v | w | x0 | ... | x7 | y0 | ... | y7 | z |

0   80

```
struct Coefficients
{
  float u, v, w;
  float x[8], y[8], z;
};

__global__ void kernel(Coefficients *data)
{
  int i = cg::this_grid.thread_rank();

  data[i].u = data[i].u + 10.f;
  data[i].y[0] = data[i].y[0] + 10.f;
}
```

Impact of Data Layout

Array-of-Structures (AoS) vs Structure-of-Arrays (SoA)

- When loading coefficients u and y[0]:
  - Successive threads in a warp read 4 bytes at 80-byte stride.

T0　　　　**T1**　　　　**T2**　　　　**T3**

| u | ... | y0 | ... | z | u | ... | y0 | ... | z | u | ... | y0 | ... | z | u | ... | y0 | ... | z | ... |

0　　　　　　80　　　　　160　　　　　240

```
struct Coefficients
{
  float u, v, w;
  float x[8], y[8], z;
};

__global__ void kernel(Coefficients *data)
{
  int i = cg::this_grid.thread_rank();

  data[i].u = data[i].u + 10.f;
  data[i].y[0] = data[i].y[0] + 10.f;
}
```

Impact of Data Layout

Array-of-Structures (AoS) vs Structure-of-Arrays (SoA)

- When loading coefficients u and y[0]:
  - Successive threads in a warp read 4 bytes at 80-byte stride.
- We are reading **7x more bytes** than necessary!
  - Remember data is read in sectors of 32 bytes.
  - No potential reuse of the sectors loaded by the previous access.

```
T0              T1              T2              T3
↓               ↓               ↓               ↓

u  … y0 … z u  … y0 … z u  … y0 … z u  … y0 … z  …
44              124             204             284
```

```cpp
struct Coefficients
{
  float u, v, w;
  float x[8], y[8], z;
};

__global__ void kernel(Coefficients *data)
{
  int i = cg::this_grid.thread_rank();

  data[i].u = data[i].u + 10.f;
  data[i].y[0] = data[i].y[0] + 10.f;
}
```
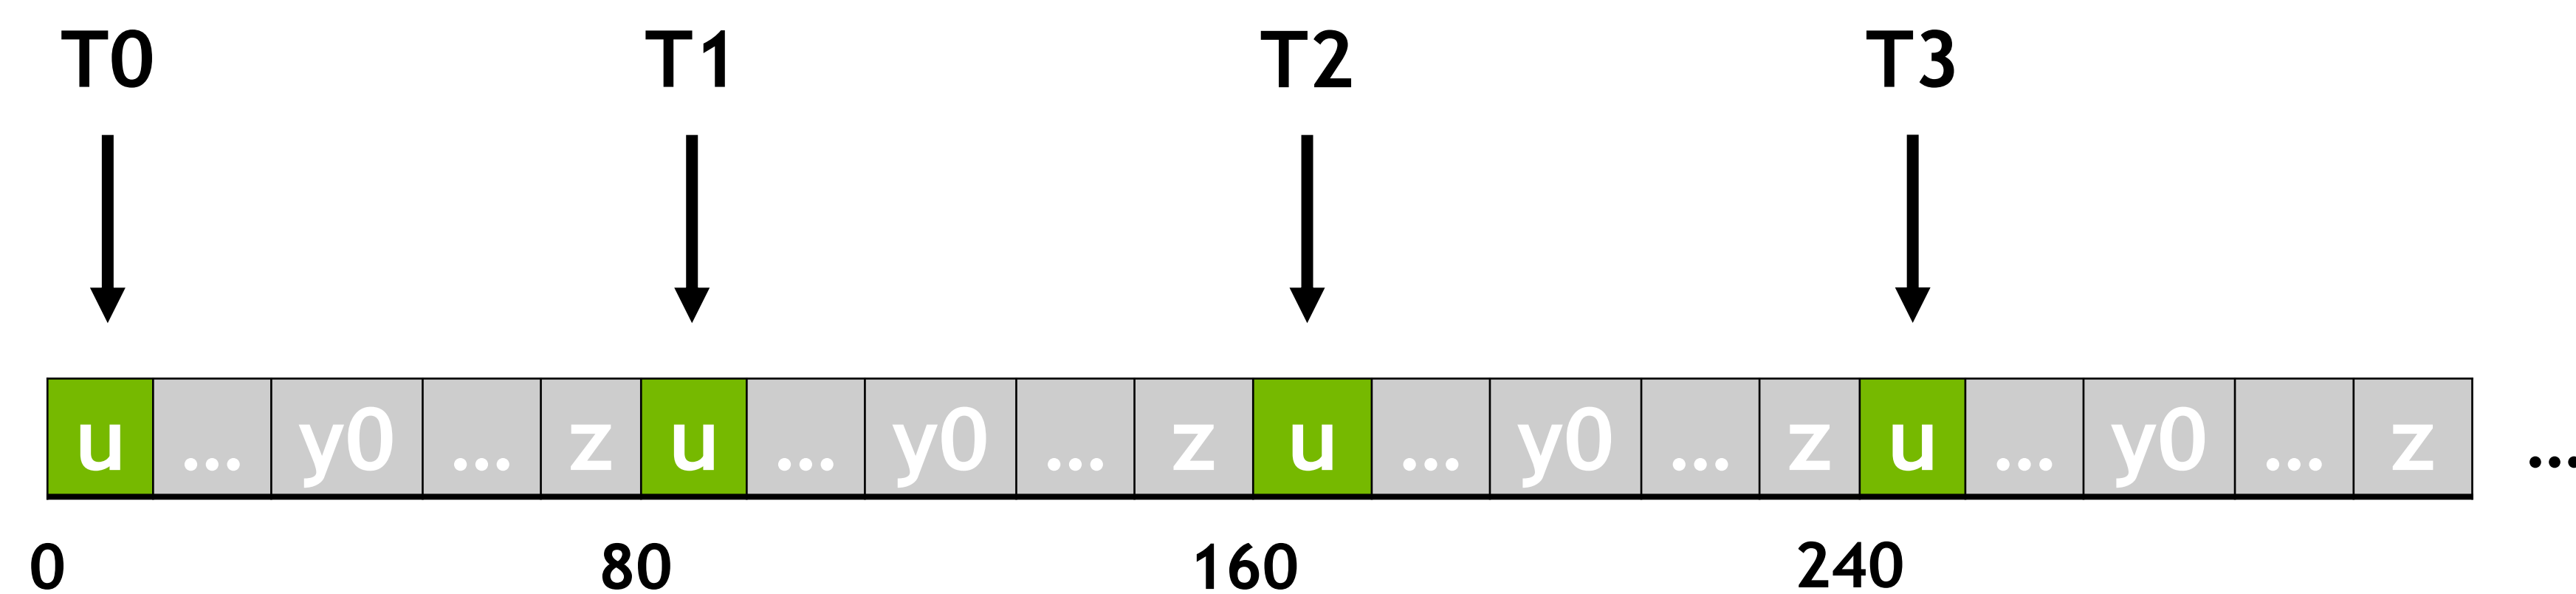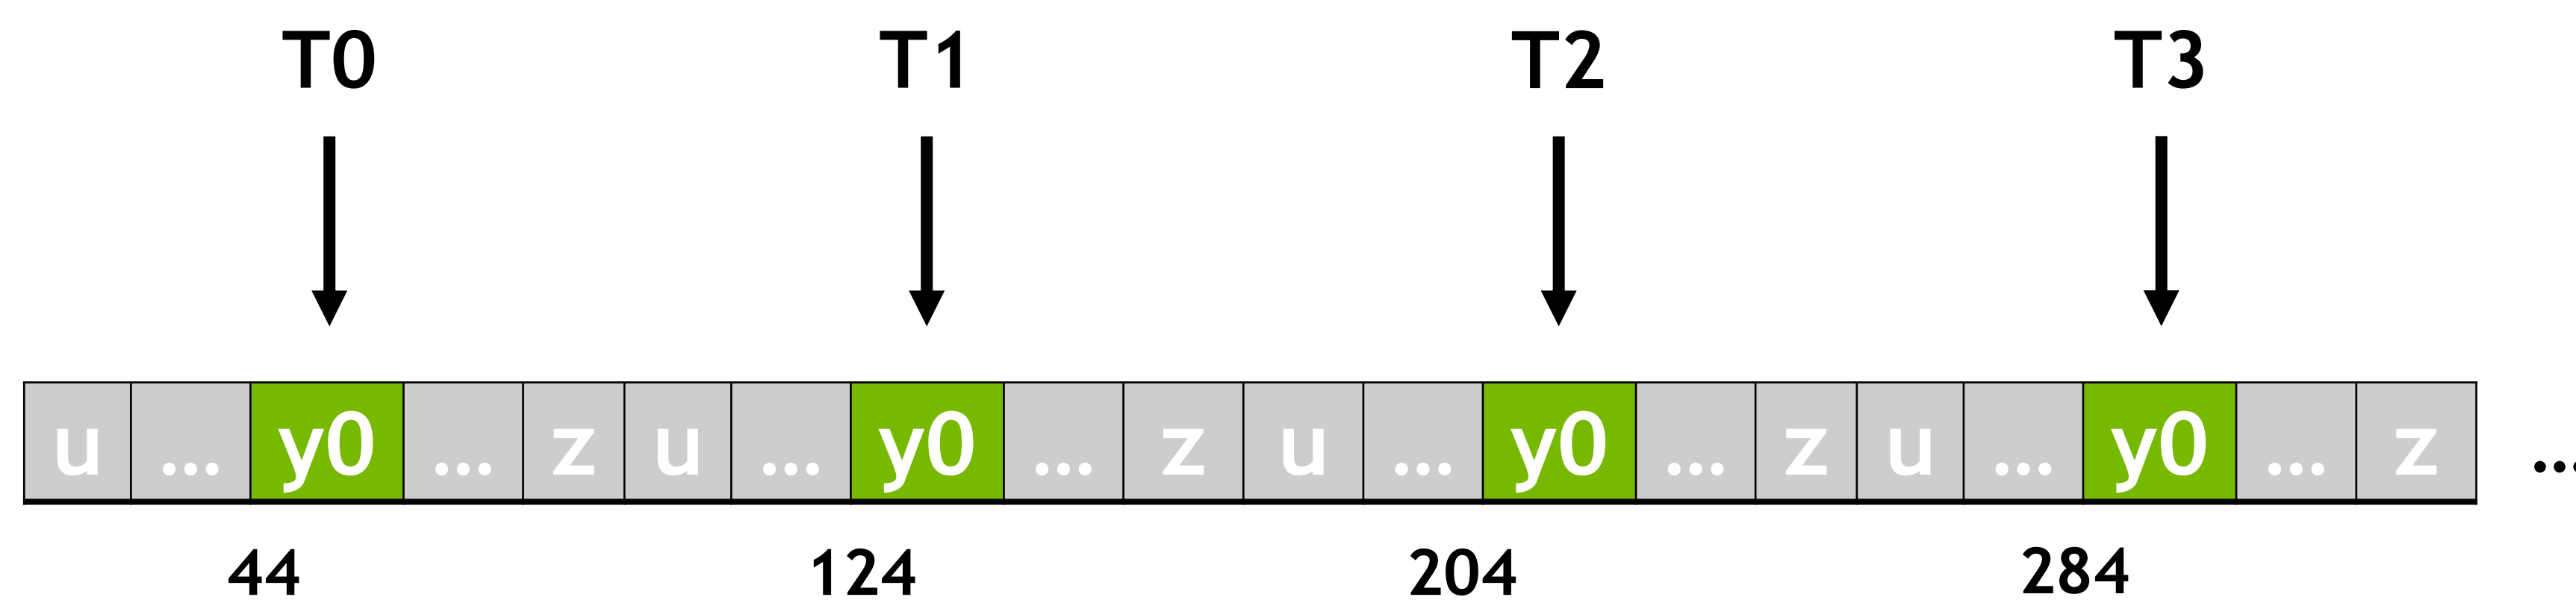
Impact of Data Layout

Array-of-Structures vs Structure-of-Arrays

- Refactoring from **AoS** to **SoA** leads to coalesced memory accesses for u and y[0].

SoA Memory Layout

T0 T1 T2 T3 T4 T5

128 bytes

```
struct Coefficients
{
  float *u, *v, *w;
  float *x0, …, *x7, *y0, … *y7, *z;
};

__global__ void kernel(Coefficients data)
{
  int i = cg::this_grid.thread_rank();

  data.u[i] = data.u[i] + 10.f;
  data.y0[i] = data.y0[i] + 10.f;
}
```

Impact of Data Layout
Array-of-Structures vs Structure-of-Arrays

- Refactoring from **AoS** to **SoA** leads to coalesced memory accesses for u and y[0].

SoA Memory Layout

T0 T1 T2 T3 T4 T5

128 bytes

```
struct Coefficients
{
  float *u, *v, *w;
  float *x0, …, *x7, *y0, … *y7, *z;
};

__global__ void kernel(Coefficients data)
{
  int i = cg::this_grid.thread_rank();

  data.u[i] = data.u[i] + 10.f;
  data.y0[i] = data.y0[i] + 10.f;
}
```

Impact of Data Layout
Performance Analysis

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
  - Problem size = 2^28
  - Thread block size = 256

Implementation	Load Efficiency (%)	Store Efficiency (%)	Main Memory Bandwidth Utilization (%)	GPU Time (ms)
AoS	12.5	12.5	13.50	28.497
SoA	100	100	79.47	**4.836**

Unified L1 and Shared Memory

- Can be used as a typical hardware managed cache (L1) and/or a user-managed memory (Shared Memory)
  - An application can configure its preferred split at runtime using `cudaFuncSetAttribute()` with the attribute `cudaFuncAttributePreferredSharedMemoryCarveout`.

- **Shared memory** can be useful for:
  - Storing frequently used data
  - Improving global memory access patterns
  - Data layout conversion
  - Communication among threads of a thread block

Shared Memory

Capacity:

- Default 48 KiB per thread block, opt-in to get more using `cudaFuncSetAttribute()` with the attribute `cudaFuncAttributeMaxDynamicSharedMemorySize`.
    - Up to 227KiB per thread block on Hopper.

Organization:

- Divided into 32 banks, each 4-byte wide.
- Successive **4-byte** words map to successive banks.
- Bank index calculation examples:
    - (4-byte word index) % 32
    - (1-byte word index / 4) % 32

Performance:

- Slower than registers, but much faster than global memory.

Logical View of Shared Memory Banks

4-byte data

With 4-Bytes data

Byte-address:

0	4	8	12	16	20	24	28	32	38	40	44		120	124	128
0	1	2	3	4	5	6	7	8	9	10			30	31	

128	132	136	140	144	148		248	252	256
32	33						•••		

256	260	264			384

•••

Bank-0 Bank-1   Bank-31
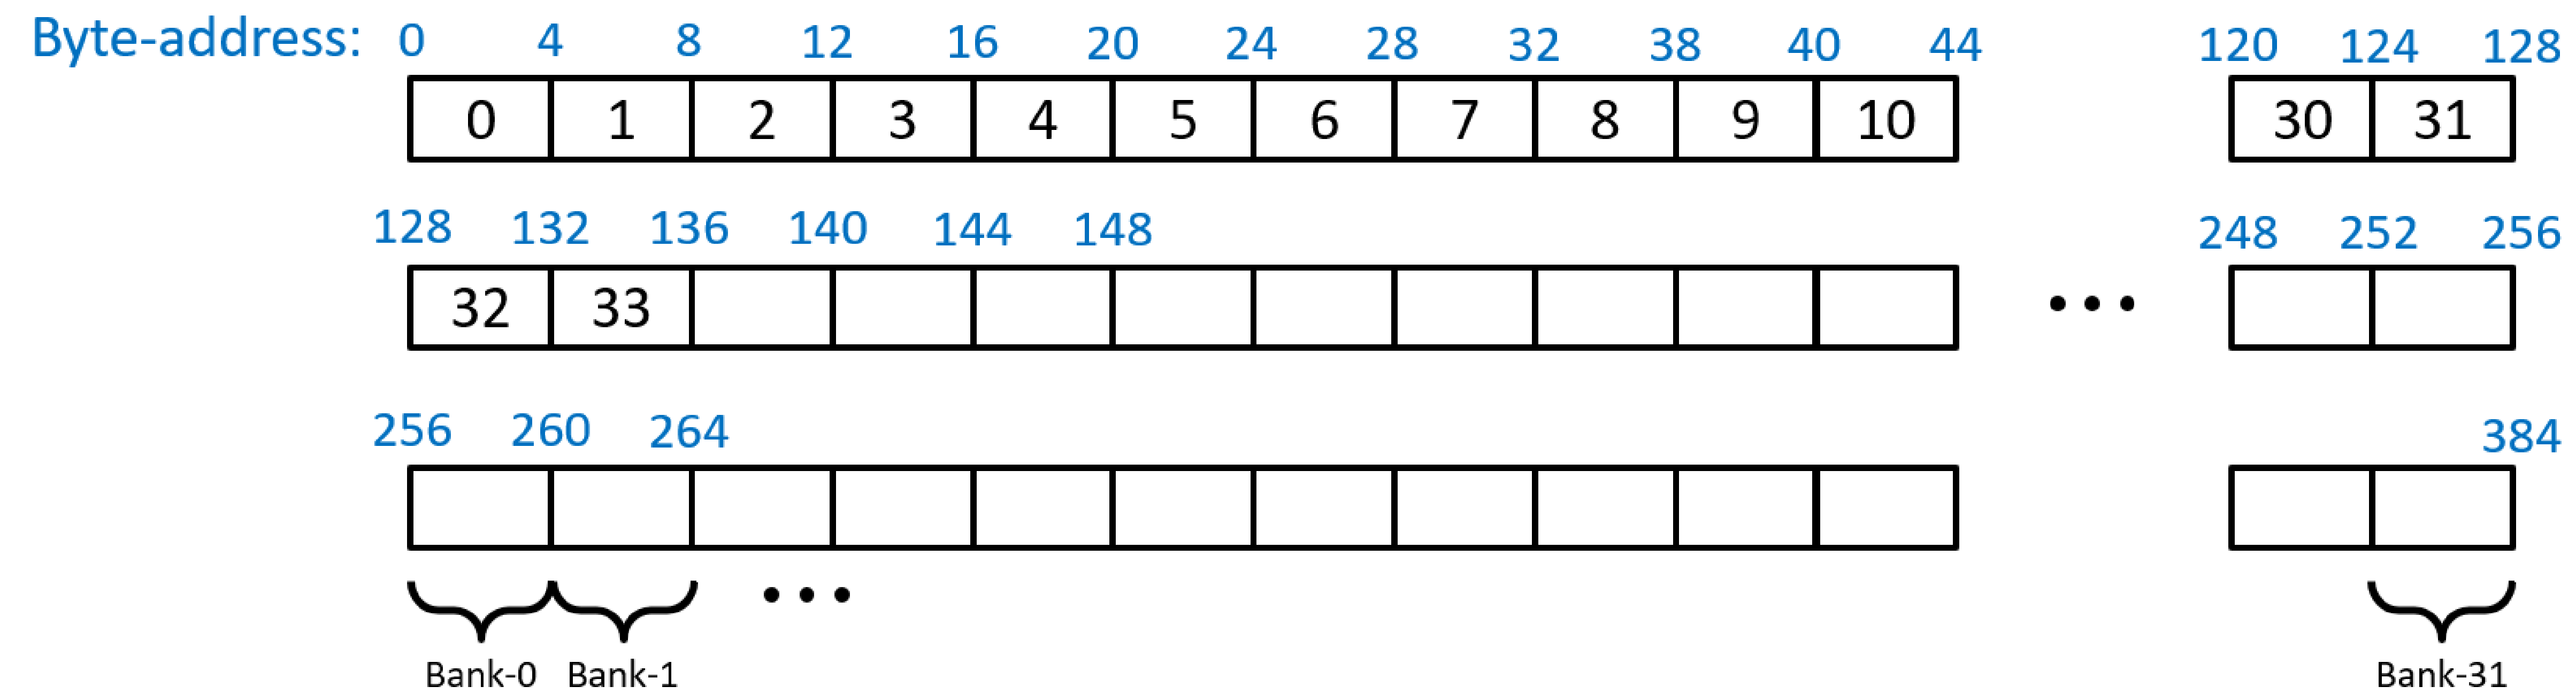
Processing Data Types of Different Sizes

- **4-byte** or smaller data types:
  - Process addresses of all threads in a warp in a single phase

- **8-byte** data types:
  - Process addresses of all threads in a warp in **2** phases
    - Each phase processes addresses of **half** of a warp

- **16-byte** data types:
  - Process addresses of all threads in a warp in **4** phases
    - Each phase processes addresses of a **quarter** of a warp

nVIDIA.

Shared Memory Access Patterns

Bank conflicts

- Bank conflicts occur when threads in the **same phase** want to access the **same bank**.

Coalesced access

(No bank conflicts)

```
shmem[threadIdx.x] = data[tid]
```

4-byte data

Threads

0   31

Shared Memory Banks

Conflict access

(2-way bank conflicts)

```
shmem[threadIdx.x * 2] = data[tid]
```

4-byte data

Threads

0   31

Shared Memory Banks

Broadcast access

(No bank conflicts)

```
data = shmem[0]
```

4-byte data

Threads

0   31

Shared Memory Banks

Bank Conflicts

Example

- 32 x 32 array of floats in shared memory
  - 4-byte data, 1 array element per bank
  - Row-major layout
  - 2D thread block

- **Access pattern:**
  - `idx := threadIdx.x*32 + threadIdx.y`
  - 32-way bank conflicts

Thread 0	(0,0)	(0,1)	(0,2)	(0,3)		(0,31)
Thread 1	(1,0)	(1,1)	(1,2)	(1,3)		(1,31)
Thread 2	(2,0)	(2,1)	(2,2)	(2,3)	⋮	(2,31)
Thread 3	(3,0)	(3,1)	(3,2)	(3,3)		(3,31)
			…		…	…
Thread 31	(31,0)	(31,1)	(31,2)	(31,3)		(31,31)

Bank 0
Bank 1
Bank 2
Bank 3
...
Bank 31

All threads in a warp access the same bank!

Resolving Bank Conflicts

Padding

- 32 x **33** array of floats in shared memory
  - 4-byte data, 1 array element per bank
  - Row-major layout
  - 2D thread block

- **Access pattern:**
  - `idx := threadIdx.x*`**`33`**` + threadIdx.y`
  - No conflicts!


Each thread in a warp accesses a **distinct** bank!

Resolving Bank Conflicts

Swizzling

- 32 x 32 array of floats in shared memory
  - 4-byte data, 1 array element per bank
  - Row-major layout
  - 2D thread block

- **Access pattern:**
  - `idx = threadIdx.x*32 +`
    `threadIdx.y ^ threadIdx.x`
  - No conflicts!
  - No shared memory wasted!

Thread 0	(0,0)	(0,1)	(0,2)	(0,3)	(0,31)
Thread 1	(1,0)	(1,1)	(1,2)	(1,3)	(1,31)
Thread 2	(2,0)	(2,1)	(2,2)	(2,3)	(2,31)
Thread 3	(3,0)	(3,1)	(3,2)	(3,3)	(3,31)
	...			...	...
Thread 31	(31,0)	(31,1)	(31,2)	(31,3)	(31,31)

Bank 0
Bank 1
Bank 2
Bank 3
...
Bank 31

Each thread in a warp accesses a distinct bank!

Vectorized Memory Accesses

Multi-word as well as multi-thread

Memory

contiguous, aligned memory access

int

Threads 0-31

cache line 0

warp
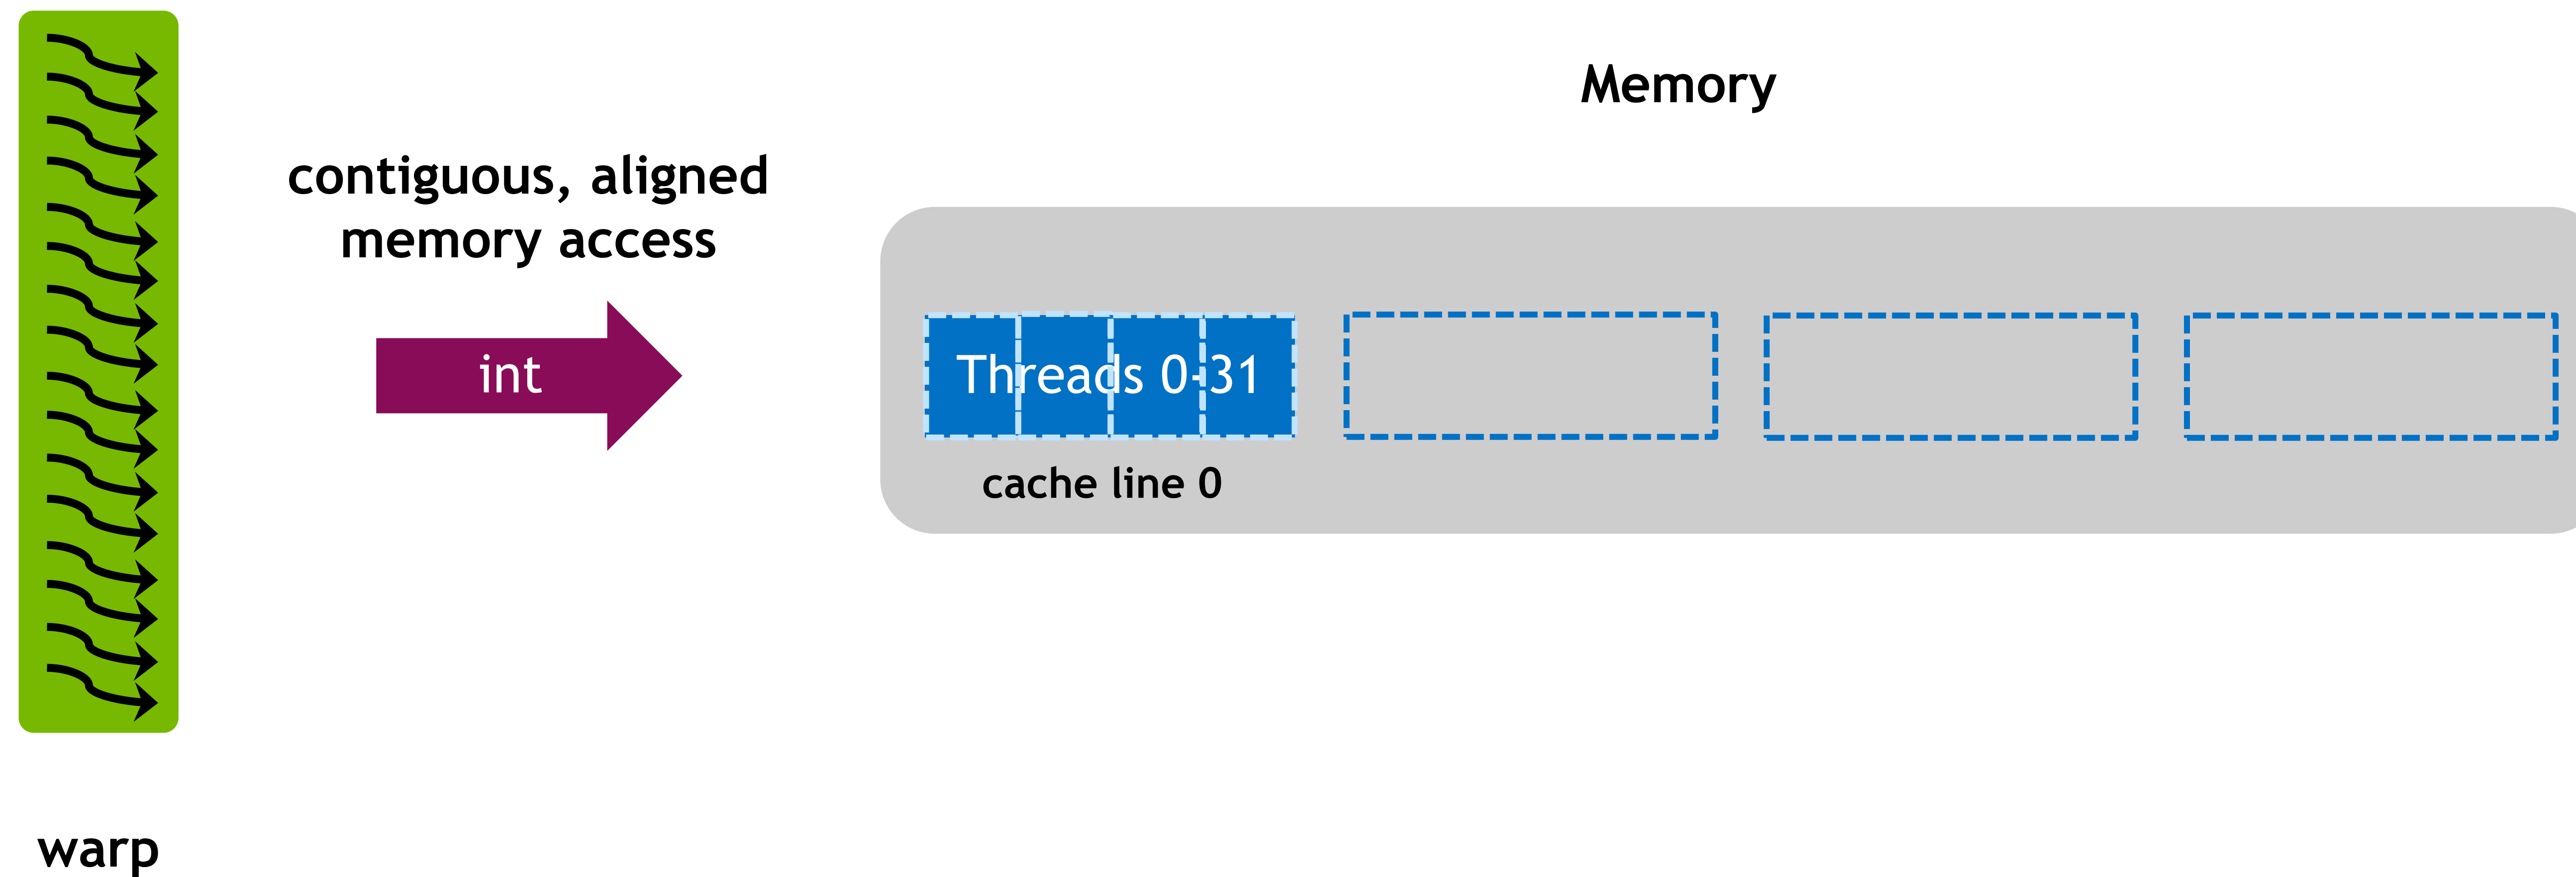
Fills 1 cache line in a single fetch.

Vectorized Memory Accesses

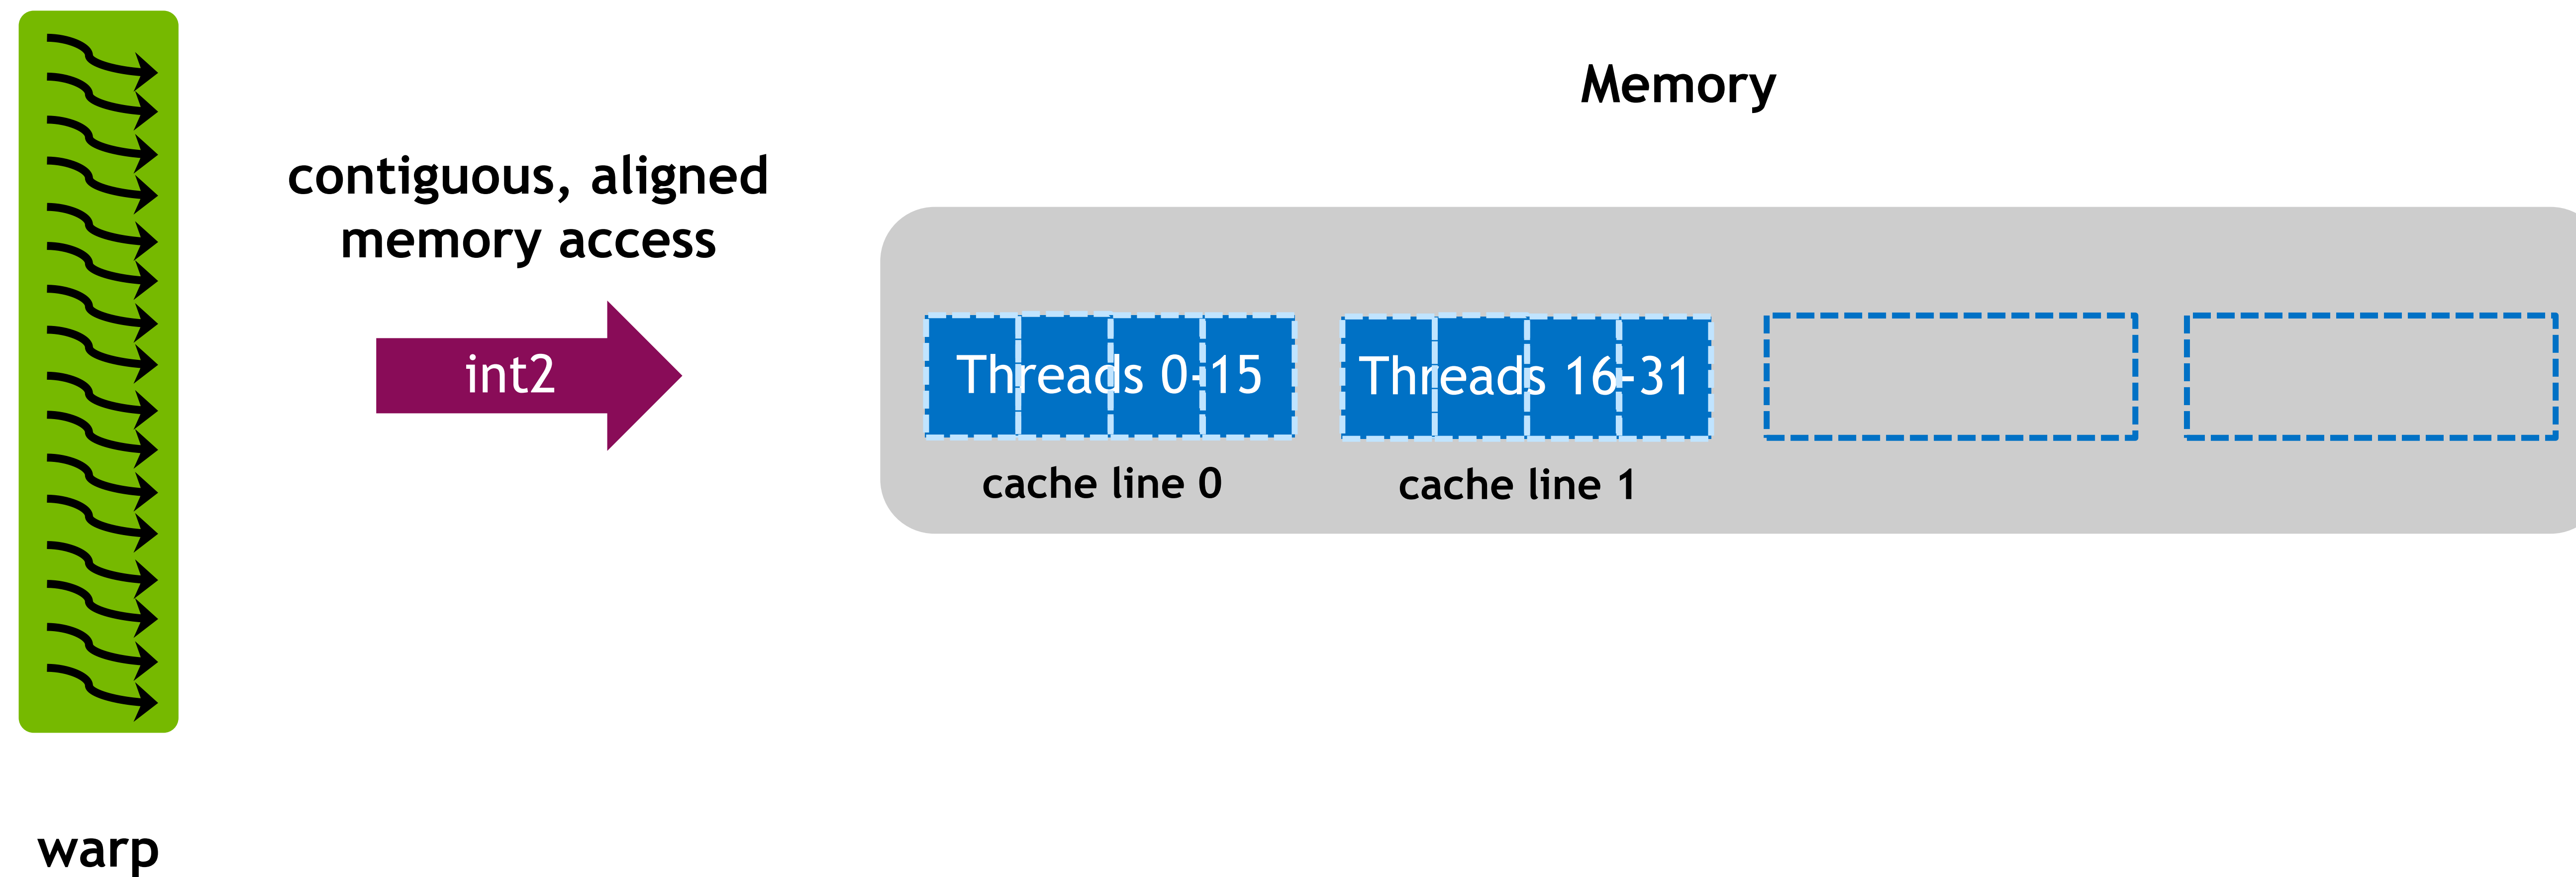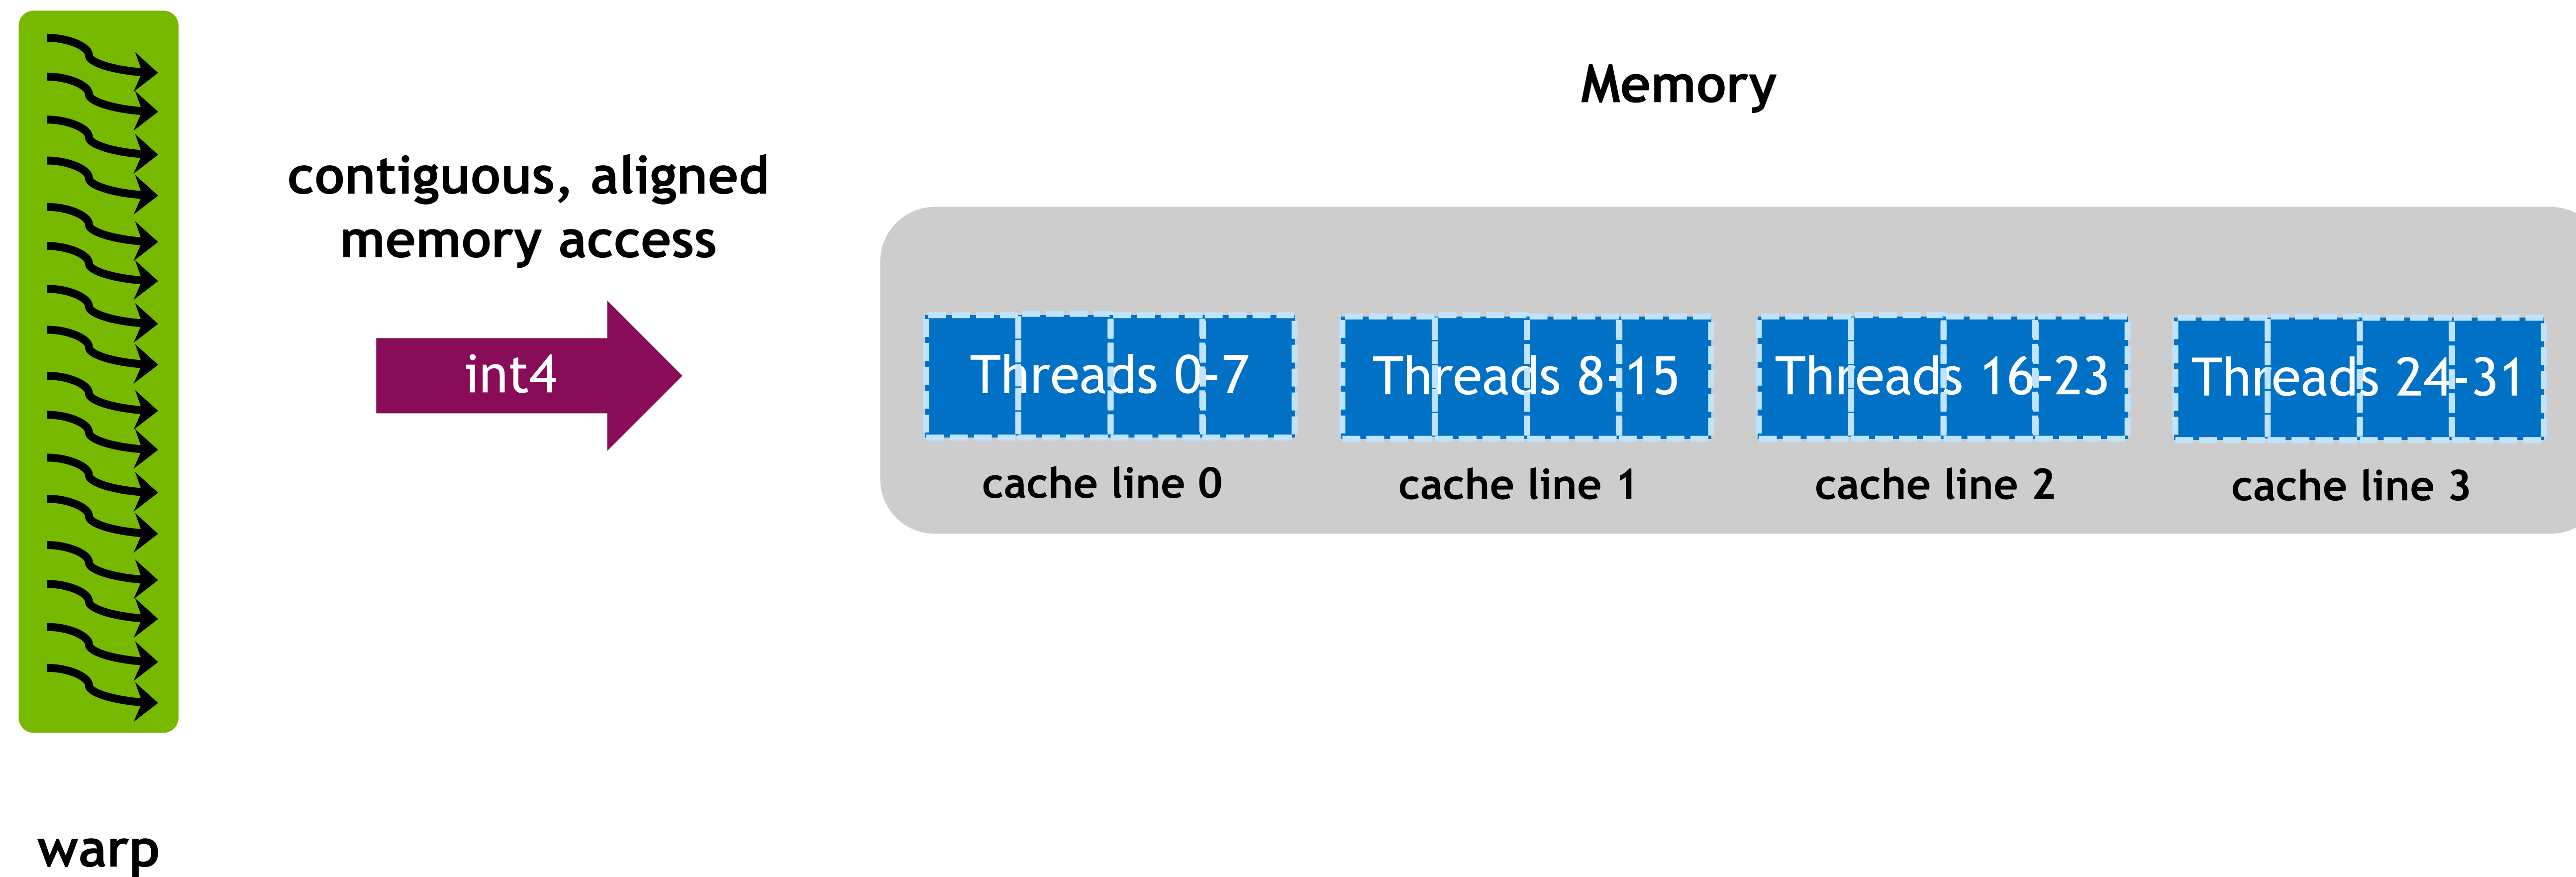## Multi-word as well as multi-thread

Fills 2 cache lines in a single fetch.

Vectorized Memory Accesses

Multi-word as well as multi-thread

contiguous, aligned
memory access

int4

Memory

Threads 0-7	Threads 8-15	Threads 16-23	Threads 24-31
cache line 0	cache line 1	cache line 2	cache line 3

warp

Fills 4 cache lines in a single fetch.

Vectorized Memory Accesses

Multi-thread, multi-word

- Vectorized **global** and **shared memory** accesses.
  - Require aligned data.
  - 64- or 128-bit width.

- Less executed instructions!

- More bytes in-flight!

- Approaches to enable vectorization:
  1) By using vector data types, e.g., `float2`, `float4`.
  2) Explicitly by casting to vector pointers.
     1) Proper alignment required.

```cpp
// Using vectors data types
__global__
void copy(const float2 * __restrict__ in,
          float2 * __restrict__ out,
          int N)
{
  auto grid = cg::this_grid();
  int tid = grid.thread_rank();
  int stride = grid.size();

  for (int i = tid; i < N / 2; i += stride) {
    out[i] = in[i];
    // Same as:
    // out[i].x = in[i].x;
    // out[i].y = in[i].y;
  }
}
```

Vectorized Memory Accesses
Performance Analysis

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
  - Problem size = 2^28
  - Thread block size = 256

Implementation	Main Memory Bandwidth Utilization (%)	GPU Time (ms)
float	60.62	1.033
float2	84.34	0.737
float4	88.82	**0.706**

Maximizing Memory Throughput
General guidelines

<u>Global memory</u>
- Strive for aligned and coalesced accesses within a warp.
- Maximize bytes in-flight to saturate memory bandwidth.
  - Process several elements per thread.
  - Use vectorized loads/stores.
  - Launch enough threads to maximize throughput.

<u>L1 and L2 caches</u>
- Cache blocking difficult, but not impossible.
- Rely on caches when you don't have a choice.

<u>Shared memory</u>
- Use it to reduce global memory traffic.
- Strive to avoid bank conflicts.
- Use vectorized loads/stores.

Atomics

Using Atomics Efficiently

Access Patterns

Same address

Serialized!
Least efficient access pattern.

Scattered

Coalesced

Most efficient access pattern.

Using Atomics Efficiently

Example #1: find the maximum value of an array

- **Problem description:** given an input array, find the maximum element in the array.

- **Naïve implementation:** every thread find its local maximum and then atomically updates the global maximum.
  - N / elements_per_thread same-address global atomics.

```
__global__
void find_max(const int * __restrict__ in, int *max, int N)
{
  int grid_tid = cg::this_grid().thread_rank();
  int grid_stride = cg::this_grid().num_threads();

  // Find my local maximum
  int local_max = INT_MIN;
  for (int i = grid_tid; i < N; i += grid_stride) {
    if (in[i] > local_max)
      local_max = in[i];
  }

  // Atomically update the global max
  atomicMax(max, local_max);
}
```

Using Atomics Efficiently

Example #1: find the maximum value of an array

- **Optimization #1**: maintain a block-level max in shared memory.
  - Reduces the number of same-address global atomics by a factor equal to the thread block size.

```cpp
__global__
void find_max(const int * __restrict__ in, int *max, int N)
{
  int grid_tid = cg::this_grid().thread_rank();
  int grid_stride = cg::this_grid().num_threads();
  auto block = cg::this_thread_block();
  int block_tid = block.thread_rank();

  __shared__ int block_max;

  // Find my local maximum
  int local_max = INT_MIN;
  for (int i = grid_tid; i < N; i += grid_stride) {
    if (in[i] > local_max)
      local_max = in[i];
  }

  // Atomically update the block-level max
  atomicMax(&block_max, local_max);
  block.sync();

  // Atomically update the global max
  if (block_tid == 0)
    atomicMax(max, block_max);
}
```

Using Atomics Efficiently

Example #1: find the maximum value of an array

- **Optimization #2**: use a parallel reduction to calculate the block-level max in shared memory.

```cpp
//Assumes a block dimension of 256
__global__
void find_max(const int * __restrict__ in, int *max, int N) {

  …
  auto tile = cg::tiled_partition<32>(block);
  extern __shared__ int sdata[];

  // Find my local maximum as before
  // Each thread puts its local max into shared memory
  sdata[block_tid] = thread_max;

  // Block-level reduction
  if (block_tid < 128) {
    if (sdata[block_tid + 128] > thread_max)
      thread_max = sdata[block_tid + 128];
    sdata[block_tid] = thread_max;
  }
  block.sync();
  if (block_tid < 64) {
    if (sdata[block_tid + 64] > thread_max)
      thread_max = sdata[block_tid + 64];
    sdata[block_tid] = thread_max;
  }
  block.sync();

  // Warp-level reduction
  if (tile.meta_group_rank() == 0) {
    thread_max = cg::reduce(tile, thread_max,
                            cg::greater<int>());

  }

  if (block_tid == 0)
    atomicMax(max, thread_max);
}
```

Performance Analysis

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
  - Problem size = 2^28
  - Uniform distribution (-50, 50).

Implementation	Thread Block Size	GPU Time (ms)
global atomics	256	6.839
shared memory atomics	256	1.334
shared memory reduction	256	**1.066**

Using Atomics Efficiently

Example #2: vector update

- **Problem description:** a = a + b * c

```
__global__
void vector(const float * __restrict__ b,
            const float * __restrict__ c,
            float * __restrict__ a,
            int N) {
  int grid_tid = cg::this_grid().thread_rank();
  int grid_stride = cg::this_grid().num_threads();

  for (int i = grid_tid; i < N; i += grid_stride) {
    a[i] += b[i] * c[i];
  }
}
```

Memory operations = 3 reads + 1 write

Using Atomics Efficiently

Example #2: vector update

- **Optimization**: use atomics to update each vector element even though atomicity is not required.

  - Offload some of the computation to the L2 cache.

  - Saves reading the value of `a[i]` in registers.

  - This reduces the latency to compute each element of the vector.

  - Can result in **more bytes in-flight**!

```
__global__
void vector(const float * __restrict__ b,
            const float * __restrict__ c,
            float * __restrict__ a,
            int N) {
  int grid_tid = cg::this_grid().thread_rank();
  int grid_stride = cg::this_grid().num_threads();

  for (int i = grid_tid; i < N; i += grid_stride) {
    atomicAdd(&a[i], b[i] * c[i]);
  }
}
```

Memory operations = 2 reads + 1 write

Performance Analysis
NVIDIA H100 SXM

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
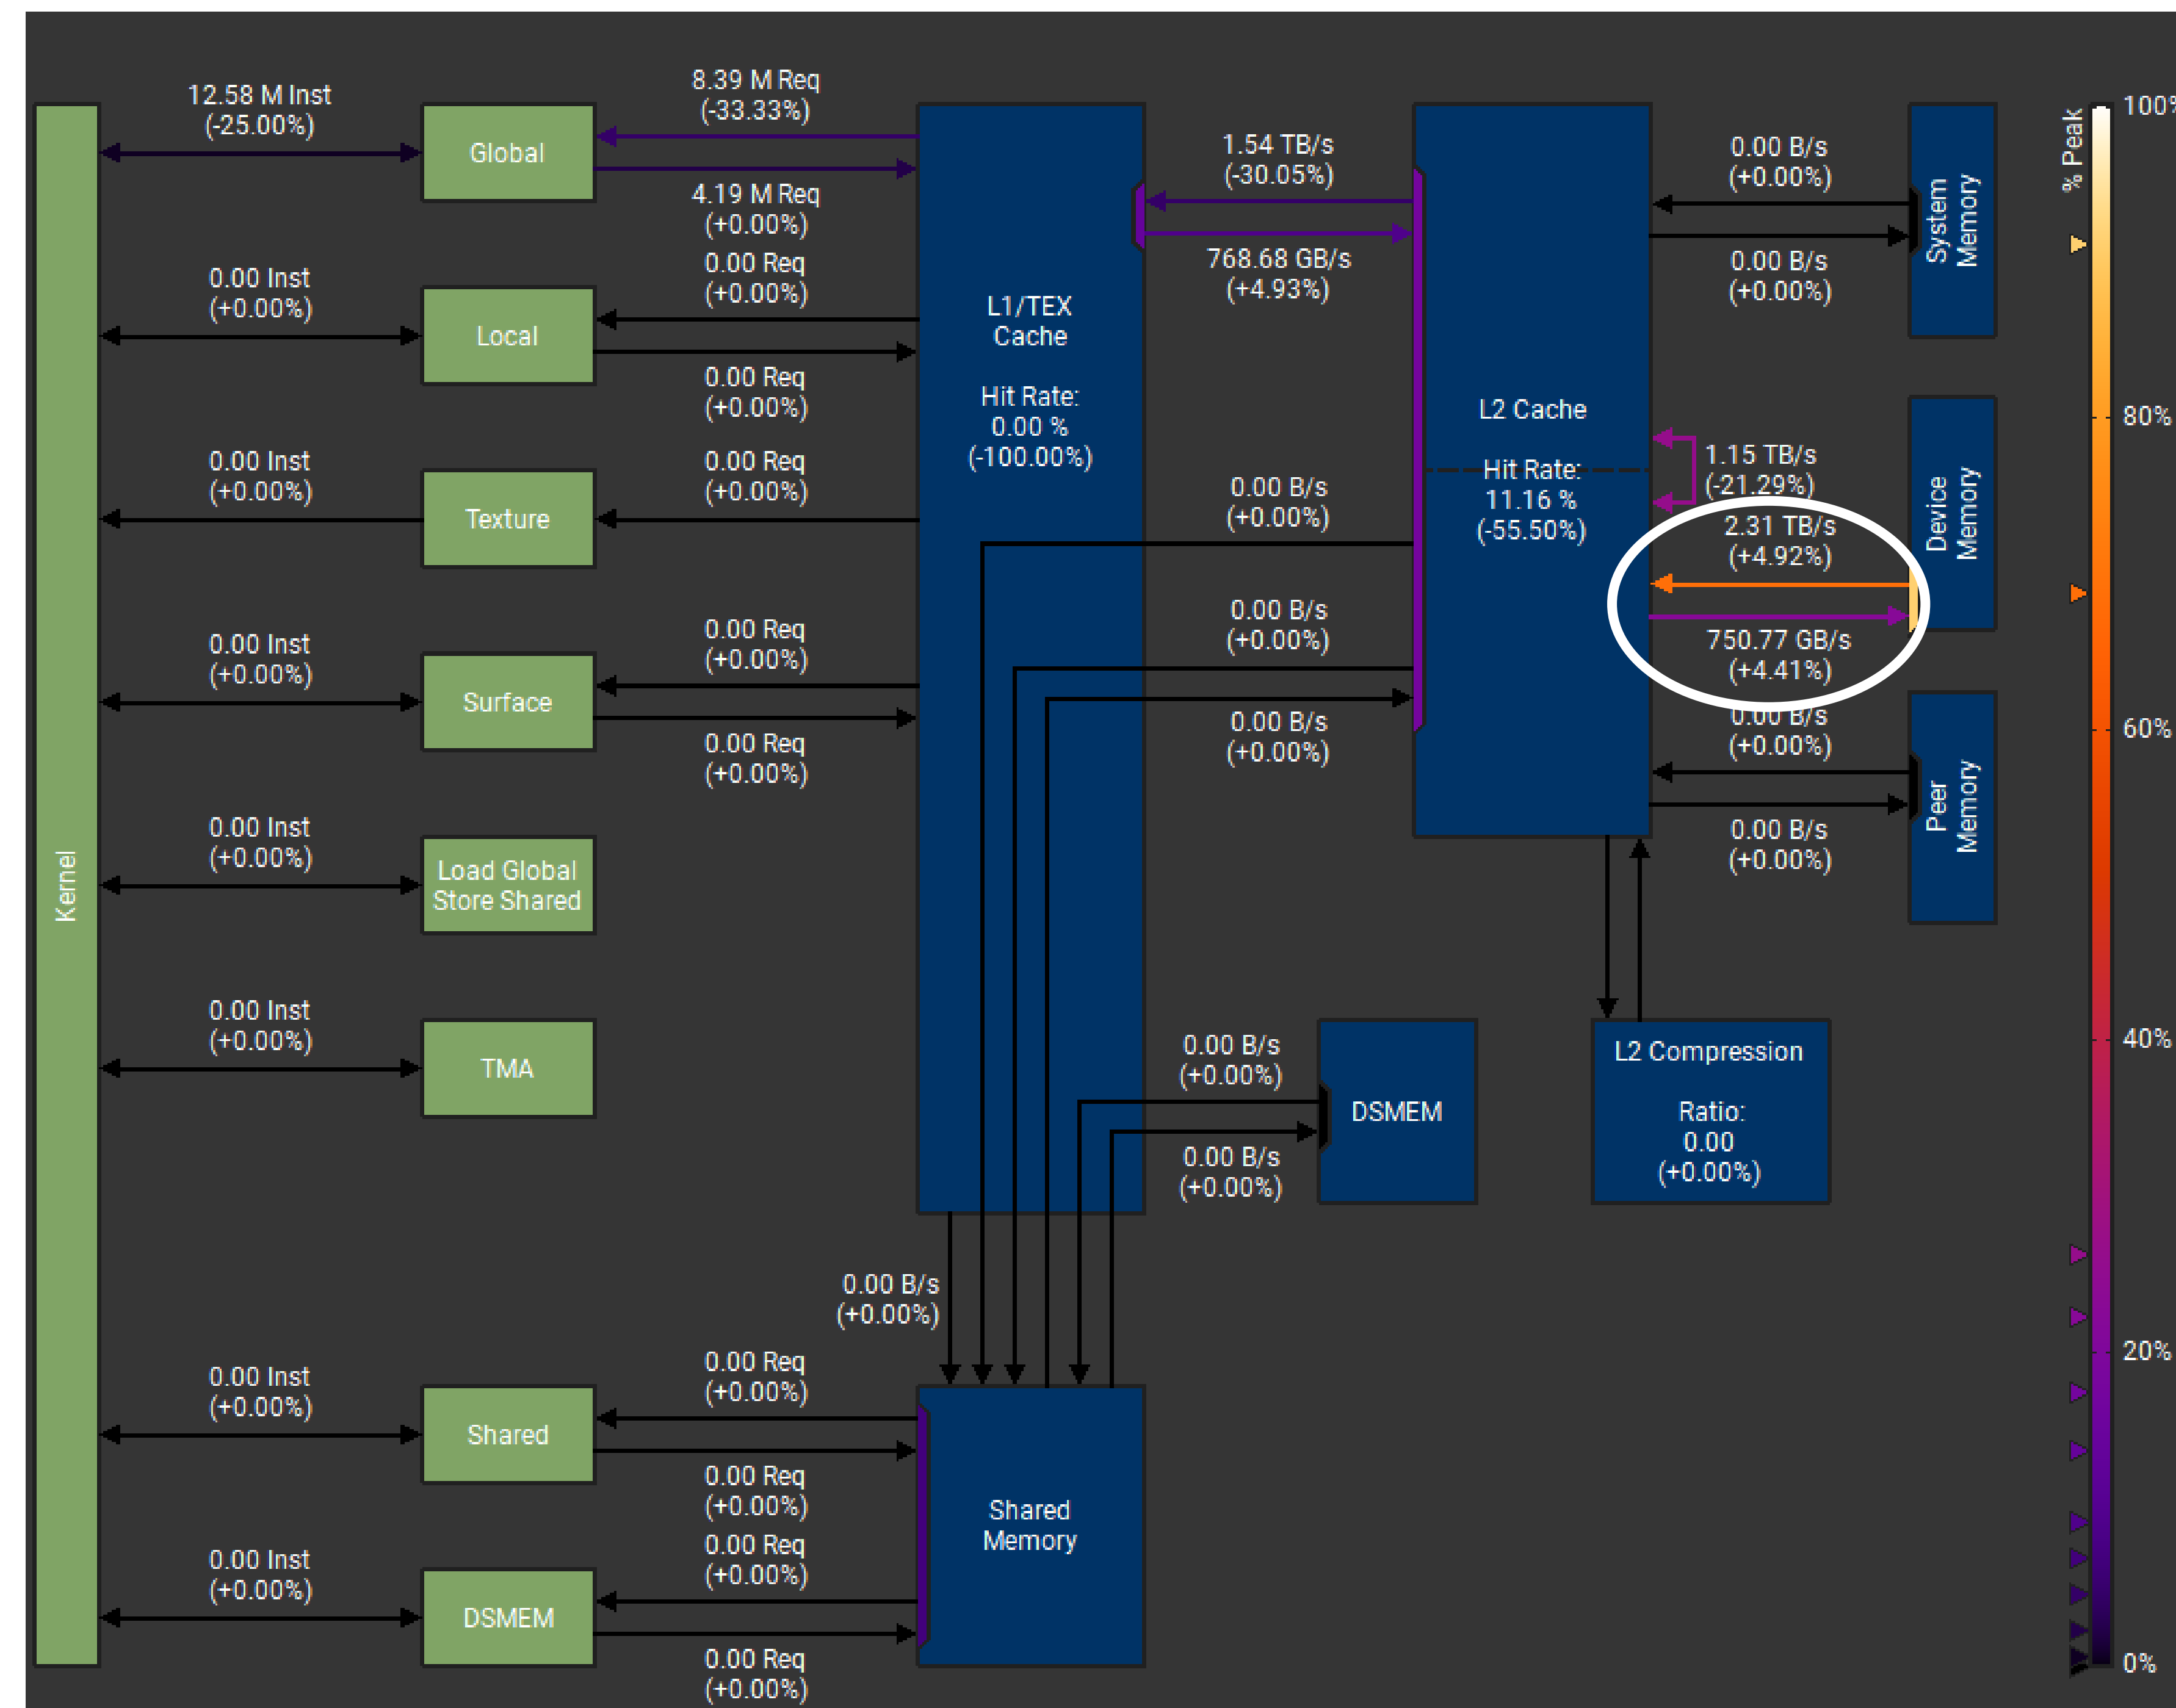 - Problem size = 2^27

NCU Memory Chart (Transfer Size)

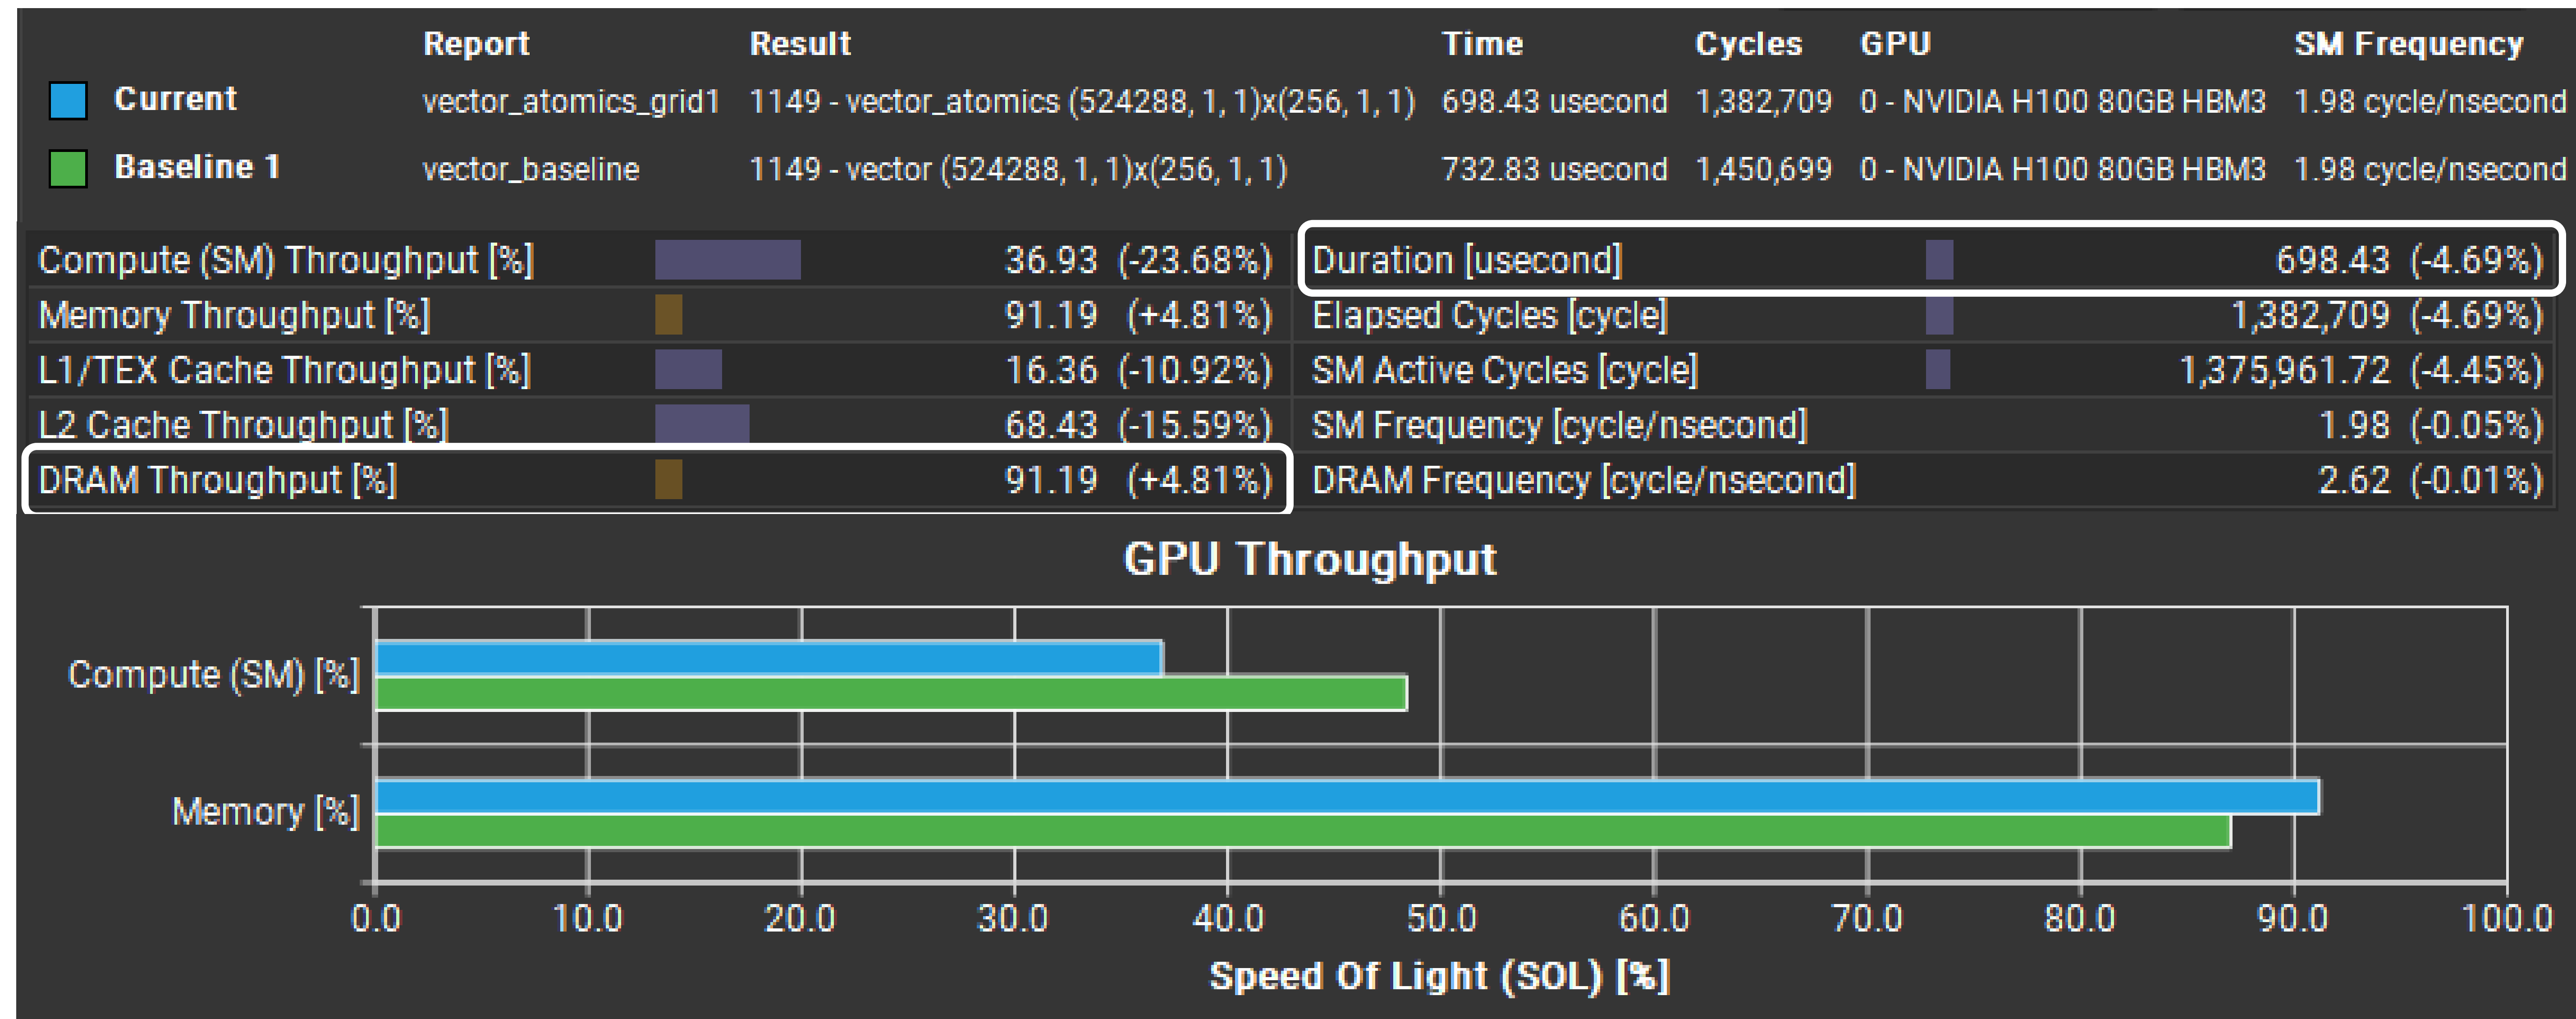# Performance Analysis
NVIDIA H100 SXM

- **Experimental setup:**
  - NVIDIA H100 SXM, 1980 MHz
  - Problem size = 2^27

NCU Memory Chart (Throughput)

Performance Analysis
NVIDIA H100 SXM

	Report	Result	Time	Cycles	GPU	SM Frequency
■ Current	vector_atomics_grid1	1149 - vector_atomics (524288, 1, 1)x(256, 1, 1)	698.43 usecond	1,382,709	0 - NVIDIA H100 80GB HBM3	1.98 cycle/nsecond
■ Baseline 1	vector_baseline	1149 - vector (524288, 1, 1)x(256, 1, 1)	732.83 usecond	1,450,699	0 - NVIDIA H100 80GB HBM3	1.98 cycle/nsecond

Compute (SM) Throughput [%]	36.93 (-23.68%)	Duration [usecond]	698.43 (-4.69%)
Memory Throughput [%]	91.19 (+4.81%)	Elapsed Cycles [cycle]	1,382,709 (-4.69%)
L1/TEX Cache Throughput [%]	16.36 (-10.92%)	SM Active Cycles [cycle]	1,375,961.72 (-4.45%)
L2 Cache Throughput [%]	68.43 (-15.59%)	SM Frequency [cycle/nsecond]	1.98 (-0.05%)
DRAM Throughput [%]	91.19 (+4.81%)	DRAM Frequency [cycle/nsecond]	2.62 (-0.01%)

GPU Throughput

Speed Of Light (SOL) [%]

- ~**5%** increase in memory throughput translates into a corresponding reduction in execution time. **Why?**
  - This kernel is DRAM bandwidth bound.

85

Summary

Which optimizations to focus on?
Solving the bottlenecks

- **Compute bound**
  - Reduce instruction count.
    - E.g., use vector loads/stores.
  - Use tensor cores.
  - Use lower precision arithmetic, fast math intrinsics.

- **Bandwidth bound**
  - Reduce the amount of data transferred
    - Optimize memory access patterns.
    - Lower precision datatypes.
    - Kernel fusion.

- **Latency bound**
  - Increase number of instructions and memory accesses in-flight.
  - Increase parallelism, occupancy.

Resources/Further Study

- **CUDA best practices guide:** https://docs.nvidia.com/cuda/cuda-c-best-practices-guide/

- **CUDA samples:** https://github.com/NVIDIA/cuda-samples

- **GTC'24 sessions:** https://www.nvidia.com/gtc/sessions/performance-optimization/
  - Advanced Performance Optimization in CUDA [S62192]
  - Performance Optimization for Grace CPU Superchip [S62275]
  - Grace Hopper Superchip Architecture and Performance Optimizations for Deep Learning Applications [S61159]
  - Multi GPU Programming Models for HPC and AI [S61339]
  - More Data, Faster: GPU Memory Management Best Practices in Python and C++ [S62550]
  - Harnessing Grace Hopper's Capabilities to Accelerate Vector Database Search [S62339]
  - From Scratch to Extreme: Boosting Service Throughput by Dozens of Times with Step-by-Step Optimization [S62410]